

# Hadoop

## Setup and Configuration

By

**CHITRAKANT BANCHHOR**

**IT,MITCOE, Pune**

**[cobanchhor@gmail.com](mailto:cobanchhor@gmail.com)**

# **Hadoop : First step**

**awareness about Distributed computing and system**

# Cluster Computing

# What is Cluster Computing ?

- ❖ A collection of computers configured in such a way that they can be used to solve a problem by means of parallel processing.

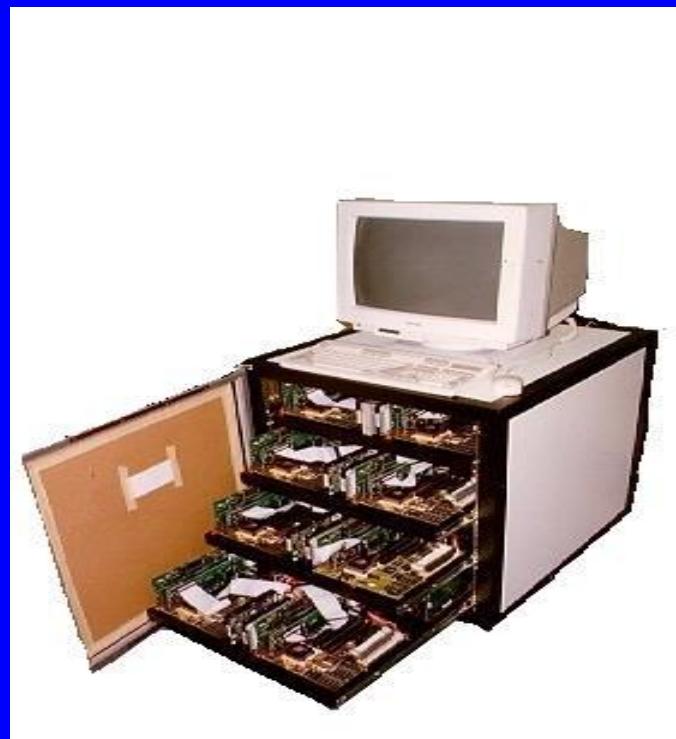
This environment is provided where program part's execute in all the nodes.





# The First Beowulf

- ❖ 16 x 486DX4, 100MHz processors
- ❖ 16MB of RAM each, 256MB in total
- ❖ Ethernet (2 x 10Mbps)



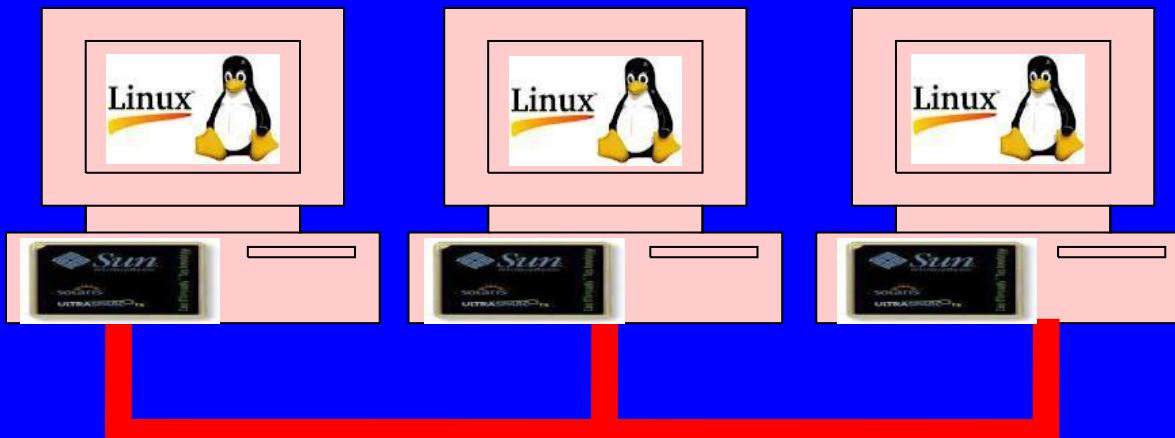
# Current Beowulf

- ❖ Faster processors, faster interconnect, but the idea remains the same
- ❖ Top cluster: 1.433 TFLOPS peak



# Homogeneous Cluster

- ❖ The cluster in which every single node is exactly same, from the motherboard and the memory, to the disk drives and the NIC .

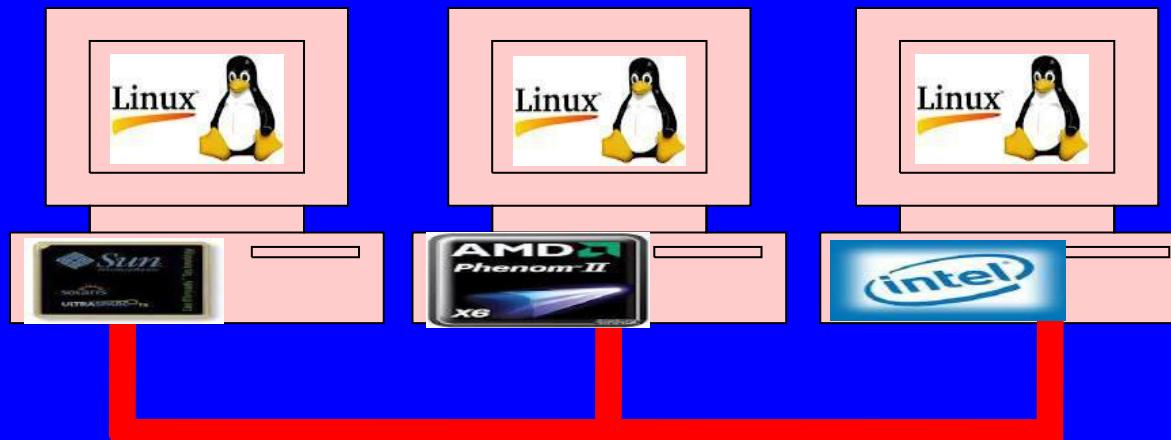


- ❖ Advantages:

1. The use of homogeneous system will ensure that software will operate same on each node.
2. The nodes within the system are interchangeable

# Heterogeneous Cluster

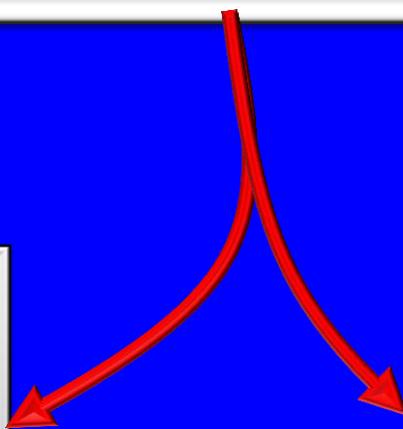
- Made from different kinds of computers, SPARC, DEC ALPHA



## **Basic Parallelism techniques**

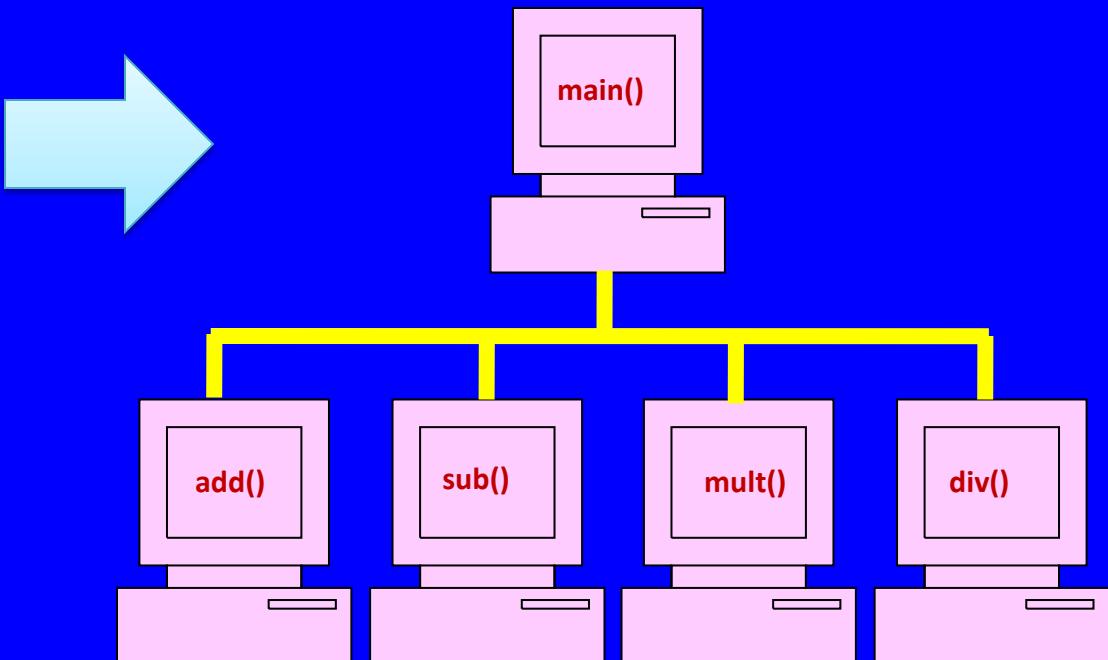
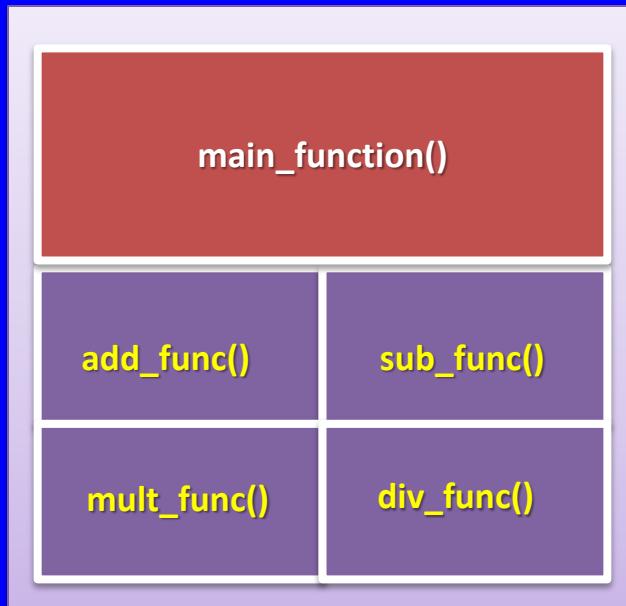
**Function – Level  
parallelism**

**Data- Level  
parallelism**



# Functional parallelism

- ❖ An application is parallelized along its functions



## Data Parallelism (SPMD)

- ❖ All the tasks are same, but each one only solves the small part of the data.

**MATRIX - A (100 X 100)**

A(0,0), .....A(0,99)

|

|

A(99,0).....A(99,99)

X

**MATRIX - B (100 X 100)**

B(0,0), .....B(0,99)

|

|

B(99,0).....B(99,99)

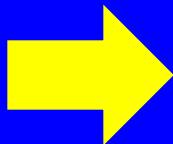
## MATRIX - A (100 X 100)

A(0,0), .....A(0,99)

|

|

A(99,0).....A(99,99)



A(0,0) .. A(0,24)

A(24,0) .. A(24,24)

A(50,50) .. A(50,74)

A(74,50) .. A(74,74)

A(25,25) .. A(25,49)

A(49,25) .. A(49,49)

A(75,75) .. A(75,99)

A(99,75) .. A(99,99)

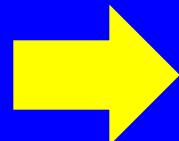
## MATRIX - B (100 X 100)

B(0,0), ..... B(0,99)

|

|

B(99,0).....B(99,99)



B(0,0) .. B(0,24)

B(24,0) .. B(24,24)

B(50,50) .. B(50,74)

B(74,50) .. B(74,74)

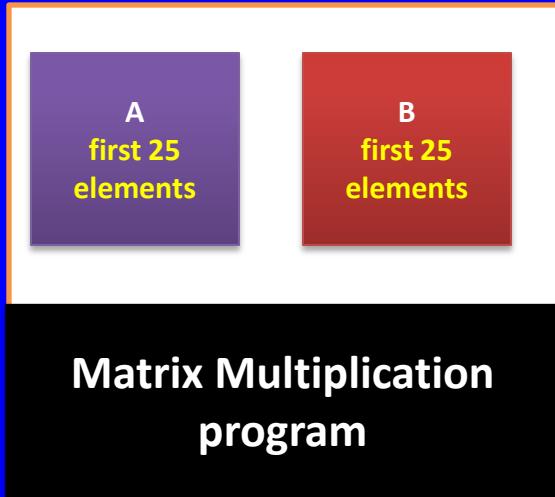
B(25,25) .. B(25,49)

B(49,25) .. B(49,49)

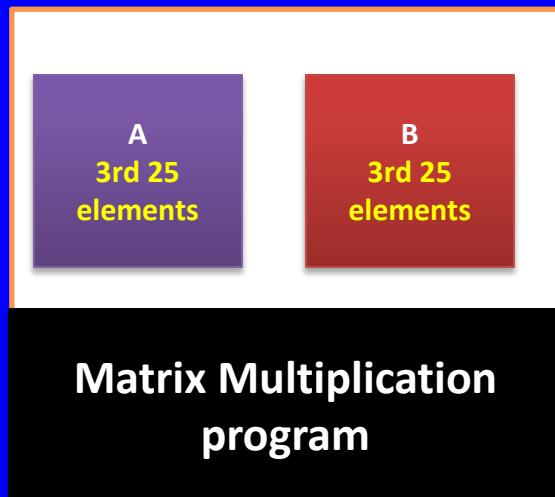
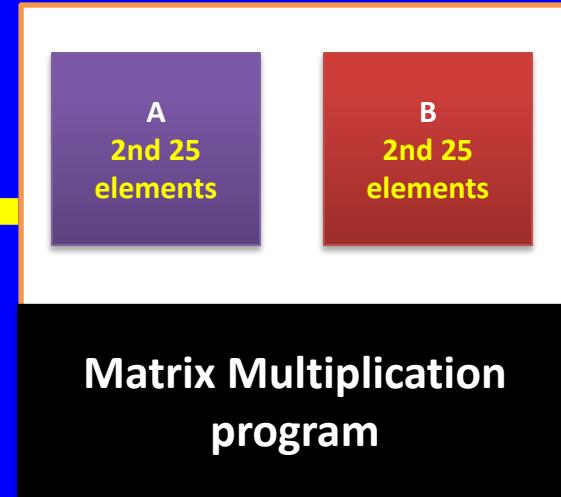
B(75,75) .. B(75,99)

B(99,75) .. B(99,99)

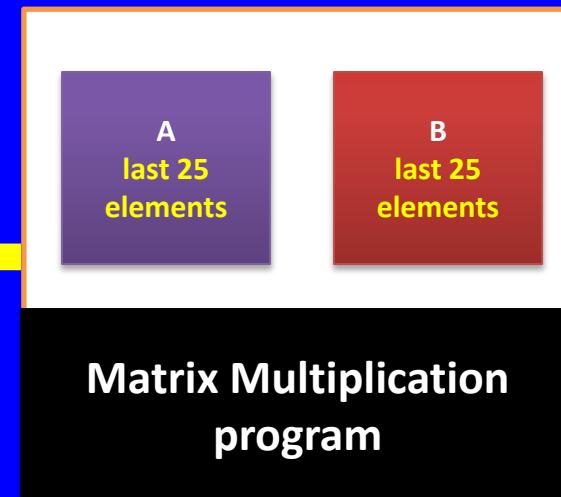
**Node - 1**



**Node - 2**



**Node - 3**



**Node - 4**

# Common parallel programming Paradigm

## ❖ Master – slave model

**master.c**

1. Get input data typically from file
2. Divide the data in separate chunks
3. Distribute each chunk into different slave programs in different nodes.
4. Wait for intermediate results from various slaves.
5. Combine these intermediate results and produce the final result

**slave.c**

1. Get data upon which algorithm operates from the master.
2. Performs defined task on the set of data.
3. Return intermediate result to master.

# **Step - I**

# Hadoop : Second step

## Background HDFS and MapReduce

# Hadoop

# An Introduction

# References

1. Chuck Lam, "Hadoop in Action", December, 2010, Manning publication
2. Alex Holmes, "Hadoop in Practice", October", 2012, Manning publication
3. Tom White, "Hadoop: The Definitive Guide", May 2009, O'Reilly Media
4. Eric Sammer, "Hadoop Operations", September 2012, O'Reilly Media

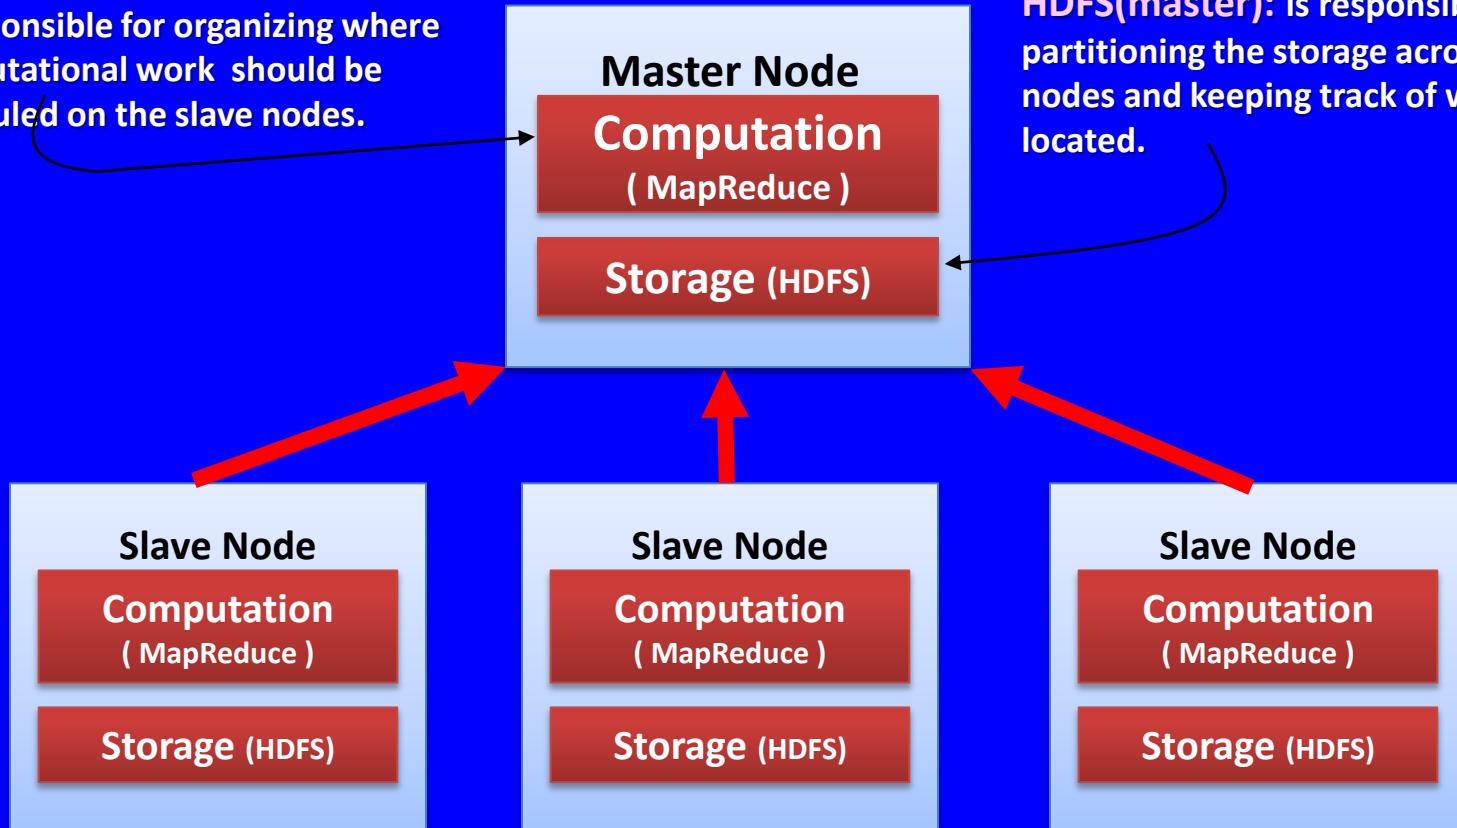
# What is Hadoop?

- ❖ **Software Library:** Hadoop is a software library for distributed computing.
- ❖ **Distributed Storage:** Hadoop platform provides distributed storage.
- ❖ **Computational Capabilities:** Hadoop platform provides distributed computational capabilities.

# High-level Hadoop architecture

## MapReduce(master):

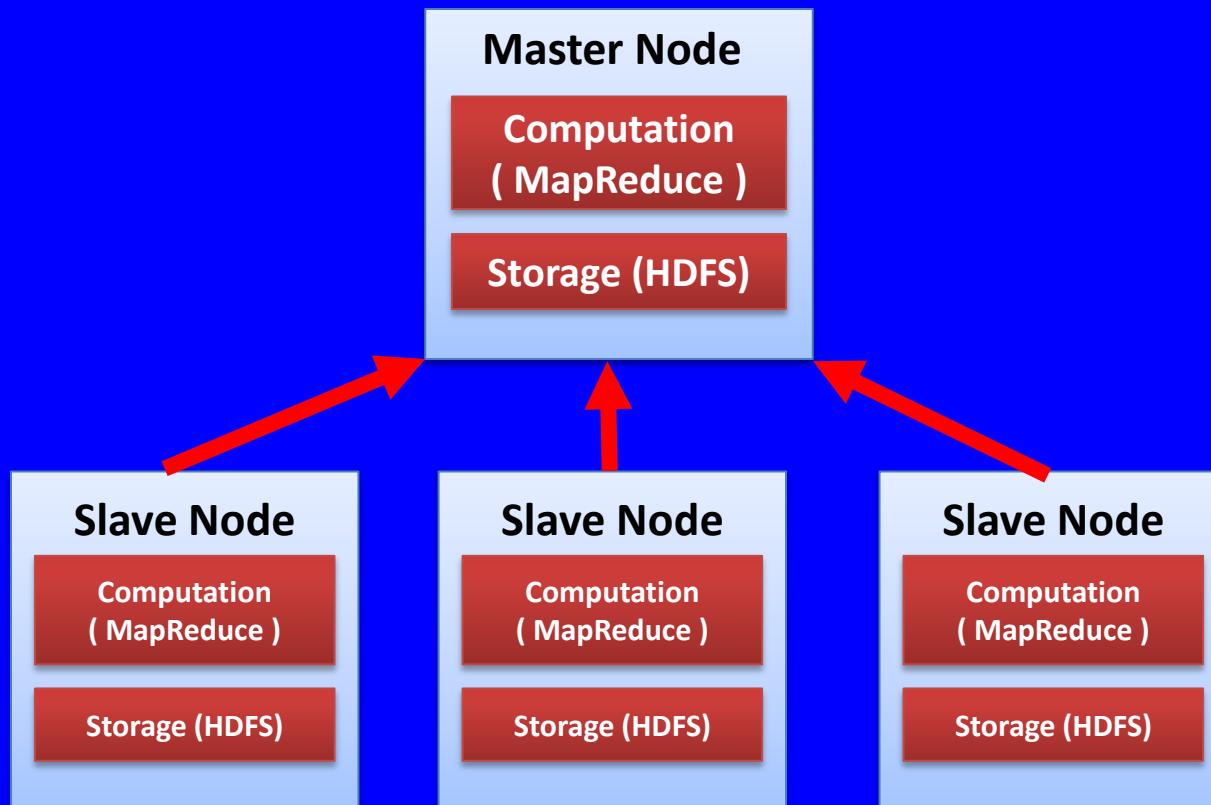
is responsible for organizing where computational work should be scheduled on the slave nodes.



- ❖ Hadoop is a distributed master-slave architecture.

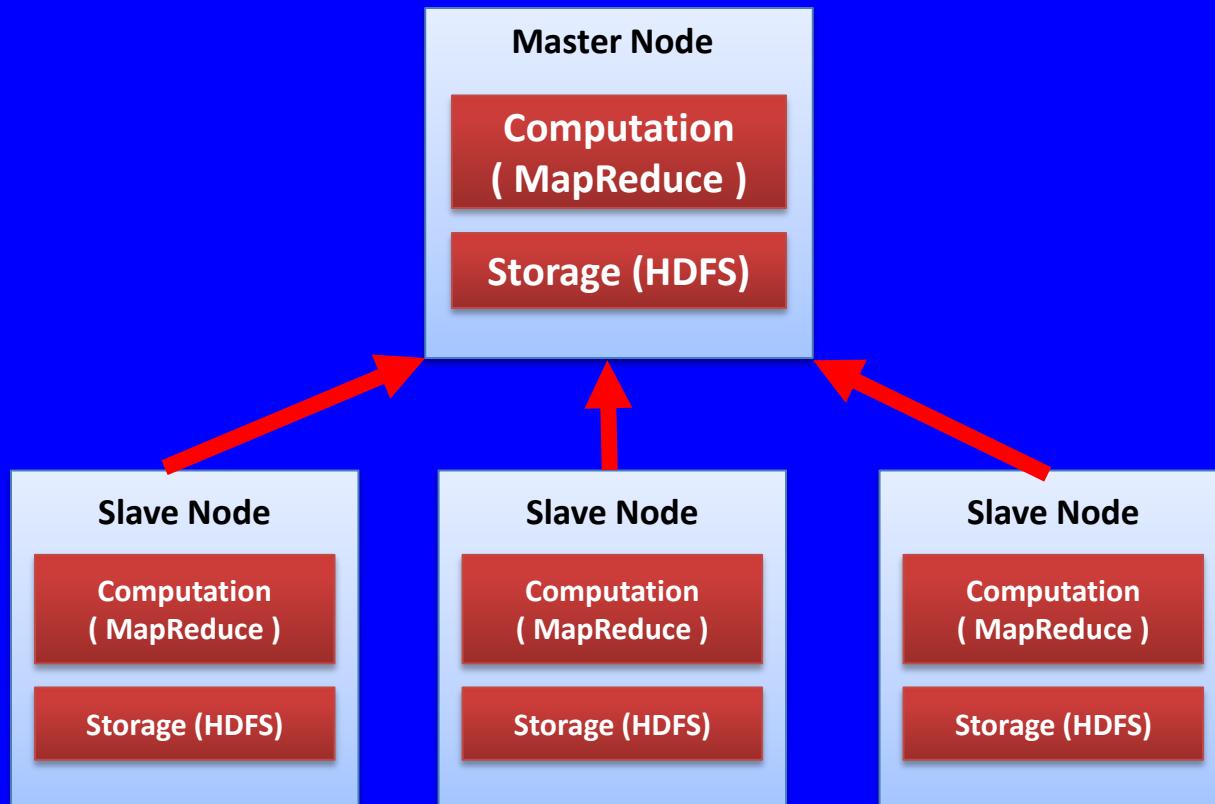
## ❖ Hadoop consists :

- HDFS for storage
- Map-Reduce for computational capabilities.



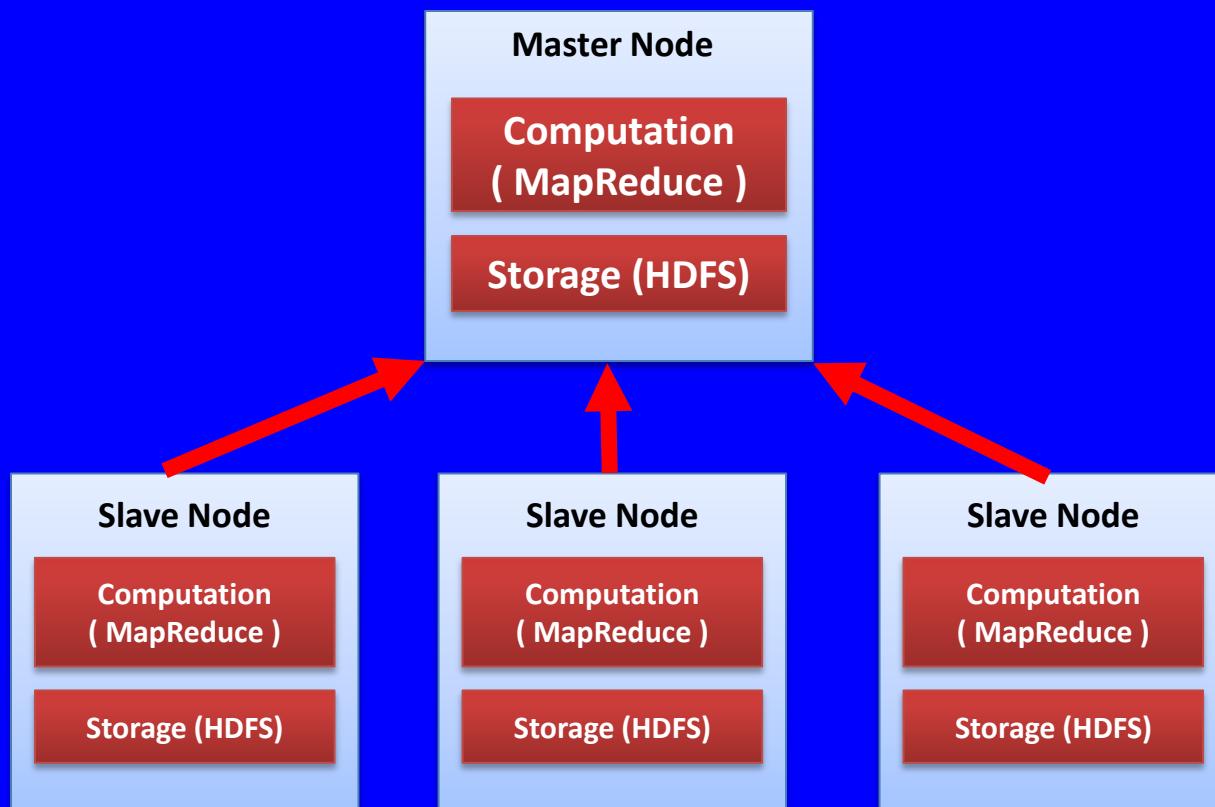
## ❖ HDFS(master):

- **Partitions the overall storage across the slave nodes.**
- **Keeps track of locations of data.**

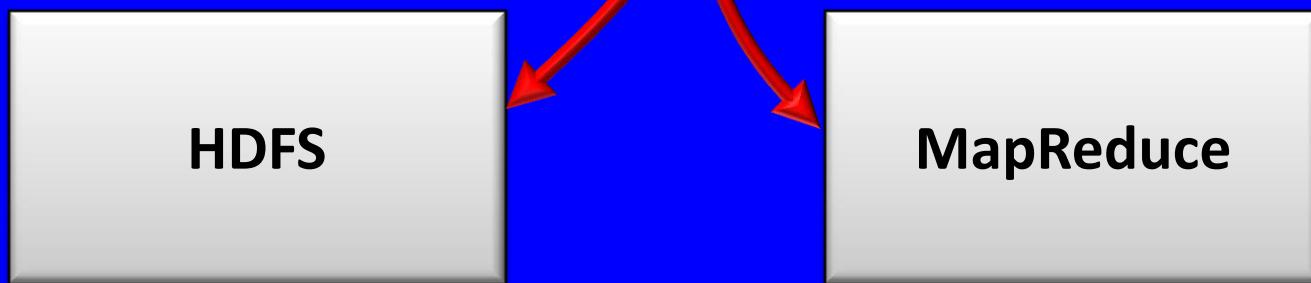


## ❖ MapReduce(master):

- Organizing slave nodes to schedule computational tasks.



## Hadoop: Core components



# Hadoop Distributed File System ( HDFS )

# HDFS

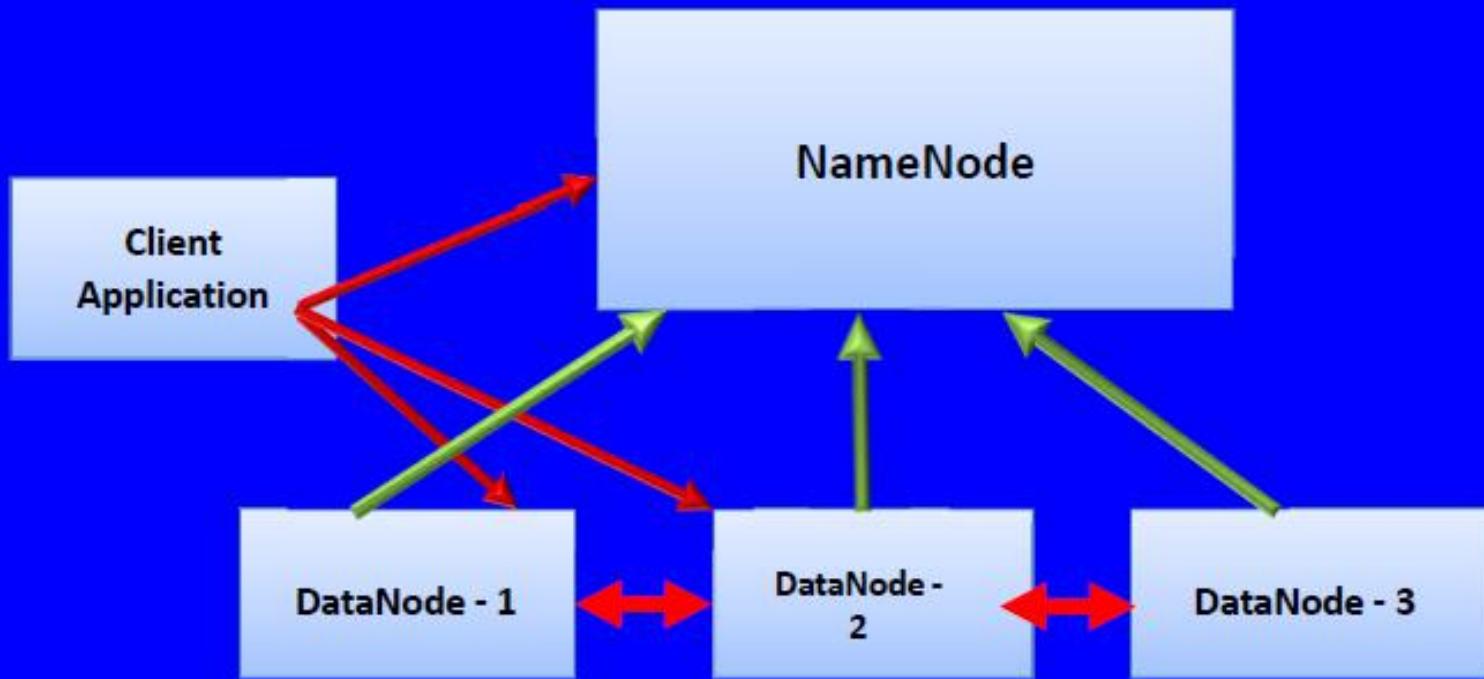
- HDFS: **Hadoop's implementation of distributed file system.**
- **File system facilities:** **hdfs provides all facilities of a file system**

# HDFS Characteristics

1. **large files:** works best when reading and writing large files (gigabytes and larger).
2. **Scalability:** achieved in part due to data replication.
3. **High availability:** achieved in part due to fault tolerance.

# HDFS architecture

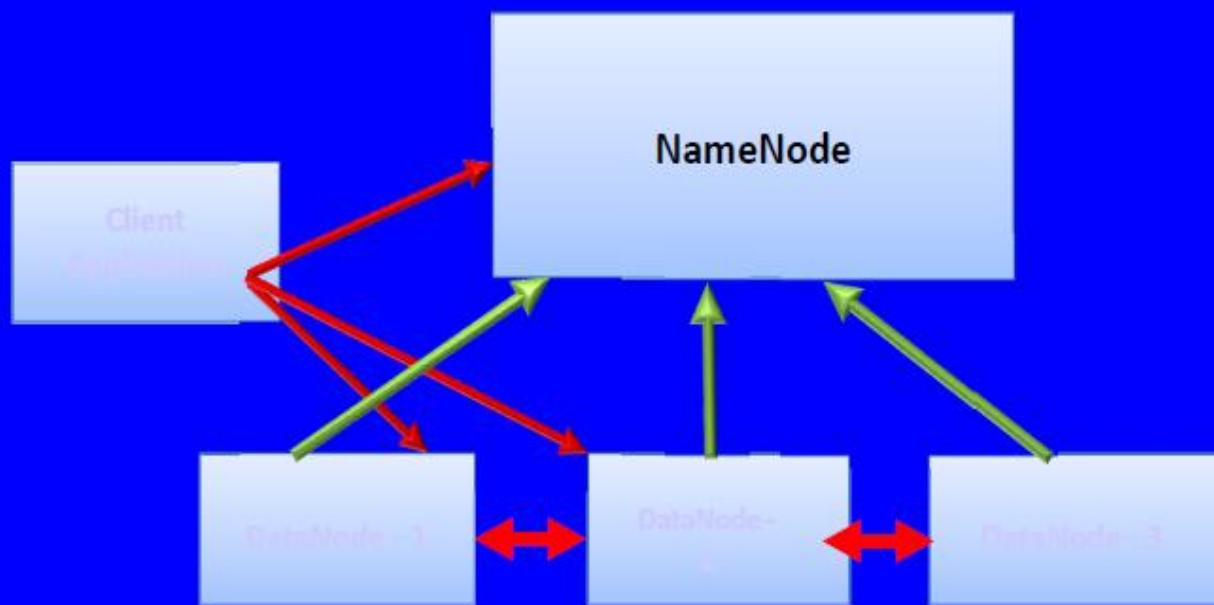
- Consider an HDFS client communicating with the master NameNode and slave DataNodes.

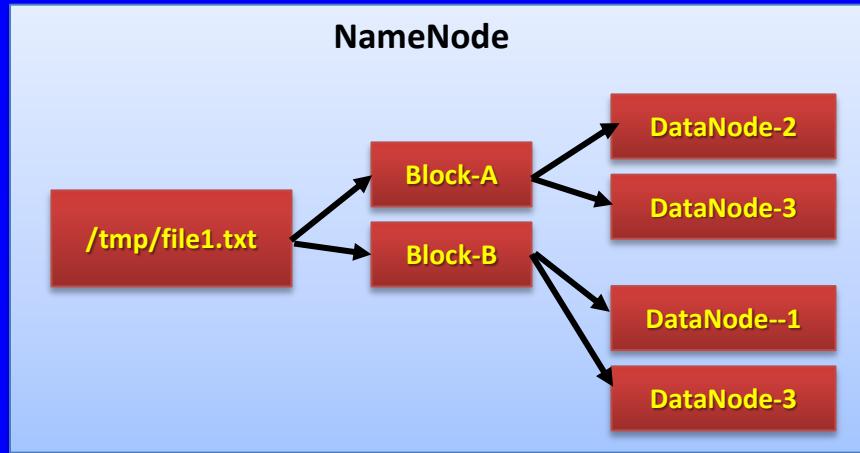


- Master-Slave Architecture

# HDFS NameNode

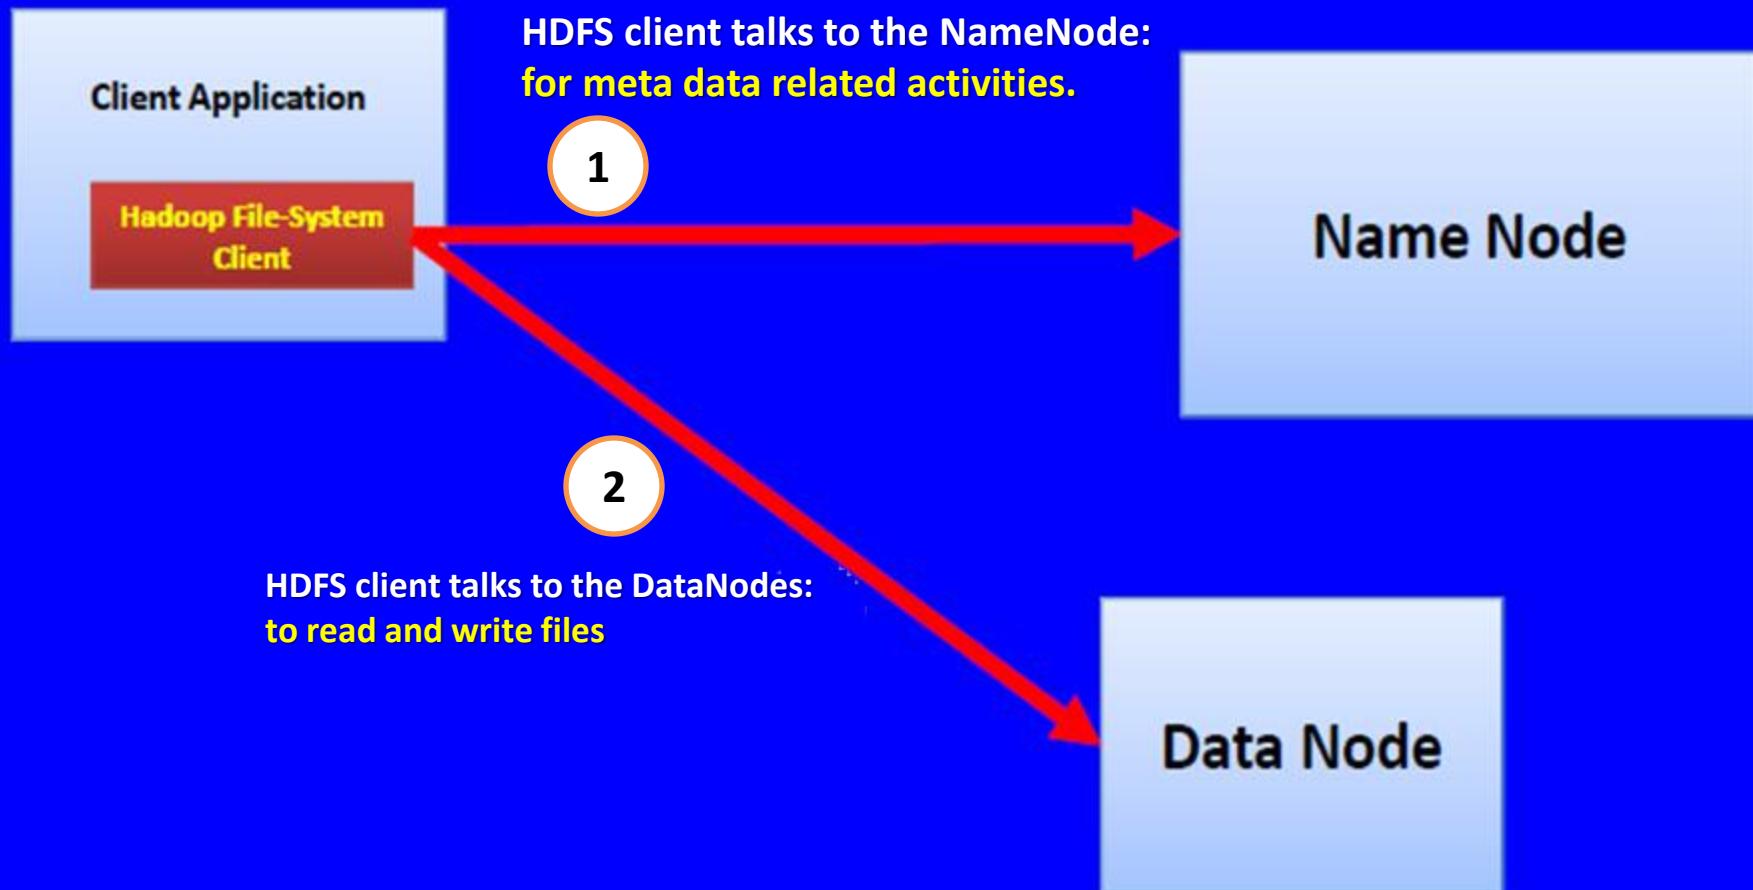
- HDFS cluster consists of a single Namenode.
- It is a master server.
- ❖ Single point of failure : if namenode goes down - whole filesystem is offline



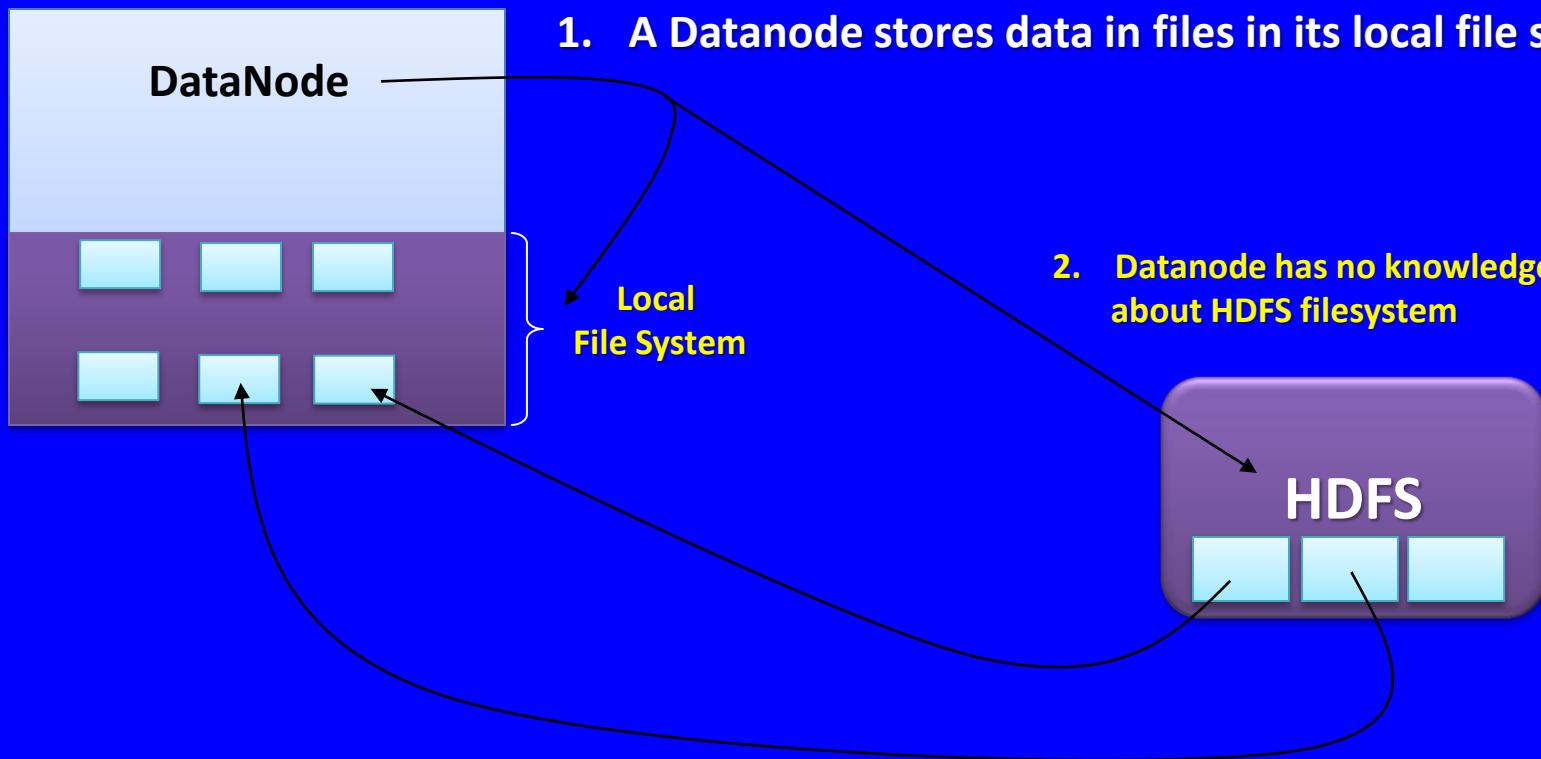


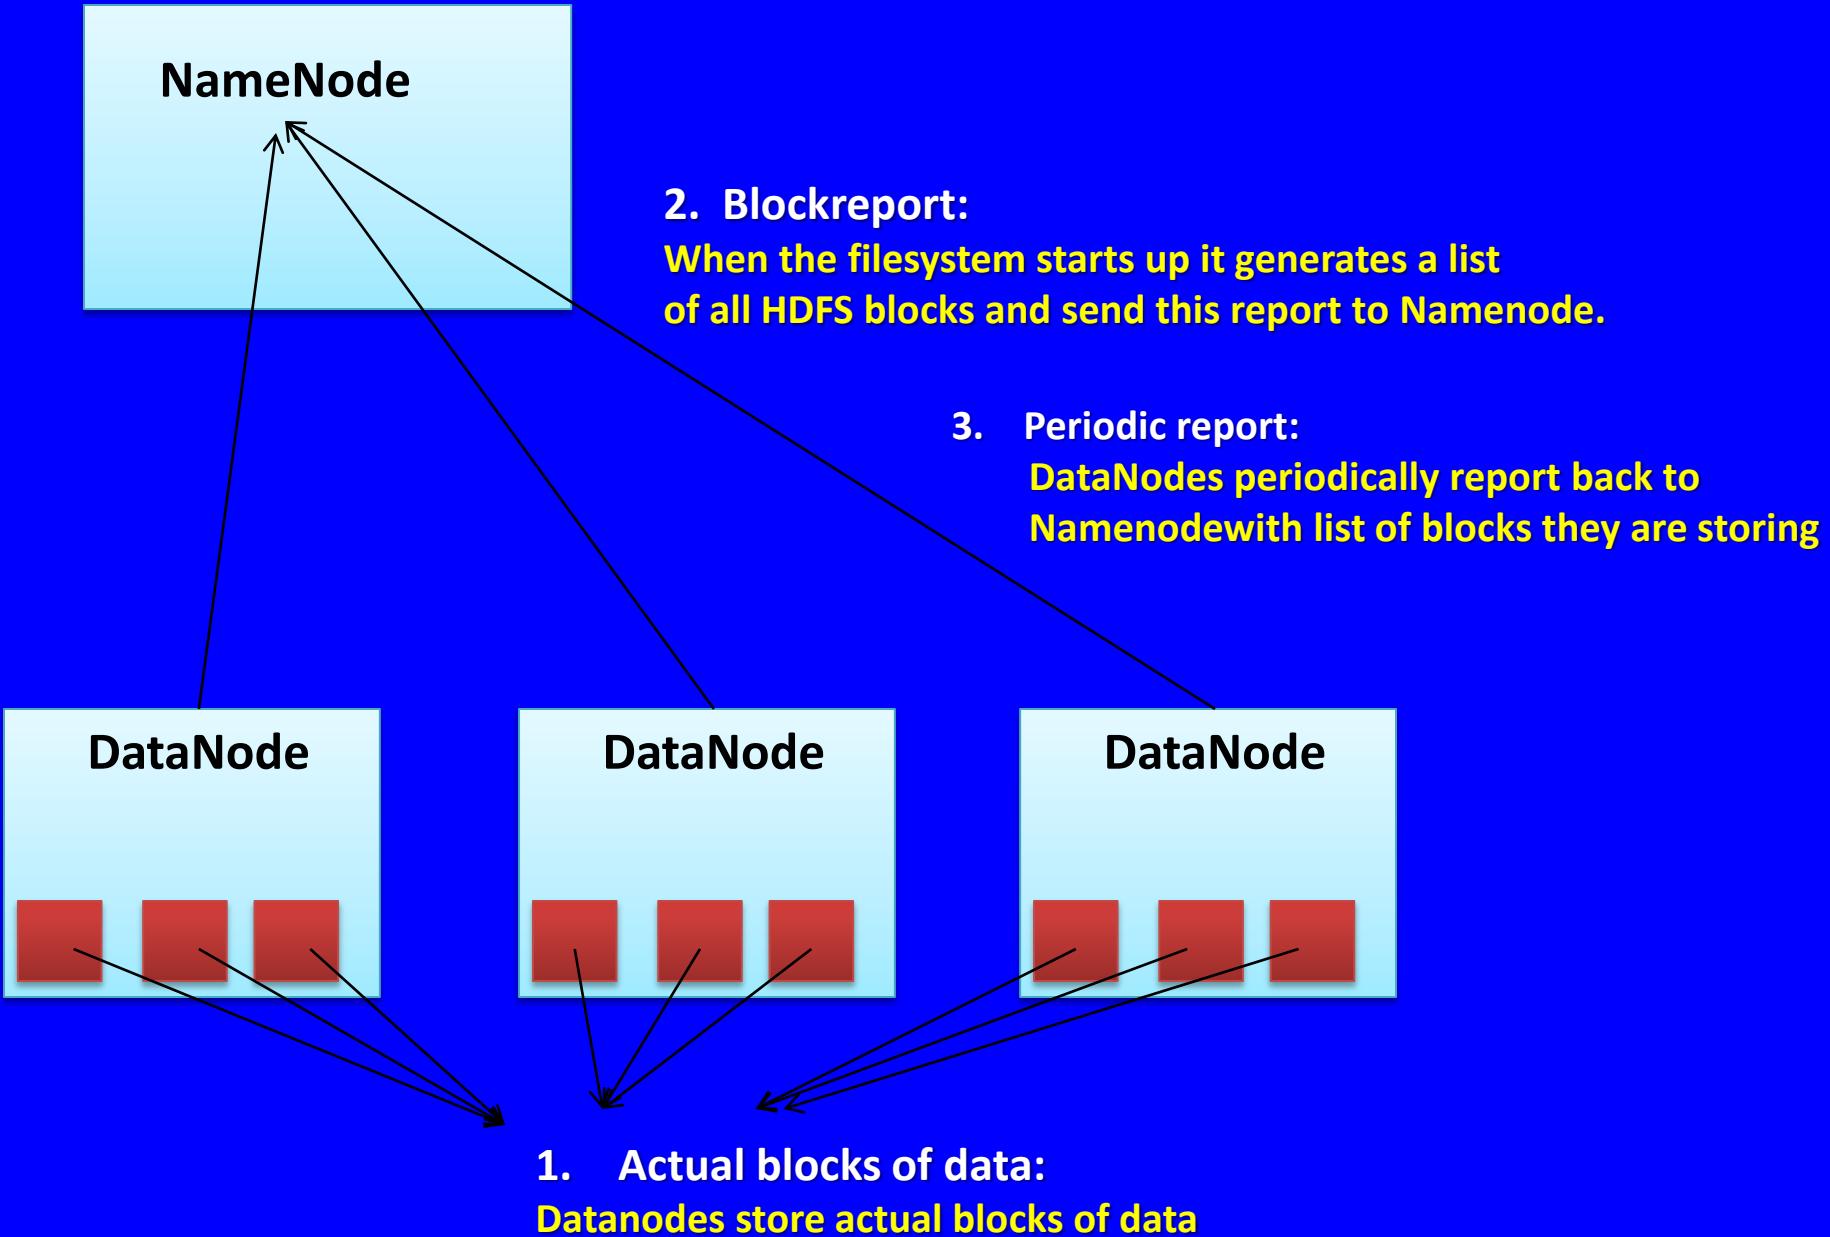
- ❖ NameNode contains: (The meta data about the file system )
  1. Information regarding a block's location.
  2. The information of entire directory structure and files
- ❖ Meta data about the file system: which data nodes manage blocks for each file

## ❖HDFS Client

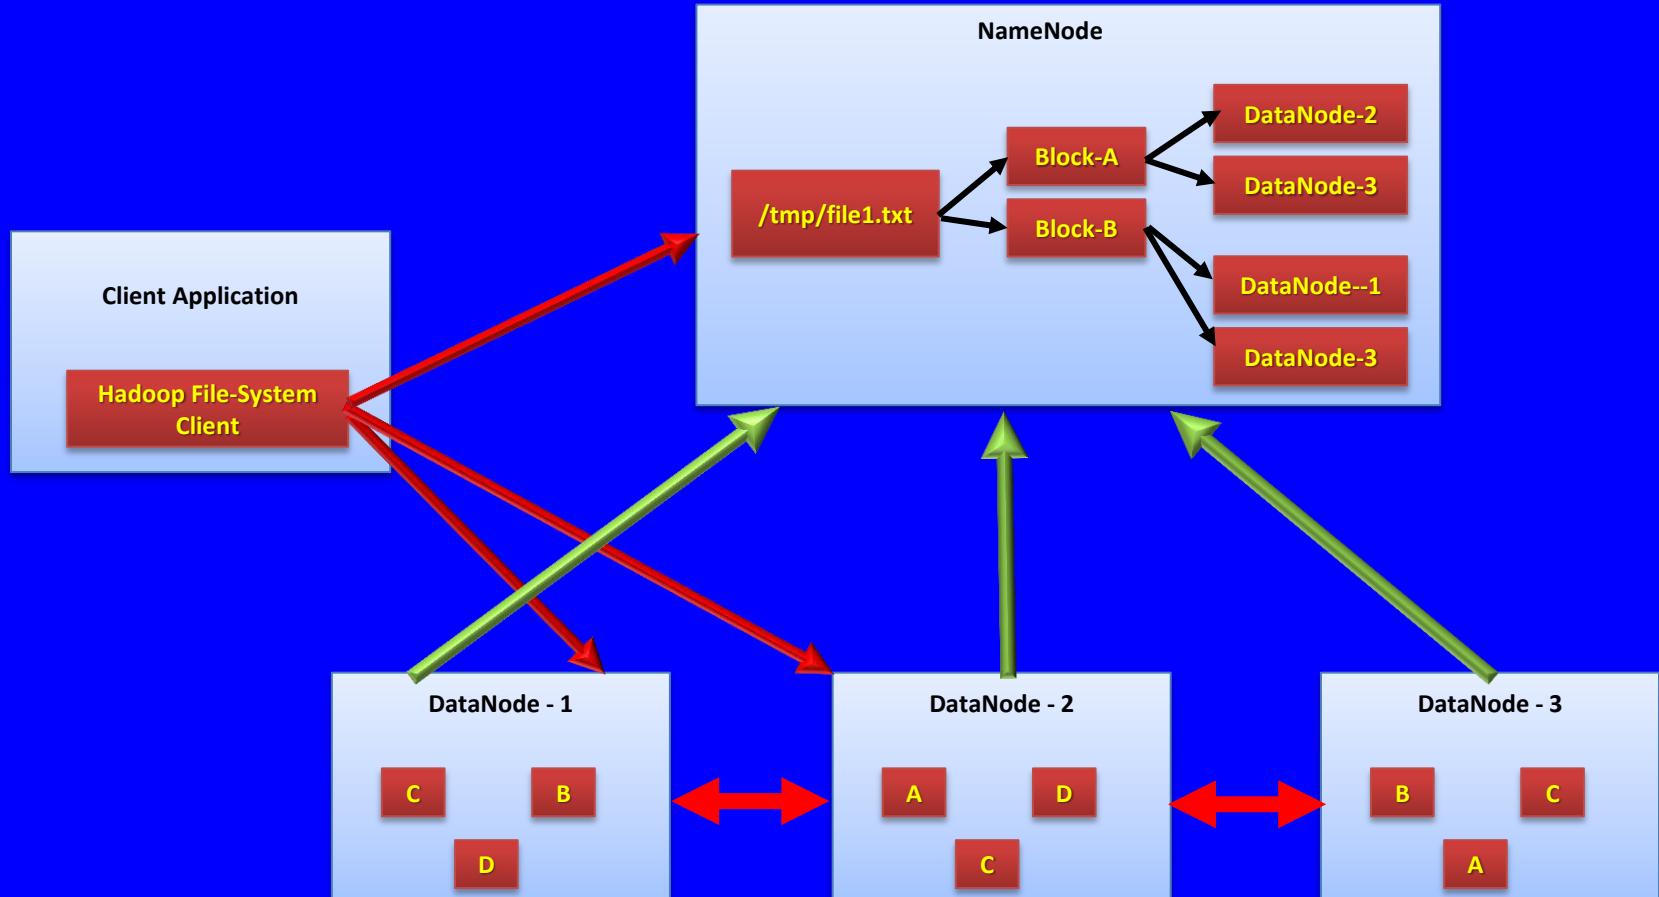


# HDFS DataNodes





- ❖ Logical representation of components: (**NameNode**, **DataNode**, **ClientApp**)
  - Figure shows a logical representation of the components in HDFS:



# Map-Reduce

# MapReduce Algorithm

# Overview

## MapReduce is a Distributed Data Processing Algorithm

### MapReduce Algorithm

- Useful to process huge amount of data in parallel
- Reliable and efficient way in cluster environments

# MapReduce: Internals

## Divide & Conquer

- MapReduce technique for processing of large amount of data

- ❖ Divides input task into smaller and manageable sub-tasks
- ❖ Subtasks execute independently in-parallel

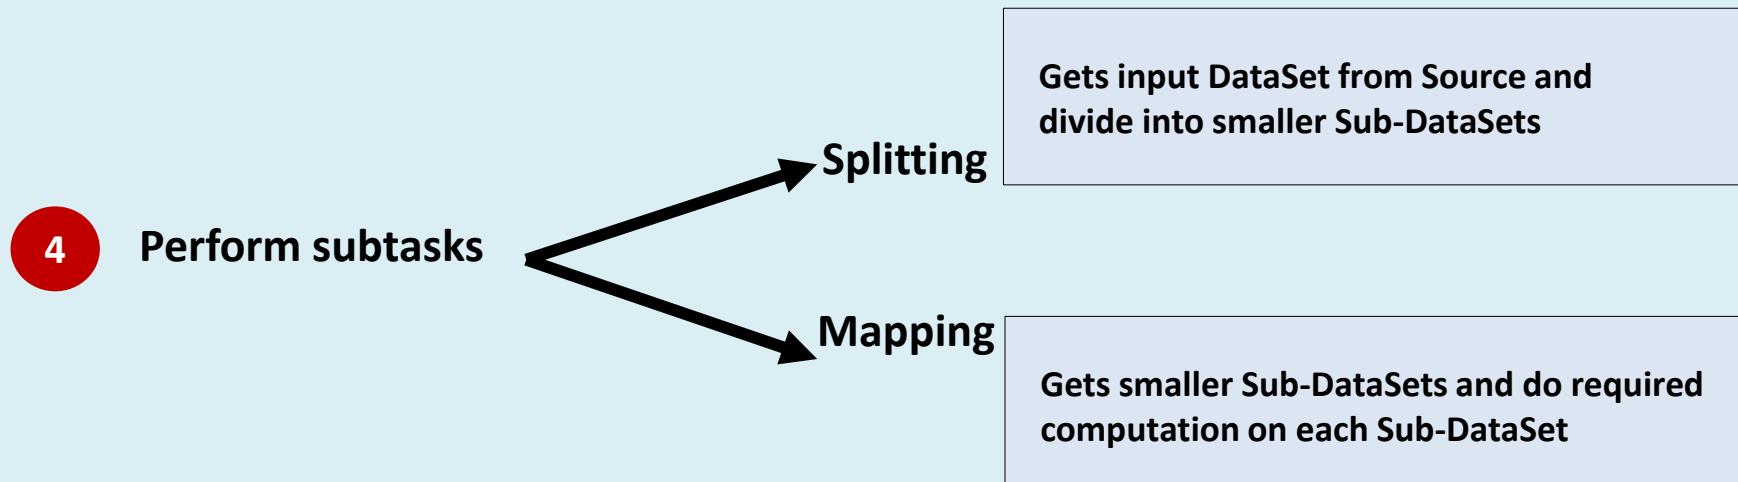
# MapReduce Algorithm Steps

## Main steps

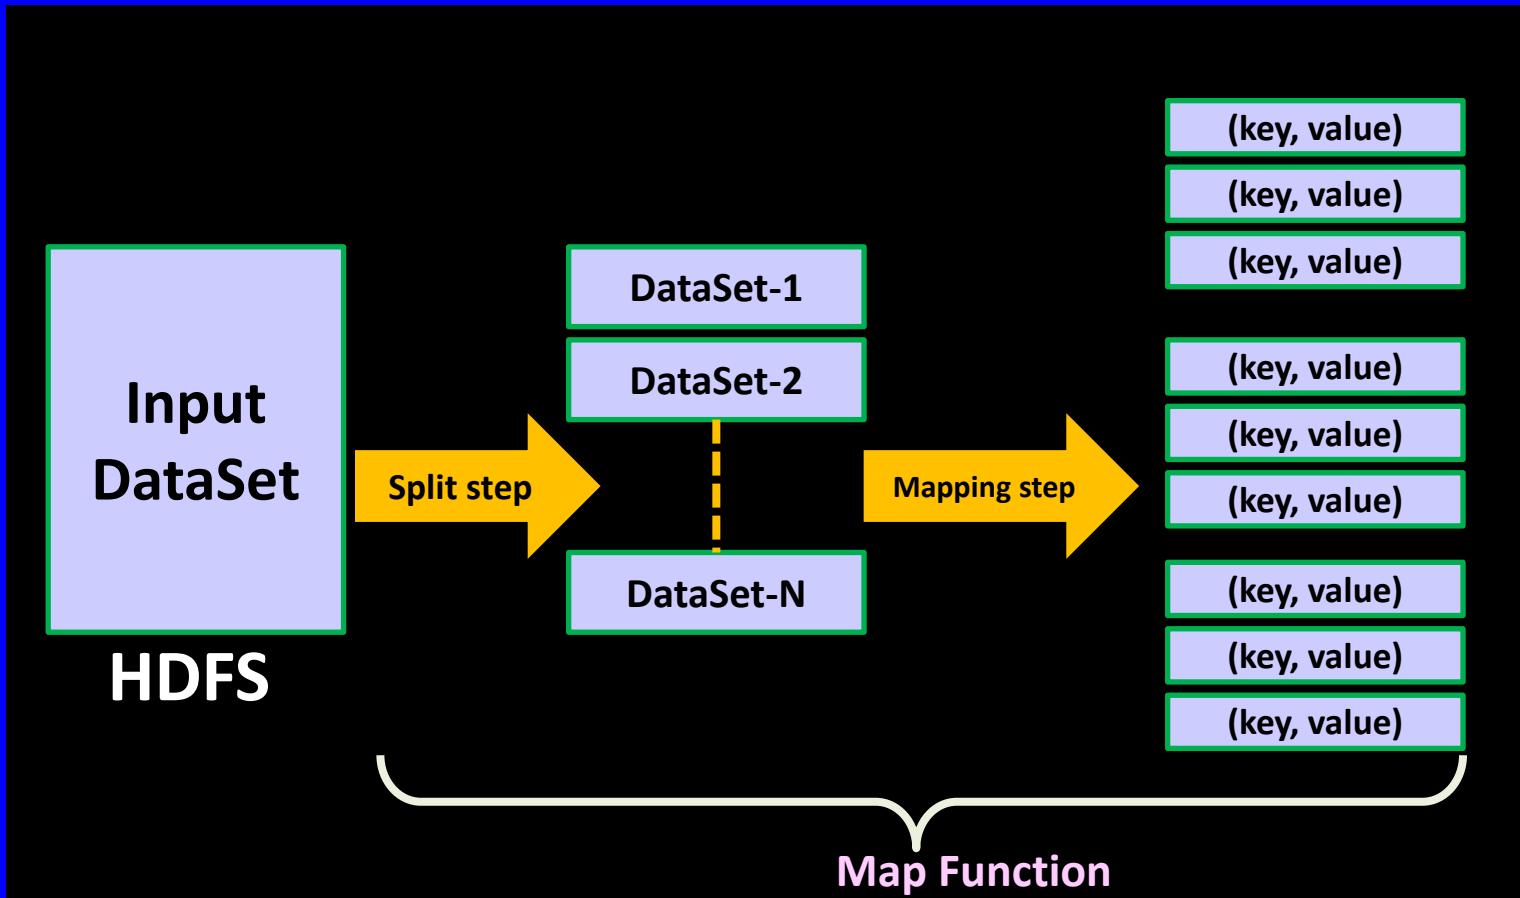
- 1 Map Function**
- 2 Shuffle Function**
- 3 Reduce Function**

# Map Function

- 1 Accepts input tasks (DataSets)
- 2 Divides input tasks into smaller sub-tasks
- 3 Perform required computation on each sub – task in parallel



# Map function output → A set of key and value pairs as <Key, Value>



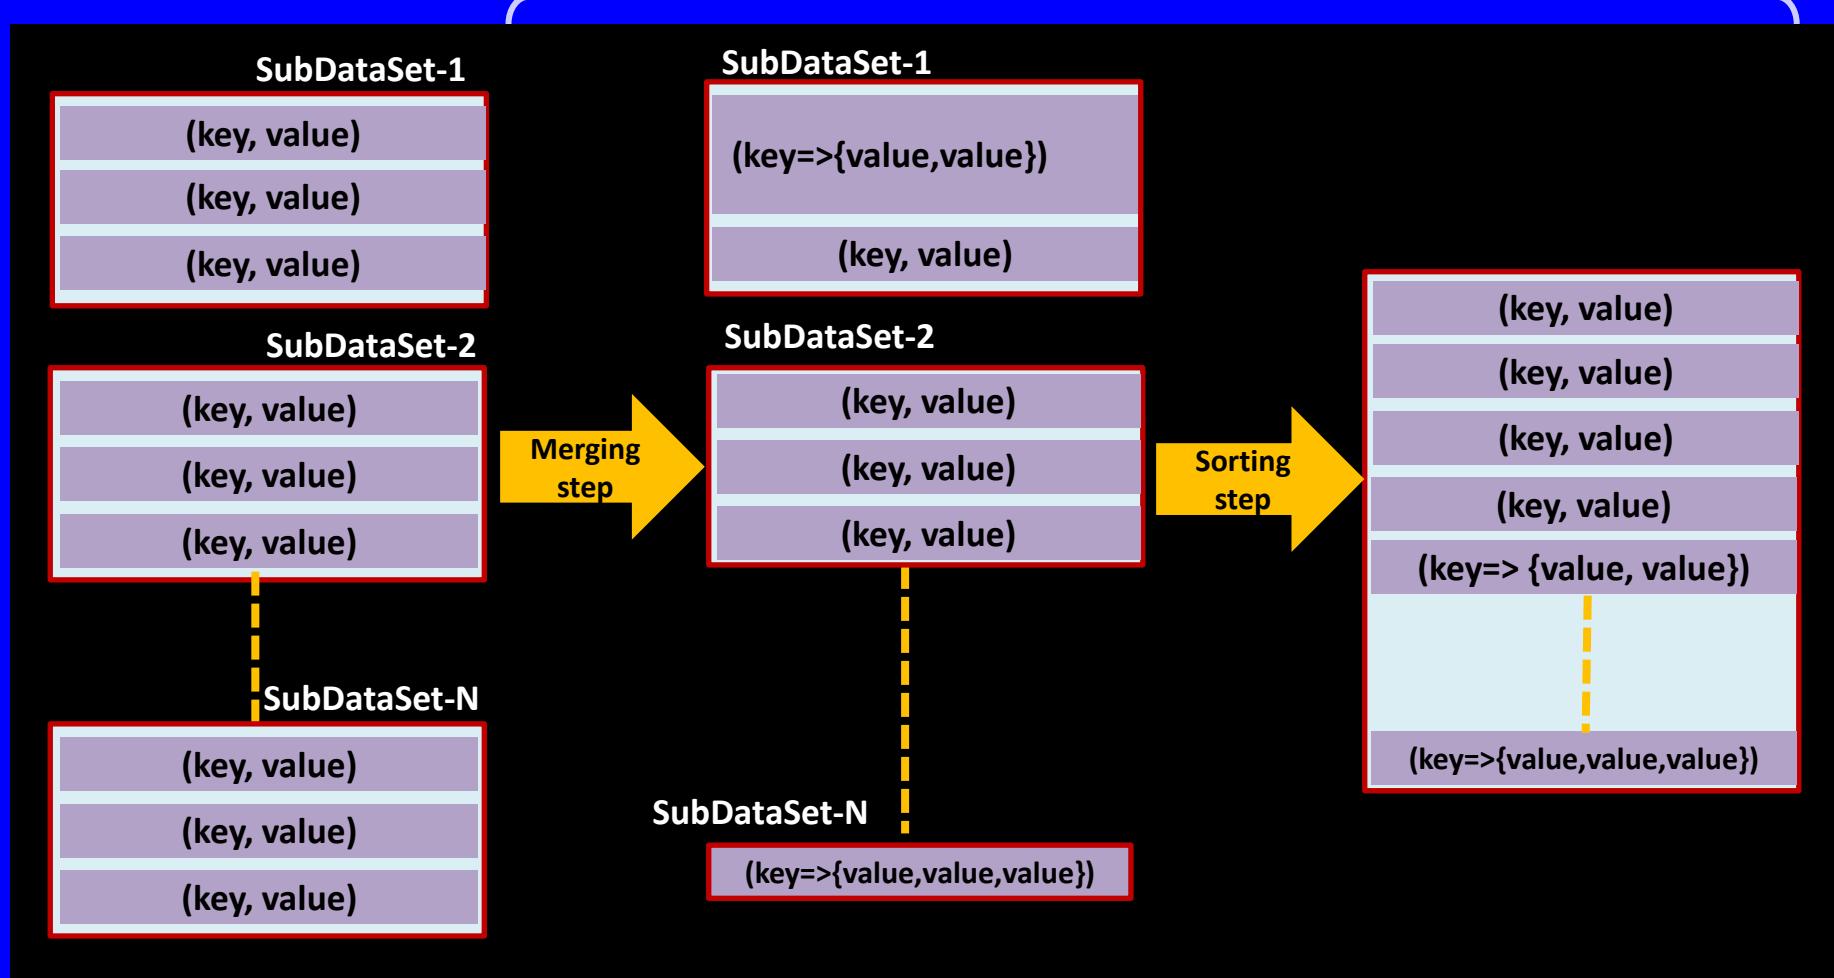
MapReduce First Step Output: **Output of Map Function = List of <key, Value> Pairs**

# Shuffle Function (Combine Function)



- Takes a list of outputs coming from “Map Function”
- Perform two sub-steps (merging, sorting) on each and every key-value pair

## Shuffle Function



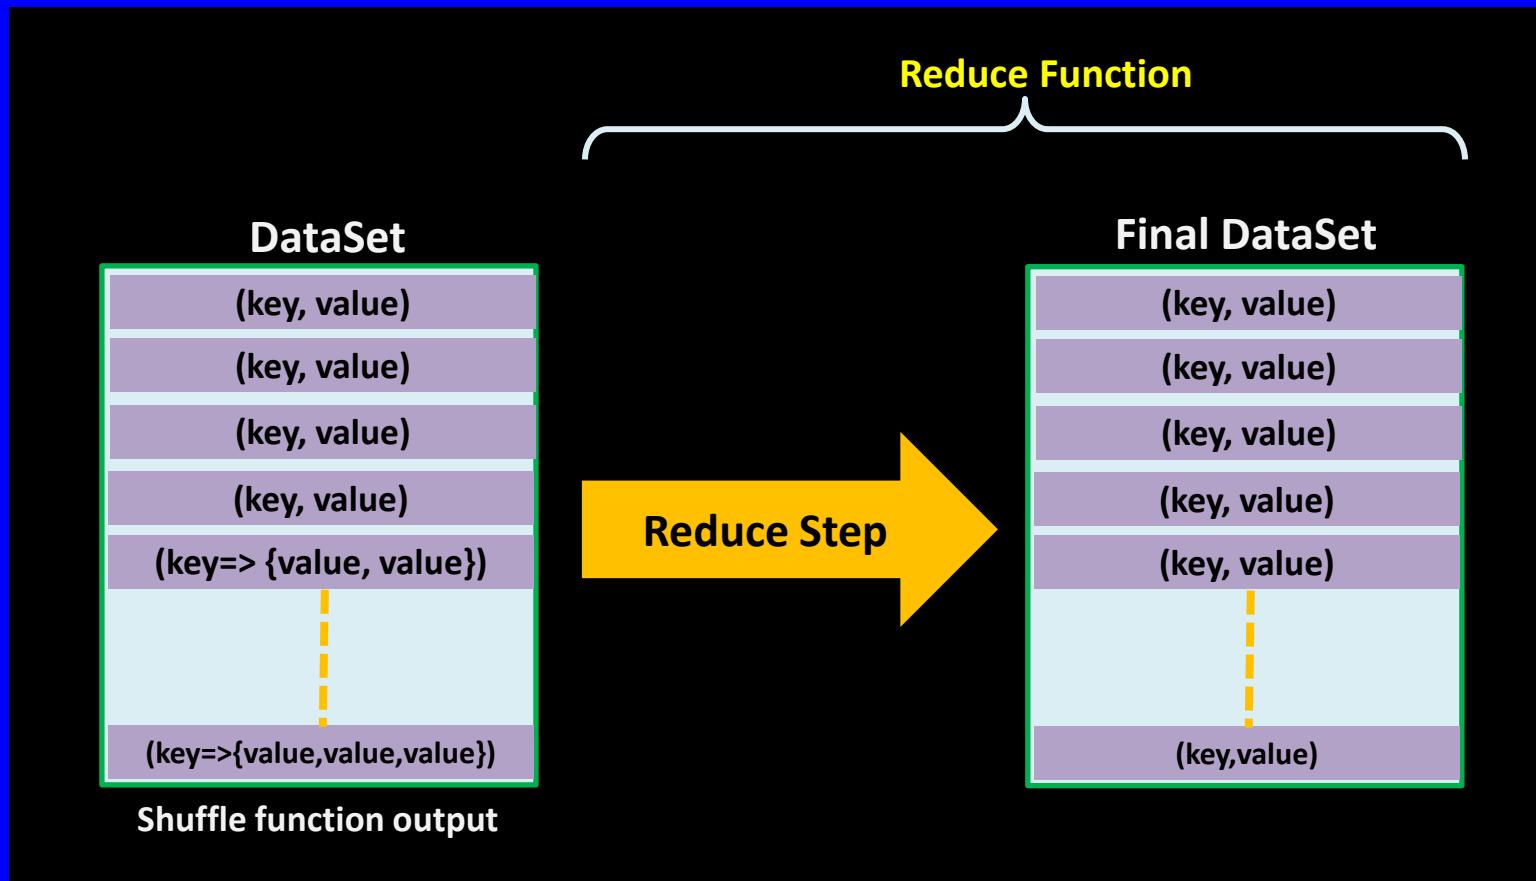
MapReduce Second Step Output: **shuffle function output = List of<key, List<value>> pairs**

# Reduce Function

## Reduce

- ❑ Final step in MapReduce Algorithm
- ❑ Performs Only one step: Reduce step

- Takes a list of <Key, List<Value>> sorted pairs from shuffle function
- Performs reduce operation



MapReduce final Step Output: **Reduce function output = List of<key, value> pairs**

# MapReduce with word count

# Problem statement

- Counting: the number of occurrences of each word available in a Data-Set.

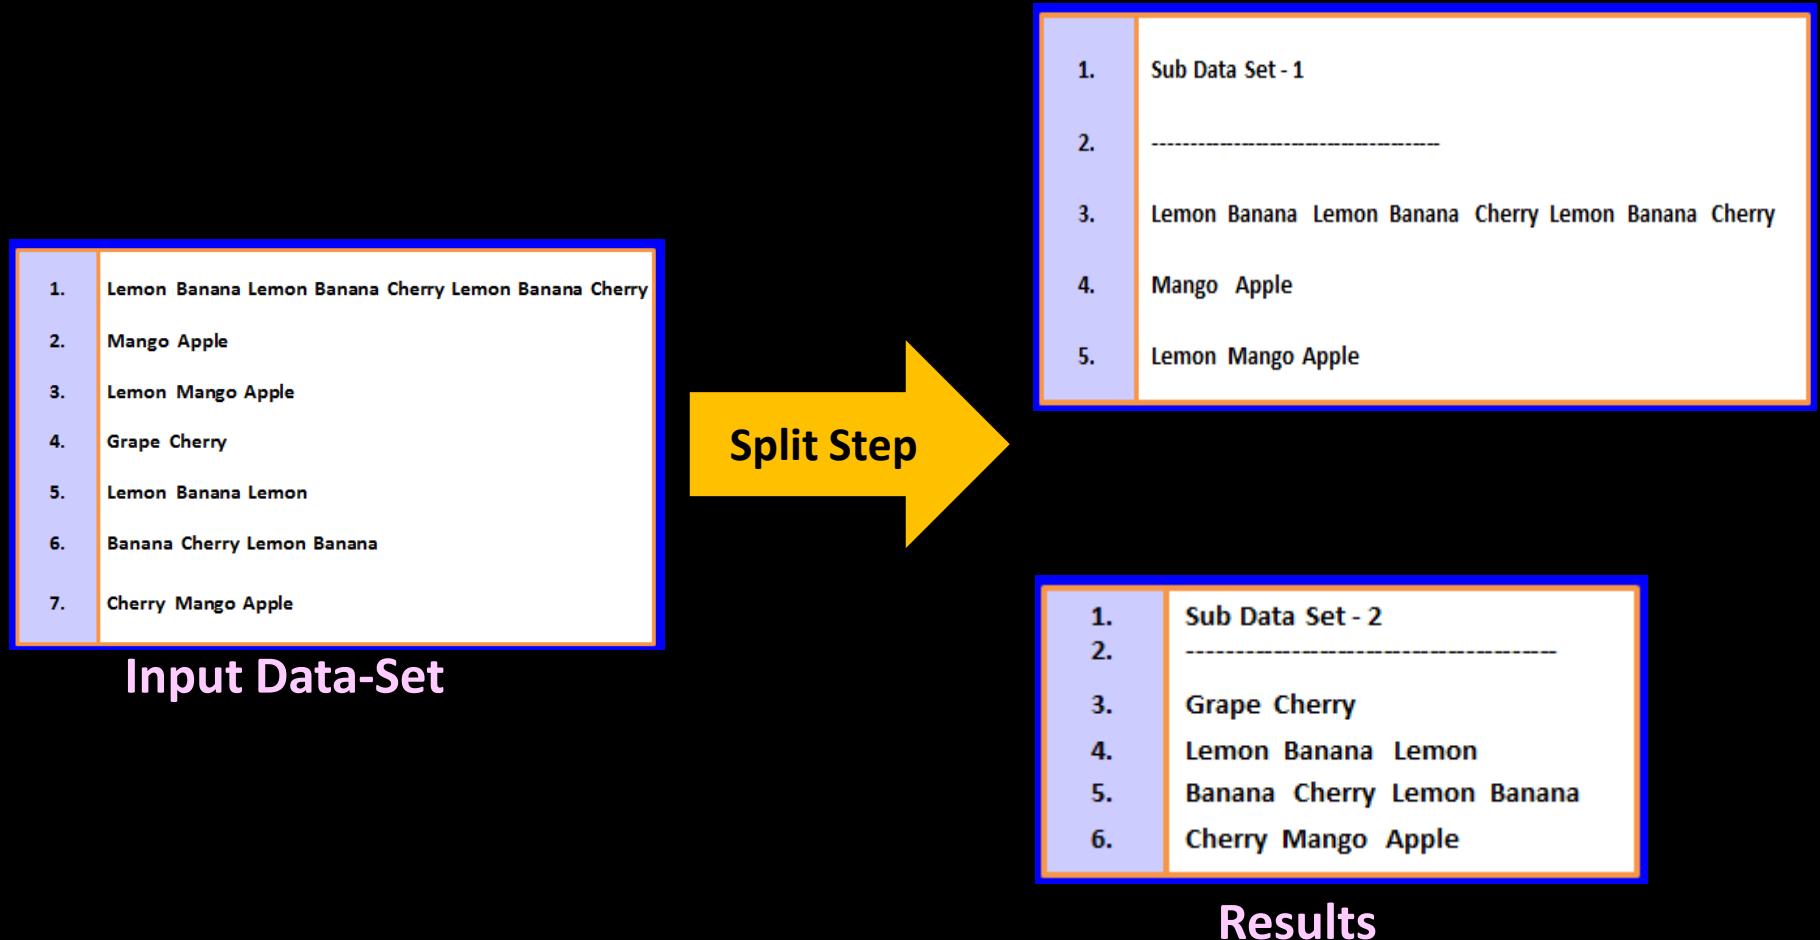
## **Input Data-Set**

1. Lemon Banana Lemon Banana Cherry Lemon Banana Cherry
2. Mango Apple
3. Lemon Mango Apple
4. Grape Cherry
5. Lemon Banana Lemon
6. Banana Cherry Lemon Banana
7. Cherry Mango Apple

## **Expected output**

1. Apple 3
2. Banana 6
3. Cherry 5
4. Grape 1
5. Lemon 7
6. Mango 3

- Map Function (Split Step)



- **Map Function (Mapping Step)**

1.	Sub Data Set - 1
2.	-----
3.	Lemon Banana Lemon Banana Cherry Lemon Banana Cherry
4.	Mango Apple
5.	Lemon Mango Apple

Mapping step

1.	SubDataSet-1
2.	.....
3.	Lemon 1
4.	Banana 1
5.	Lemon 1
6.	Banana 1
7.	Cherry 1
8.	Lemon 1
9.	Banana 1
10.	Cherry 1
11.	Mango 1
12.	Black 1
13.	Lemon 1
14.	Mango 1
15.	Apple 1

1.	Sub Data Set - 2
2.	-----
3.	Grape Cherry
4.	Lemon Banana Lemon
5.	Banana Cherry Lemon Banana
6.	Cherry Mango Apple

1.	SubDataSet-2
2.	.....
3.	Grape 1
4.	Cherry 1
5.	Lemon 1
6.	Banana 1
7.	Lemon 1
8.	Banana 1
9.	Cherry 1
10.	Lemon 1
11.	Banana 1
12.	Cherry 1
13.	Mango 1
14.	Apple 1

Output of split step

- **Shuffle Function (Merge Step-1)**

1.	SubDataSet-1
2.	.....
3.	Lemon 1
4.	Banana 1
5.	Lemon 1
6.	Banana 1
7.	Cherry 1
8.	Lemon 1
9.	Banana 1
10.	Cherry 1
11.	Mango 1
12.	Apple 1
13.	Lemon 1
14.	Mango 1
15.	Apple 1

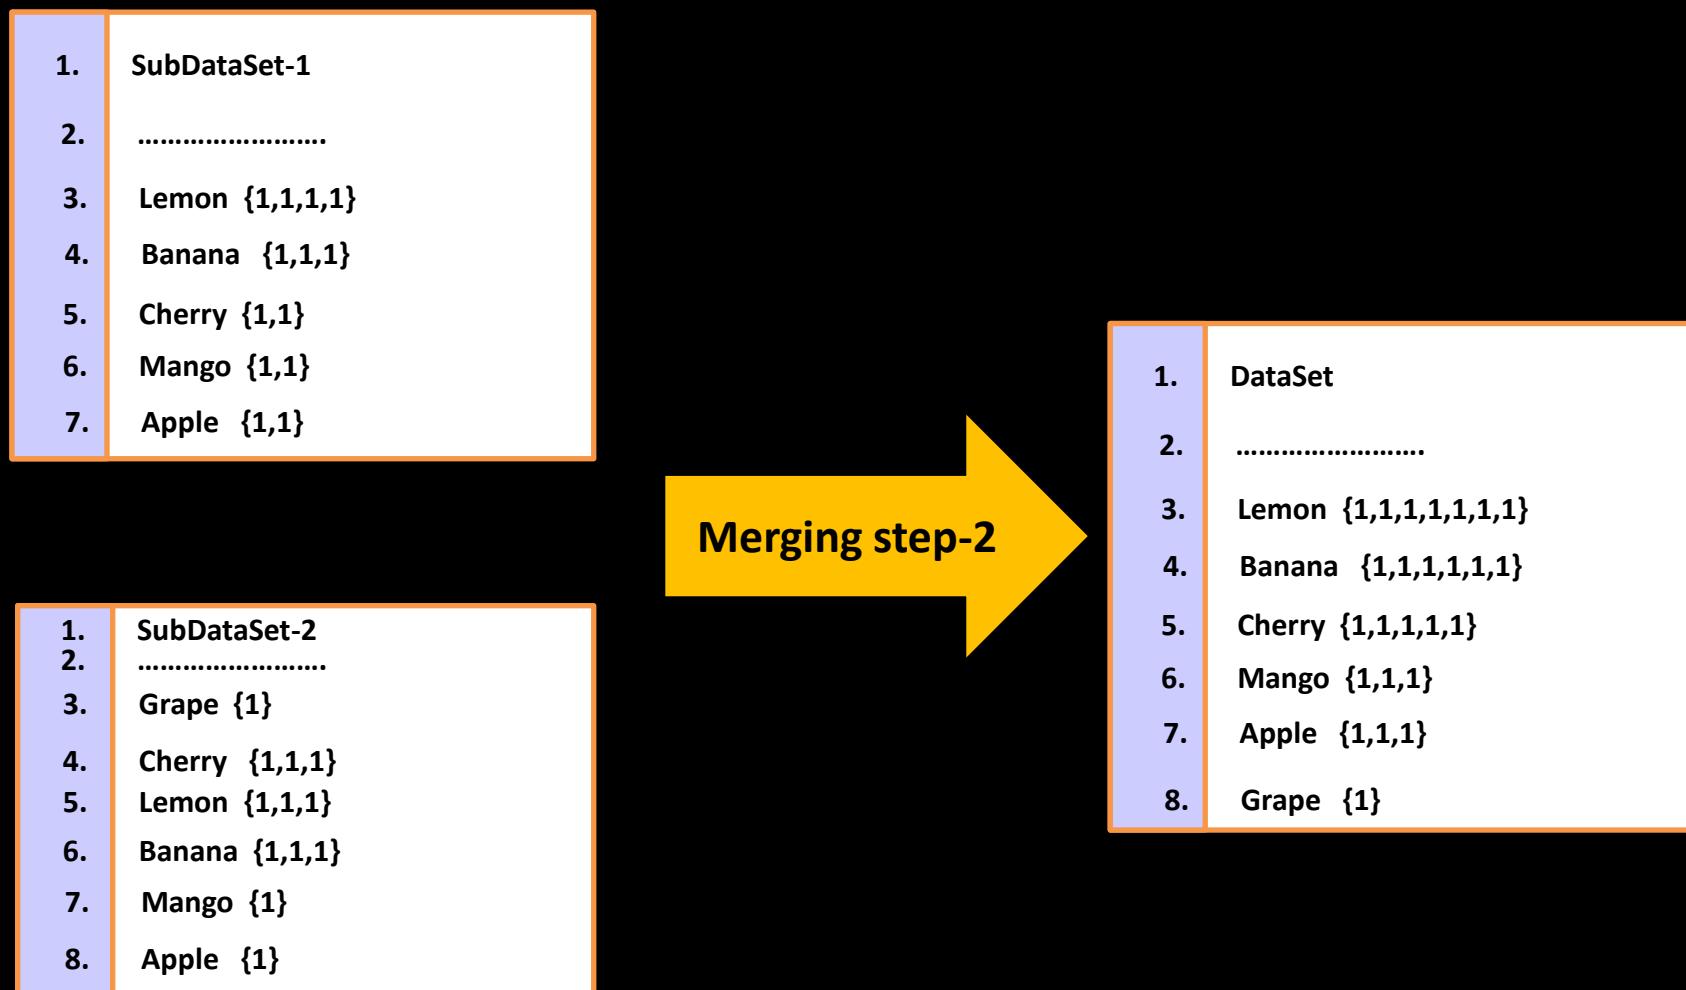
1.	SubDataSet-1
2.	.....
3.	Lemon {1,1,1,1}
4.	Banana {1,1,1}
5.	Cherry {1,1}
6.	Mango {1,1}
7.	Apple {1,1}

Merging step-1

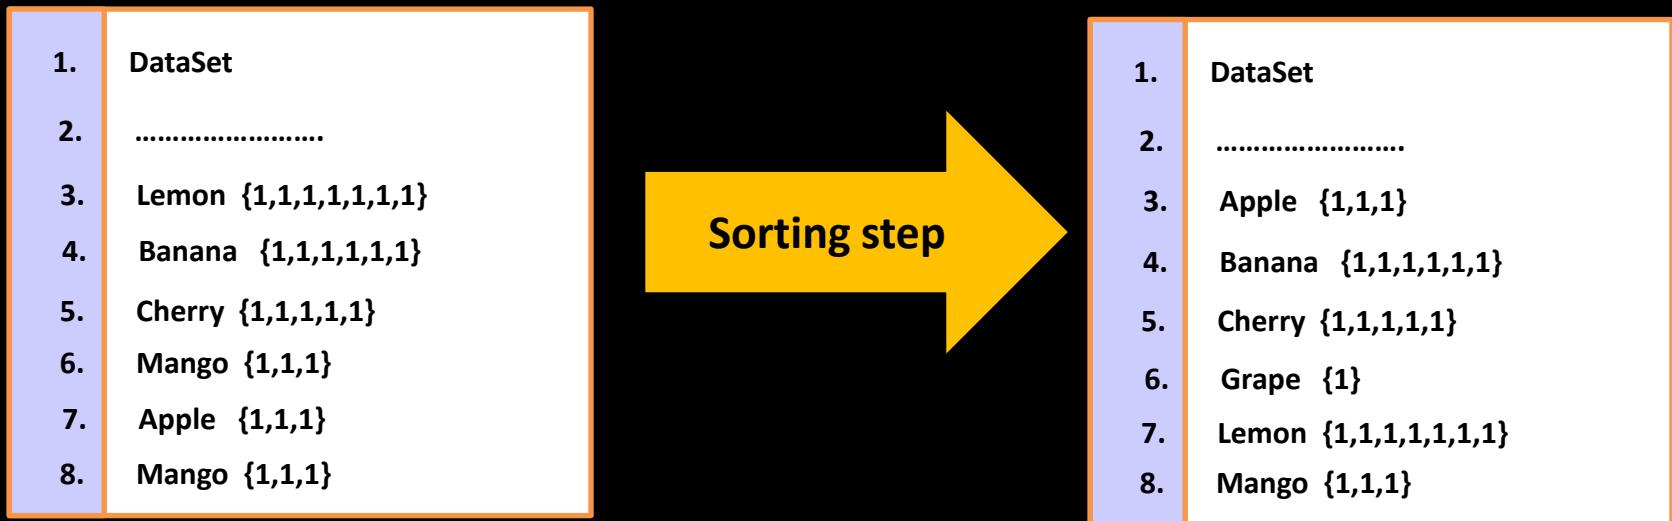
1.	SubDataSet-2
2.	.....
3.	Grape 1
4.	Cherry 1
5.	Lemon 1
6.	Banana 1
7.	Lemon 1
8.	Banana 1
9.	Cherry 1
10.	Lemon 1
11.	Banana 1
12.	Cherry 1
13.	Mango 1
14.	Apple 1

1.	SubDataSet-2
2.	.....
3.	Grape {1}
4.	Cherry {1,1,1}
5.	Lemon {1,1,1}
6.	Banana {1,1,1}
7.	Mango {1}
8.	Apple {1}

- **Shuffle Function (Merge Step-2)**



- **Shuffle Function (Sorting Step)**



- Reduce Function (Reduce Step)

1. DataSet
2. .....
3. Apple {1,1,1}
4. Banana {1,1,1,1,1,1}
5. Cherry {1,1,1,1,1}
6. Grape {1}
7. Lemon {1,1,1,1,1,1,1}
8. Mango {1,1,1}



Reduce step

1. DataSet
2. .....
3. Apple 3
4. Banana 6
5. Cherry 5
6. Grape 1
7. Lemon 7
8. Mango 3

Final Output

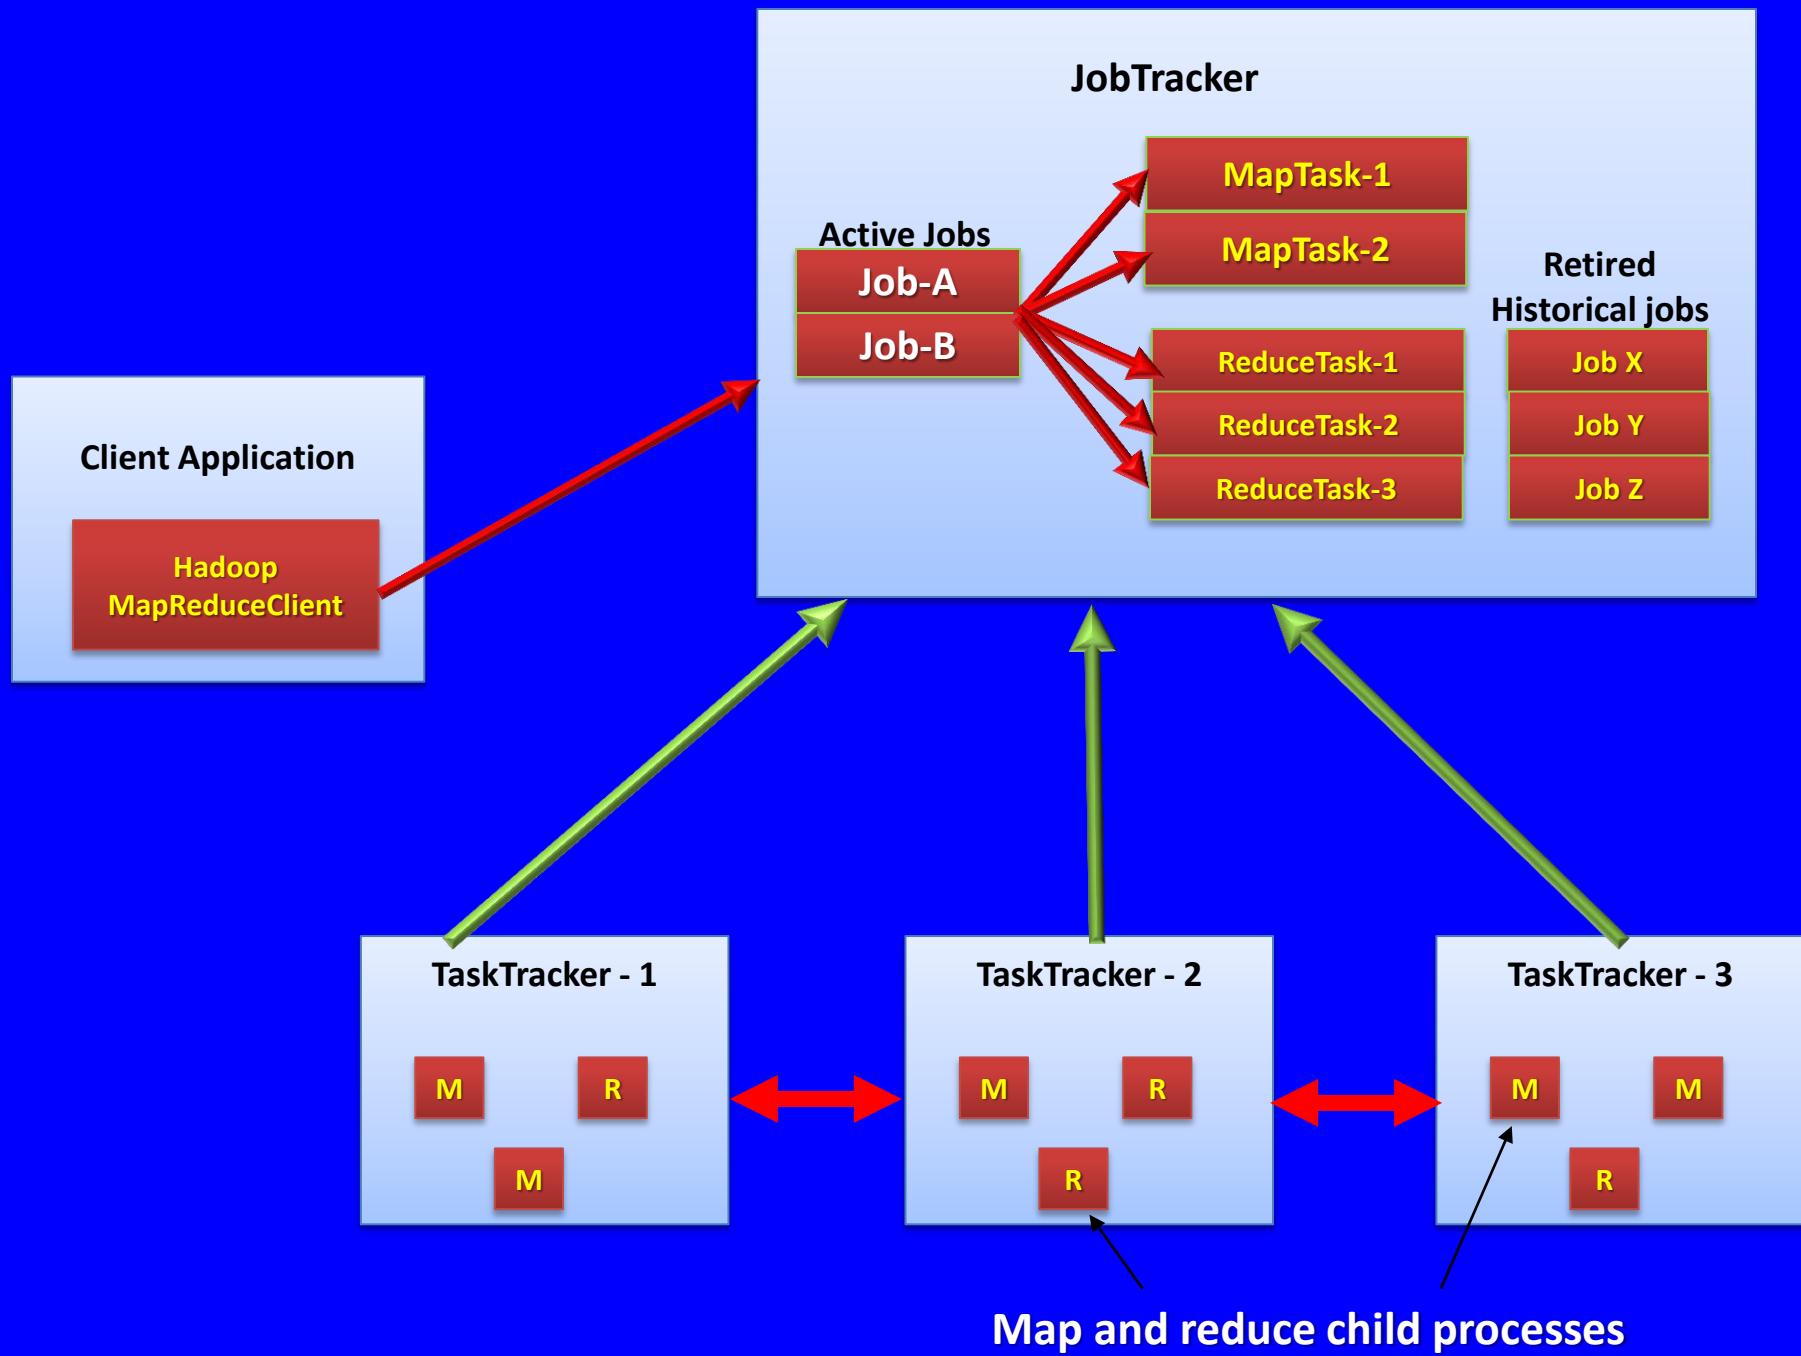
# What is Map-Reduce

Map  
Reduce

- ❖ **Distributed computing framework:** batch based.
- ❖ **Parallel works:** over a large amount of raw data.

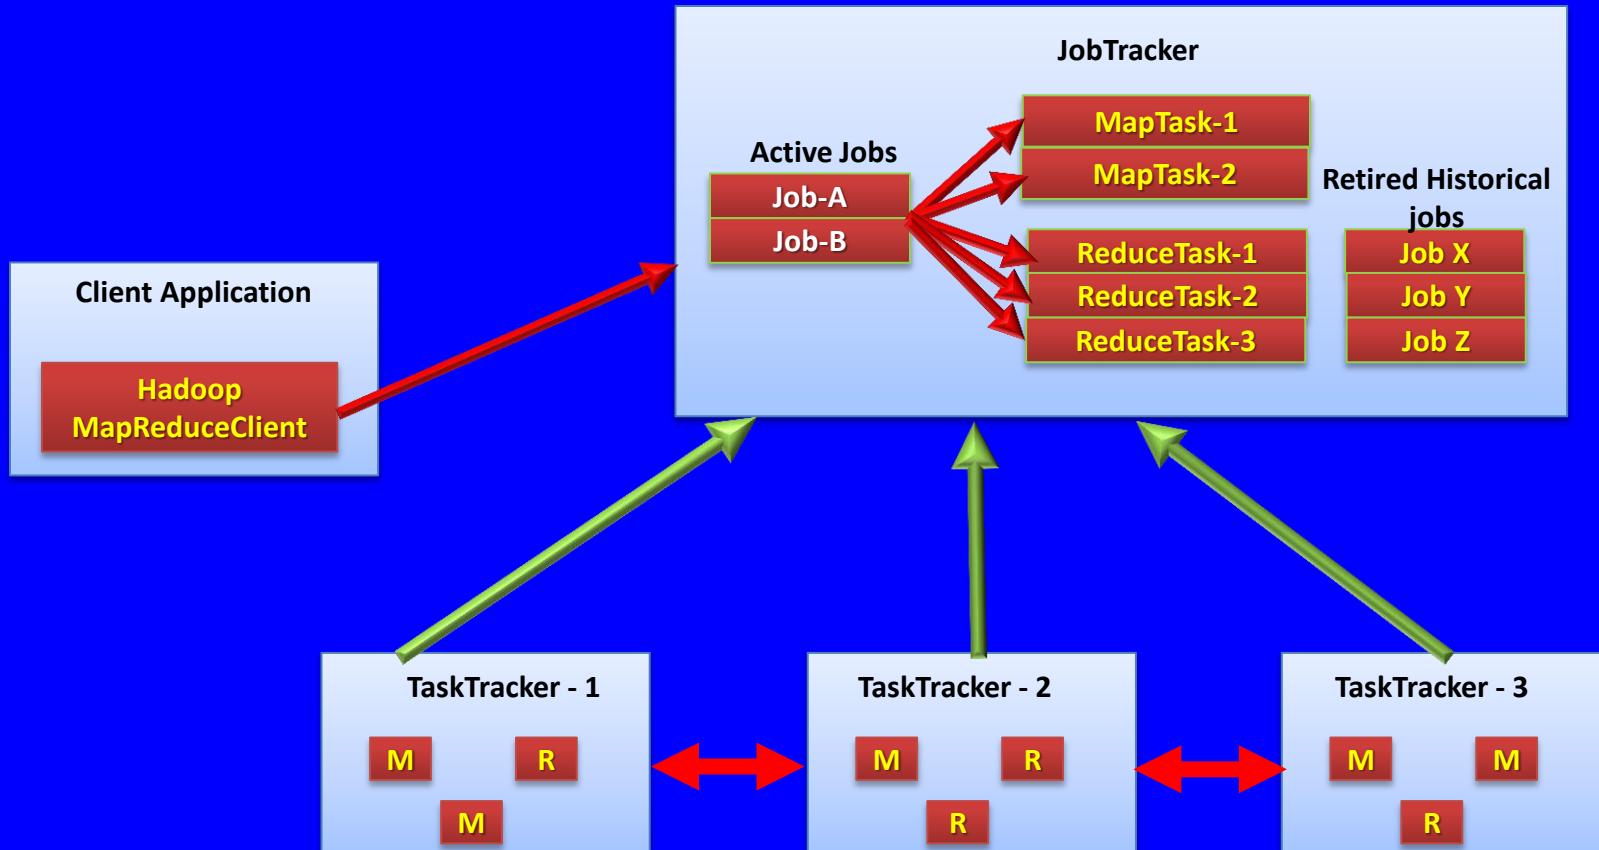
# MapReduce logical architecture

- ❖ Hadoop's MapReduce architecture is similar to the master-slave model in HDFS.
- The main components of MapReduce:
  1. Hadoop Map Reduce client
  2. Task Tracker
  3. Job Tracker



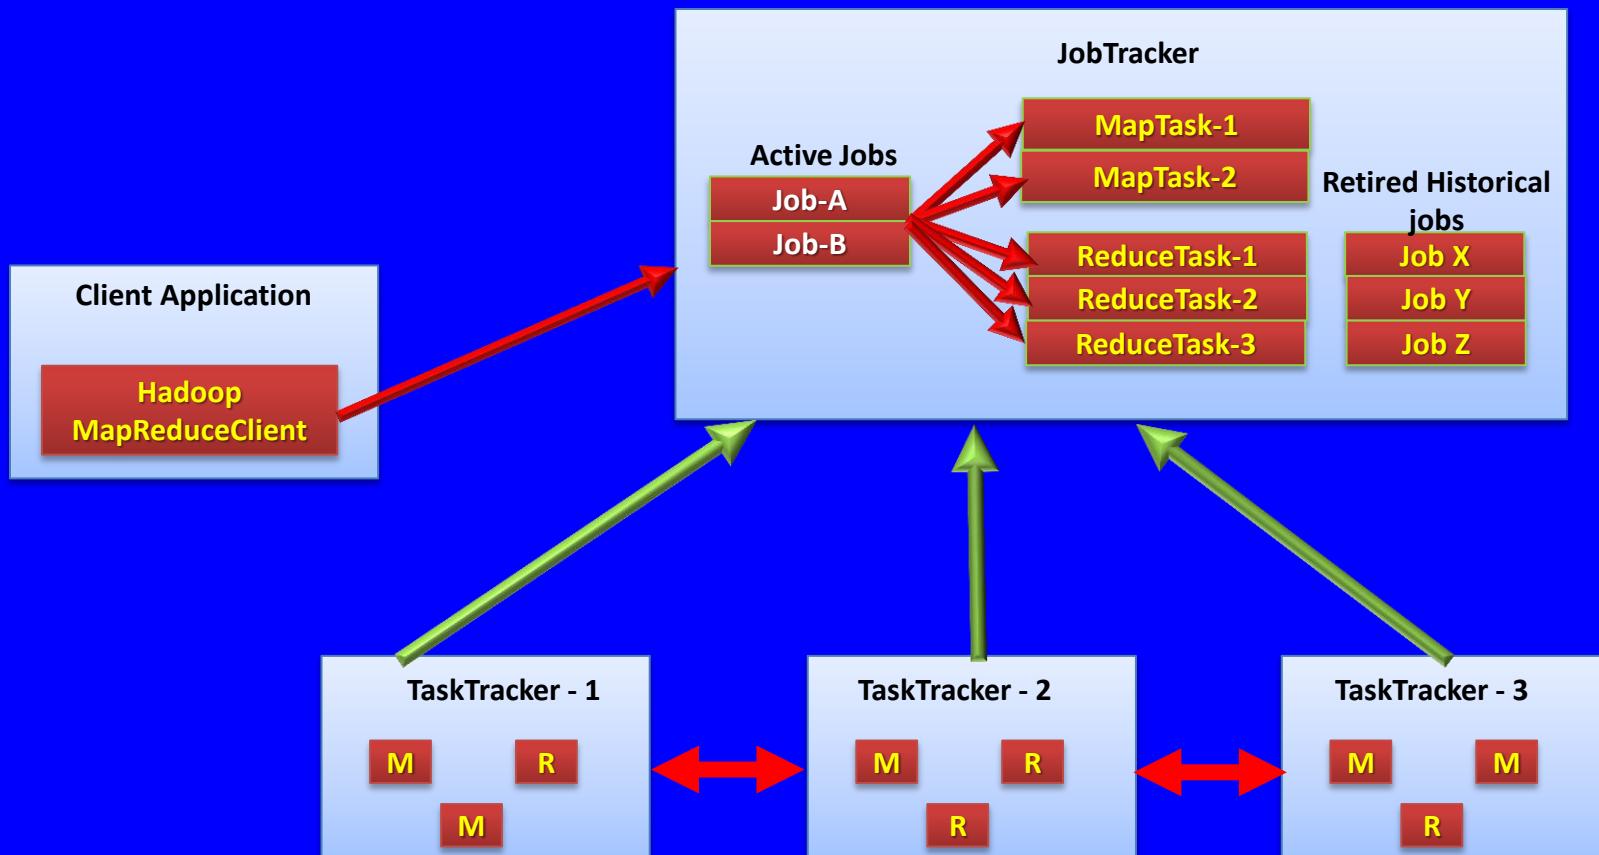
# JobTracker

- ❖ Number of JobTrackers: There is one JobTracker running on a master node.
- ❖ Single point of failure: It is also single point of failure.



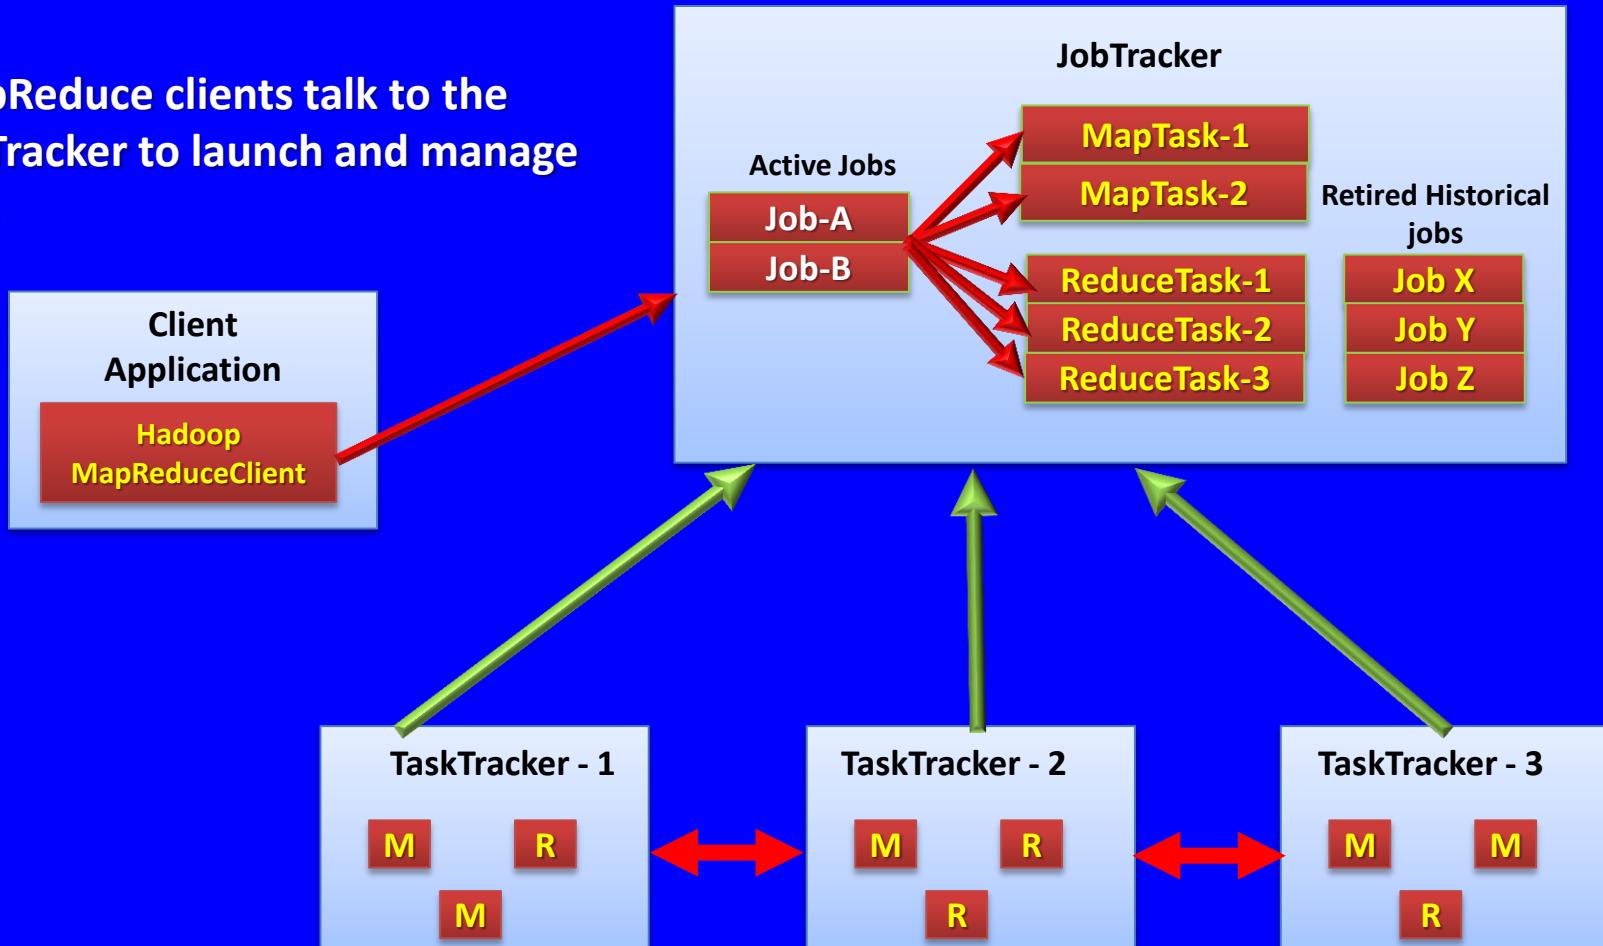
# TaskTrackers

- ❖ Number of TaskTrackers: several TaskTracker running on slave nodes.

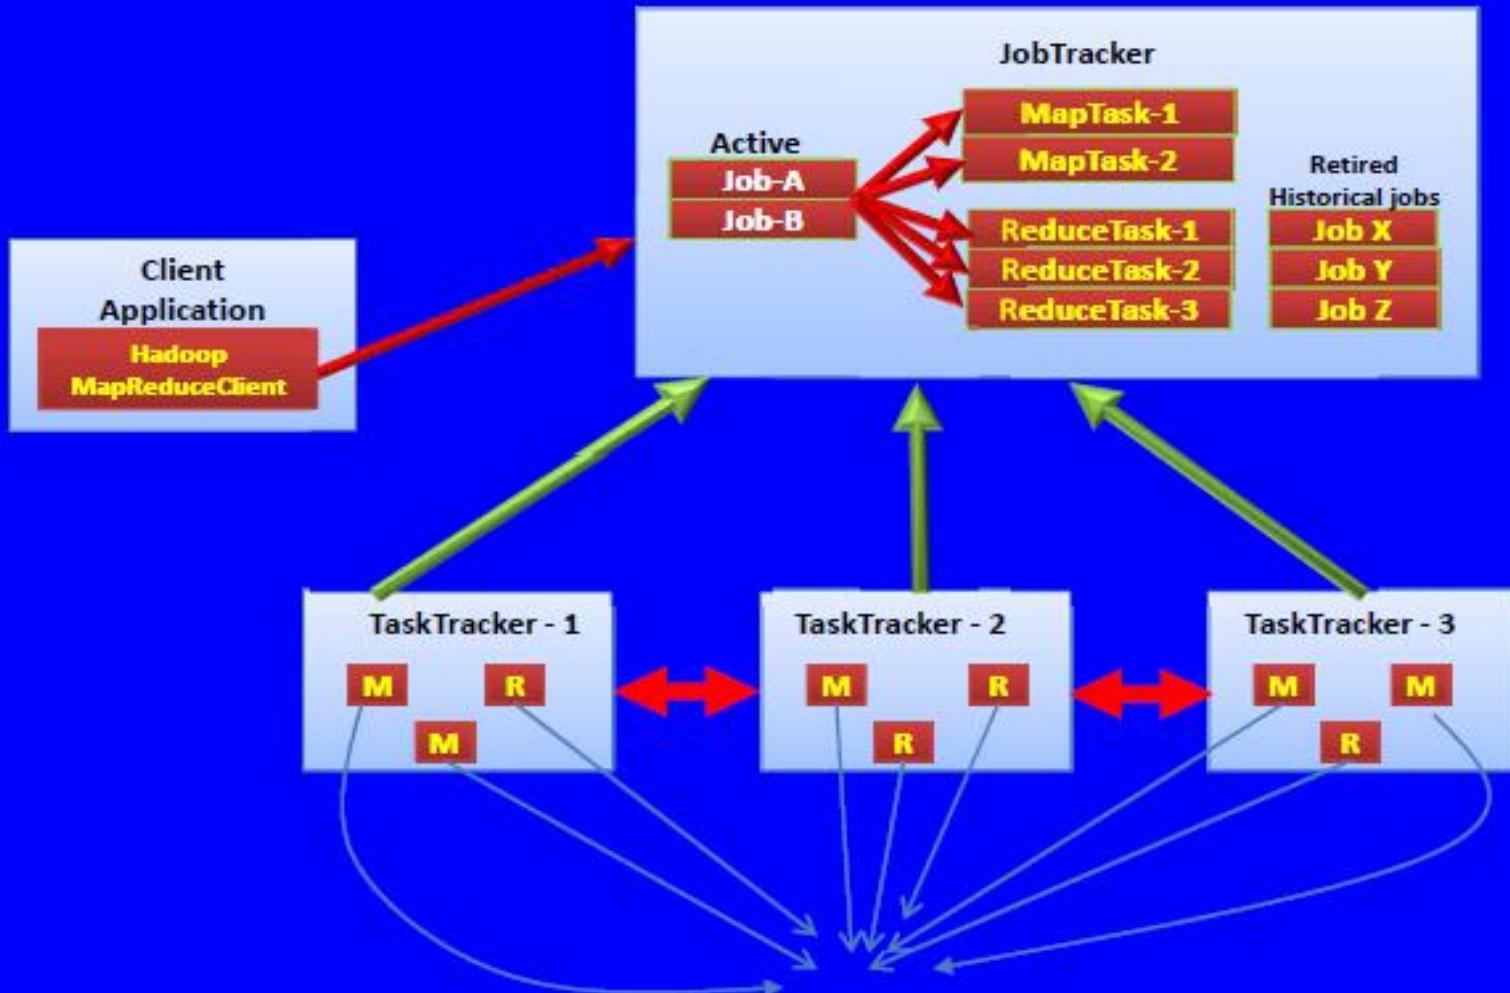


# MapReduce Clients

MapReduce clients talk to the JobTracker to launch and manage jobs

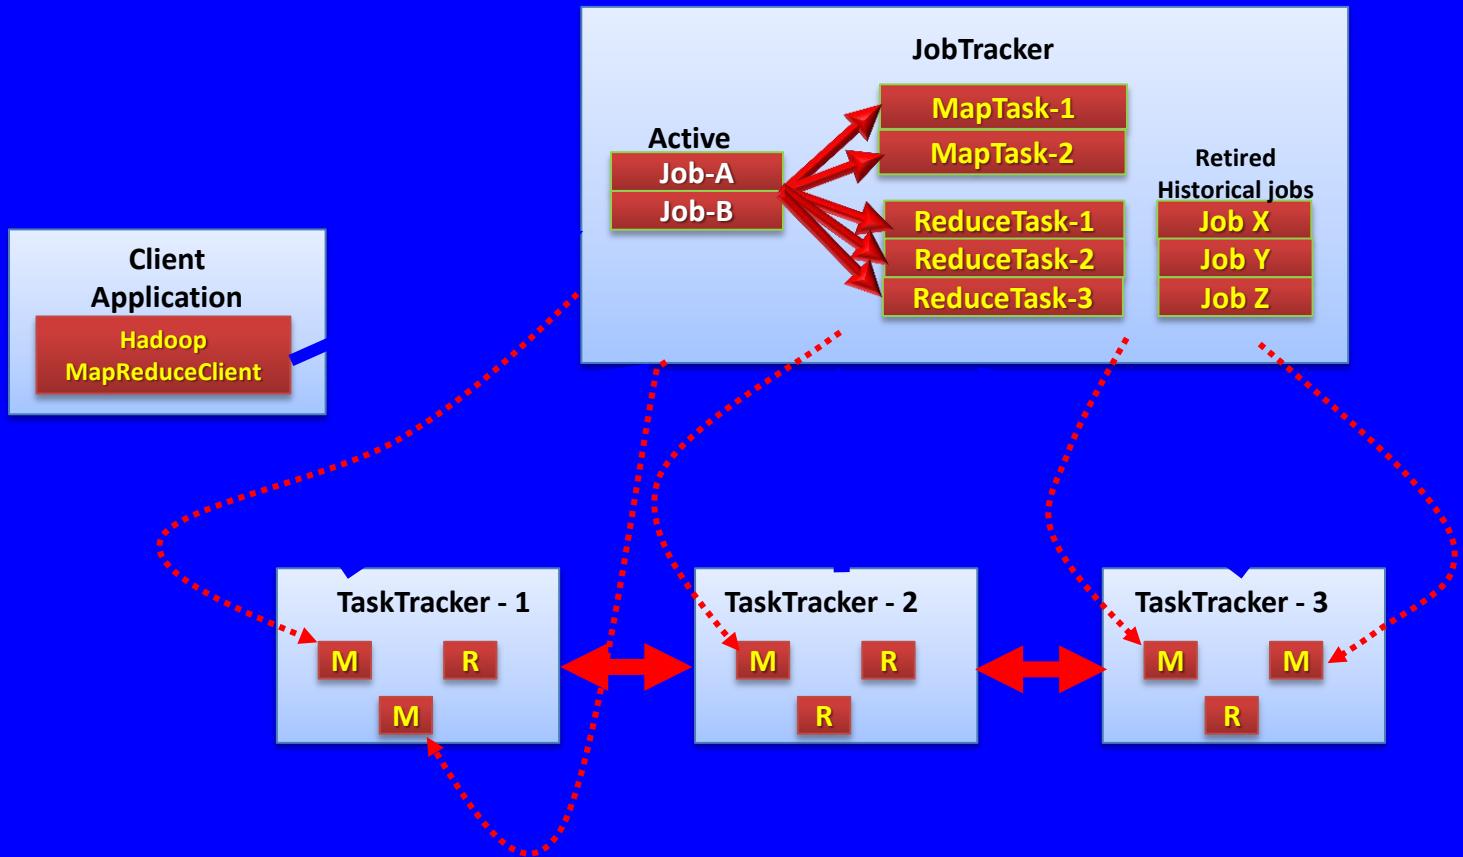


# TaskTracker: performs

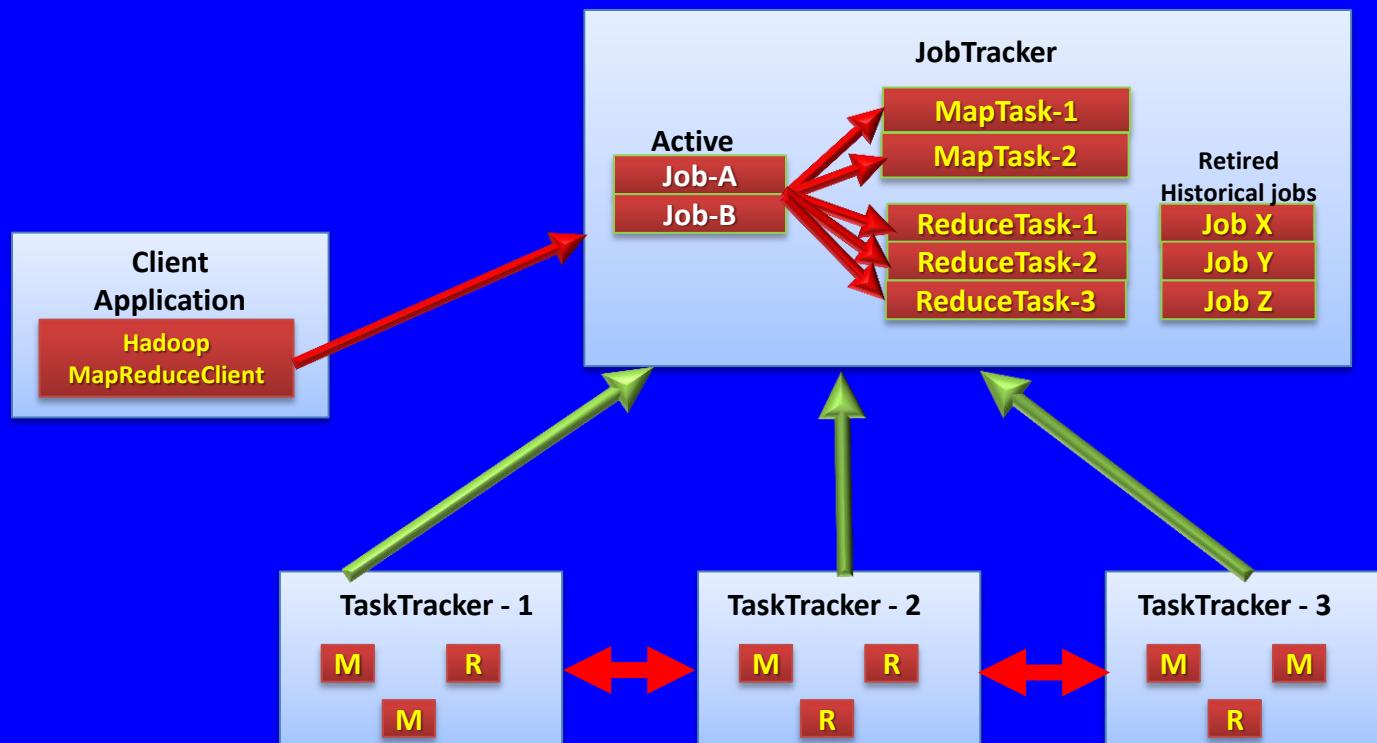


- Child processes spawned by TaskTracker:  
to perform the actual map or reduce work

- Map tasks input: read their input from HDFS
- Map task output: Write their output to the local disk.

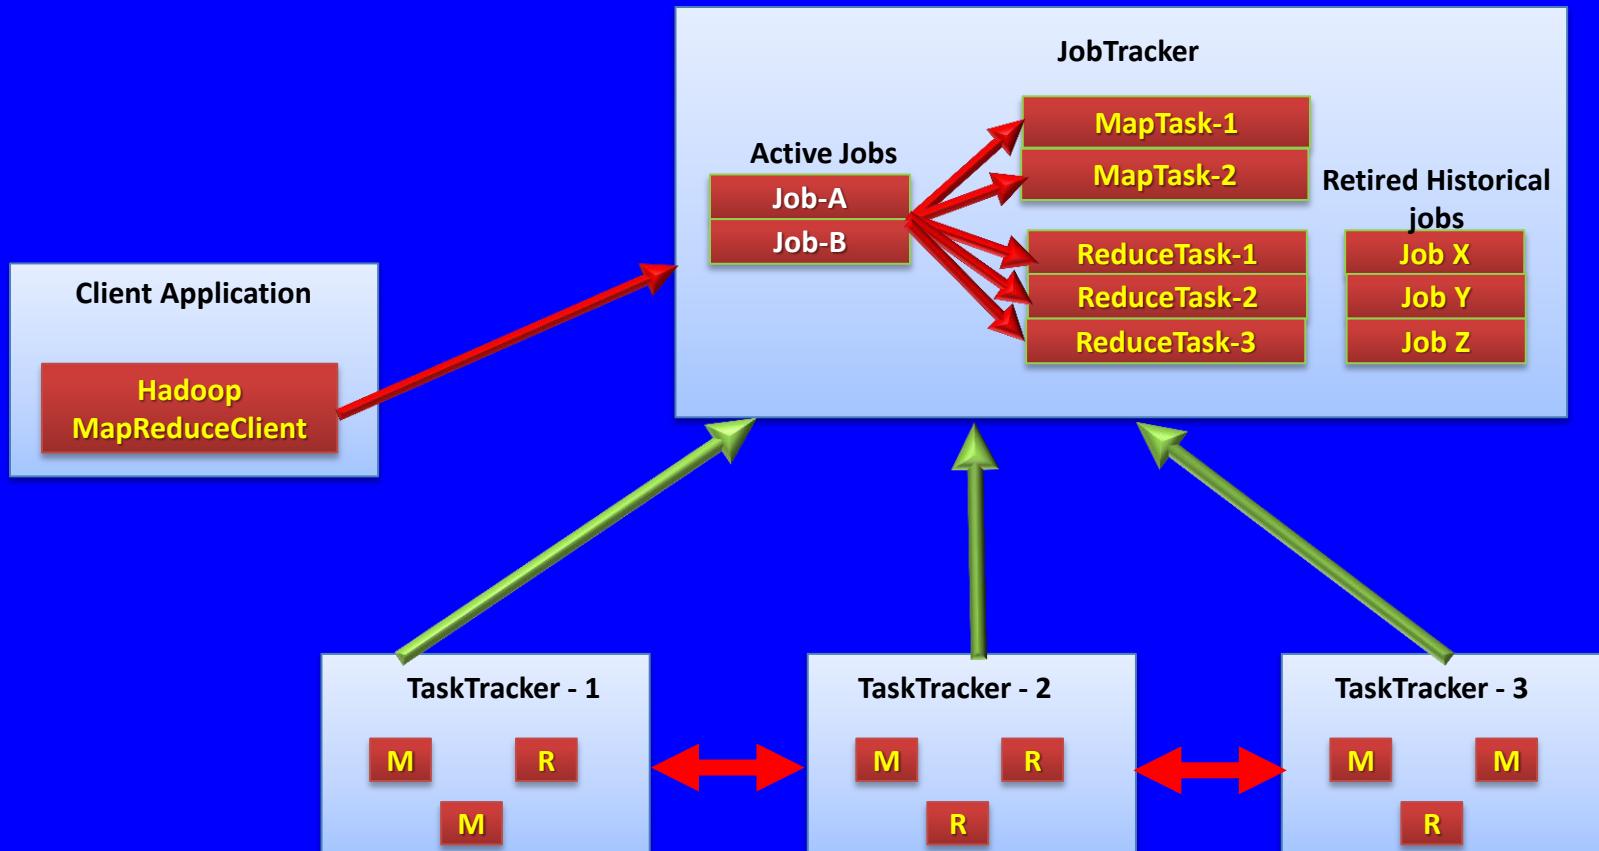


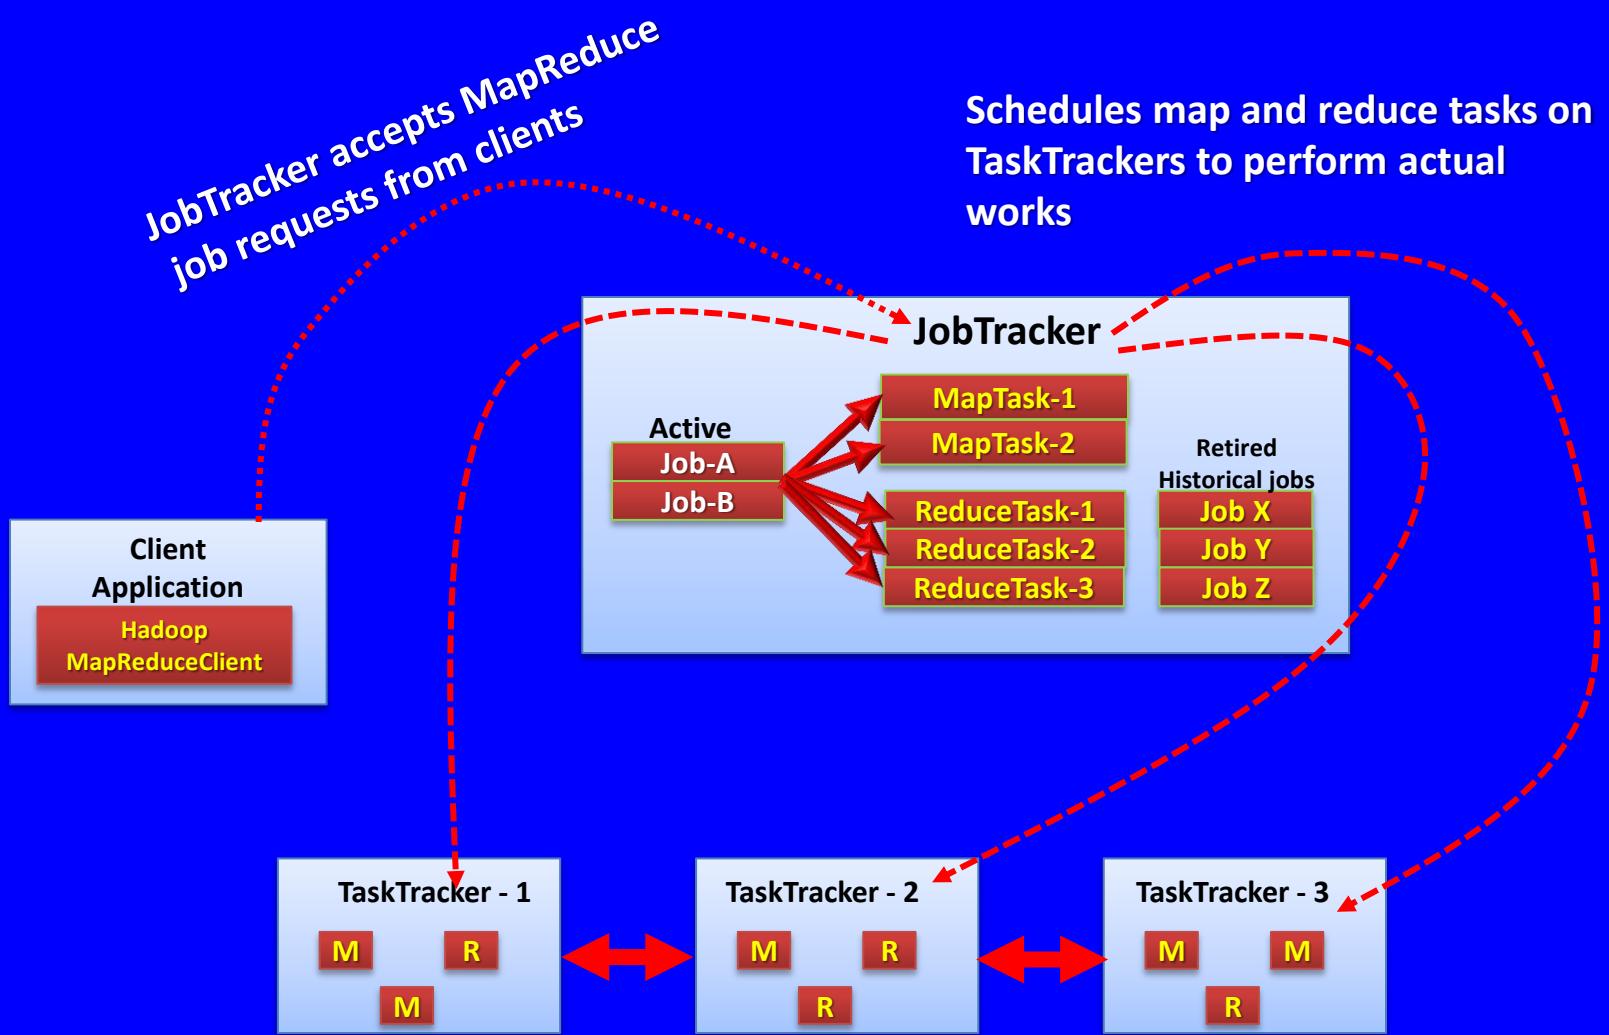
- Reduce tasks inputs: Reduce tasks read the map outputs over the network.
- Reduce tasks outputs: write their outputs back to HDFS



# Job Tracker: performs

- ❖ Coordination amongst TaskTrackers:
  - JobTracker coordinates activities across the active TaskTrackers.





## **Step - II**

# Hadoop: Third step

## Configuration and Setups

**Single node cluster**

# **Setting Hadoop cluster On Ubuntu Linux**

**Step – 1**

**Install Sun Java**

# Installing the JDK manually

- Download the adequate JDK from the Oracle Web site.



- ❖ Create temporary directory to keep the installer:

```
root@chitrakant-System-Product-Name: ~
root@chitrakant-System-Product-Name:~# mkdir -p ~/tmp/jdk-6u31
root@chitrakant-System-Product-Name:~#
```

## ❖ Copy the file to the tmp location:

```
root@chitrakant-System-Product-Name: ~/tmp/jdk-6u31
root@chitrakant-System-Product-Name:~/tmp/jdk-6u31# ls
jdk-6u31-linux-i586.bin
root@chitrakant-System-Product-Name:~/tmp/jdk-6u31# cp /root/Desktop/Required_So
ftwares/java/jdk-6u31-linux-i586.bin .
```

## ❖ Provide various permissions ( make it executable):

```
root@chitrakant-System-Product-Name: ~/tmp/jdk-6u31
root@chitrakant-System-Product-Name:~/tmp/jdk-6u31# ls
jdk-6u31-linux-i586.bin
root@chitrakant-System-Product-Name:~/tmp/jdk-6u31# chmod 777 jdk-6u31-linux-i58
6.bin
root@chitrakant-System-Product-Name:~/tmp/jdk-6u31# ls
jdk-6u31-linux-i586.bin
root@chitrakant-System-Product-Name:~/tmp/jdk-6u31#
```

## ❖ Create target directory for java:

```
root@chitrakant-System-Product-Name: /  
root@chitrakant-System-Product-Name:/# ls  
bin dev initrd.img media proc sbin sys var  
boot etc lib mnt root selinux tmp vmlinuz  
cdrom home lost+found opt run srv usr  
root@chitrakant-System-Product-Name:/# mkdir java 1  
root@chitrakant-System-Product-Name:/#
```

## ❖ Copy executable file into the created java directory:

```
root@chitrakant-System-Product-Name:~/tmp/jdk-6u31# cp jdk-6u31-linux-i586.bin /  
java/  
root@chitrakant-System-Product-Name:~/tmp/jdk-6u31# 2
```

## ❖ Start java installation :

\$ ./jdk-6u31-linux-i586.bin

3

root@chitrakant-System-Product-Name: /java

```
Creating jdk1.6.0_31/jre/lib/plugin.jar
Creating jdk1.6.0_31/jre/lib/javaws.jar
Creating jdk1.6.0_31/jre/lib/deploy.jar
```

Java(TM) SE Development Kit 6 successfully installed.

Product Registration is FREE and includes many benefits:

- \* Notification of new versions, patches, and updates
- \* Special offers on Oracle products, services and training
- \* Access to early releases and documentation

Product and system data will be collected. If your configuration supports a browser, the JDK Product Registration form will be presented. If you do not register, none of this information will be saved. You may also register your JDK later by opening the register.html file (located in the JDK installation directory) in a browser.

For more information on what data Registration collects and how it is managed and used, see:

<http://java.sun.com/javase/registration/JDKRegistrationPrivacy.html>

Press Enter to continue.....

## ❖ create a symbolic link:

```
root@chitrakant-System-Product-Name: /java
root@chitrakant-System-Product-Name:/java# ls
jdk1.6.0_31  jdk-6u31-linux-i586.bin
root@chitrakant-System-Product-Name:/java# ln -s jdk1.6.0_31/ jdk_6
root@chitrakant-System-Product-Name:/java# ls
jdk1.6.0_31  jdk_6  jdk-6u31-linux-i586.bin
root@chitrakant-System-Product-Name:/java#
```

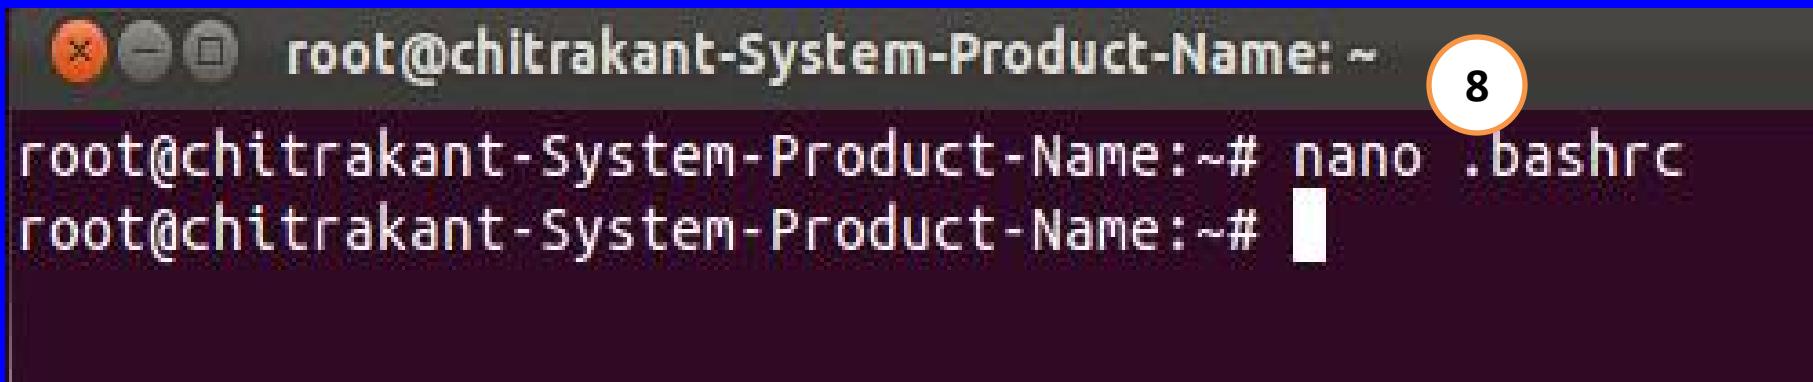
5                           6

- Add to the `~/.bashrc` the path to our JDK binary files

```
root@chitrakant-System-Product-Name: ~ 7
root@chitrakant-System-Product-Name:~# ls -a
.
..
.bash_history      Documents          .ICEauthority
.bashrc             Downloads          .libreoffice
.cache              .gconf             .local
.config             .gnome2            .mission-control
.dbus               .gstreamer-0.10   Music
Desktop            .gtk-bookmarks    Pictures
                      .gvfs             .profile
                                         Public
```

.pulse  
.pulse-cookie  
Templates  
.thumbnails  
tmp  
Videos  
.Xauthority  
.xsession-errors

## ❖ Open .bashrc file to add java installation

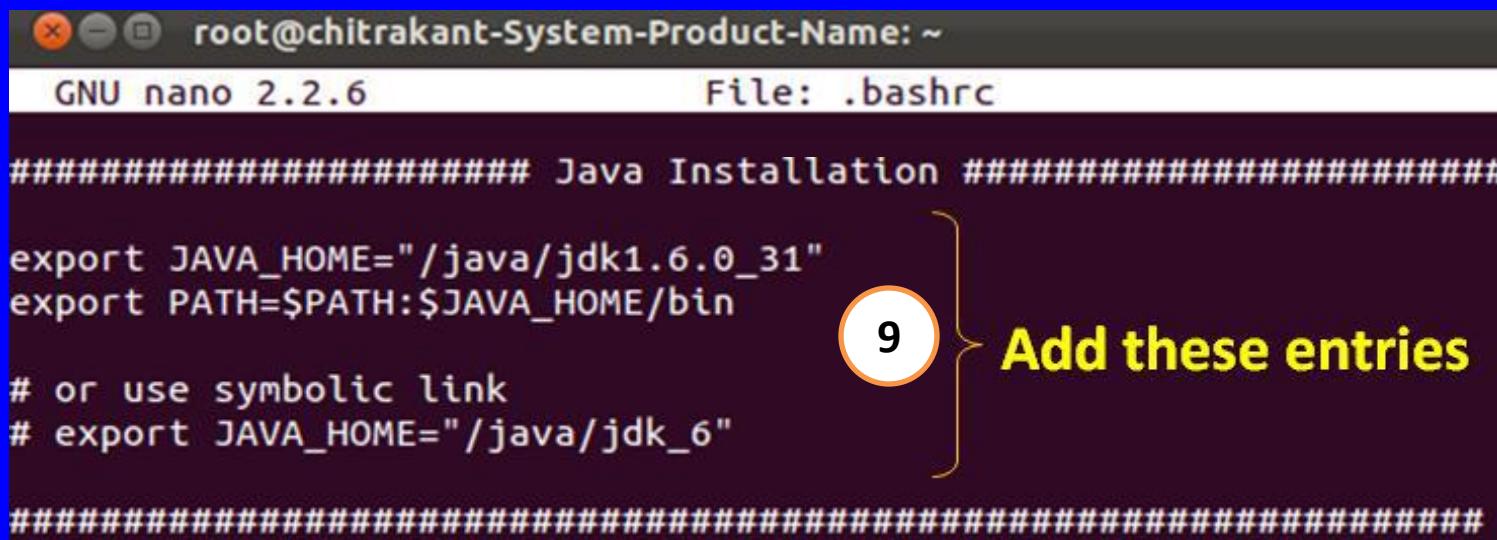


The screenshot shows a terminal window with a dark background and light-colored text. At the top, there are three icons: a red circle with an orange 'X', a grey circle with a white 'E', and a grey circle with a white 'D'. To the right of these icons, the text 'root@chitrakant-System-Product-Name: ~' is displayed. In the top right corner of the terminal window, there is a small circular badge with the number '8' inside it. Below the header, the terminal prompt shows the command 'root@chitrakant-System-Product-Name:~# nano .bashrc'. The cursor is visible at the end of the command line.

8

## Set JAVA\_HOME

- Now set the JAVA\_HOME environment variable to allow applications to find where my JDK is installed.



root@chitrakant-System-Product-Name: ~

GNU nano 2.2.6 File: .bashrc

```
#####
# Java Installation #####
export JAVA_HOME="/java/jdk1.6.0_31"
export PATH=$PATH:$JAVA_HOME/bin
# or use symbolic link
# export JAVA_HOME="/java/jdk_6"
#####
```

9 } Add these entries

- ❖ Close the current terminal and start new terminal.
- ❖ Apply the command : **java -version**

```
x root@chitrakant-System-Product-Name:~  
root@chitrakant-System-Product-Name:~# echo $JAVA_HOME 10  
/java/jdk1.6.0_31  
root@chitrakant-System-Product-Name:~# java -version 12  
java version "1.6.0_31"  
Java(TM) SE Runtime Environment (build 1.6.0_31-b04)  
Java HotSpot(TM) Client VM (build 20.6-b01, mixed mode, sharing)  
root@chitrakant-System-Product-Name:~#
```

**Step - 1**

## Step – 2

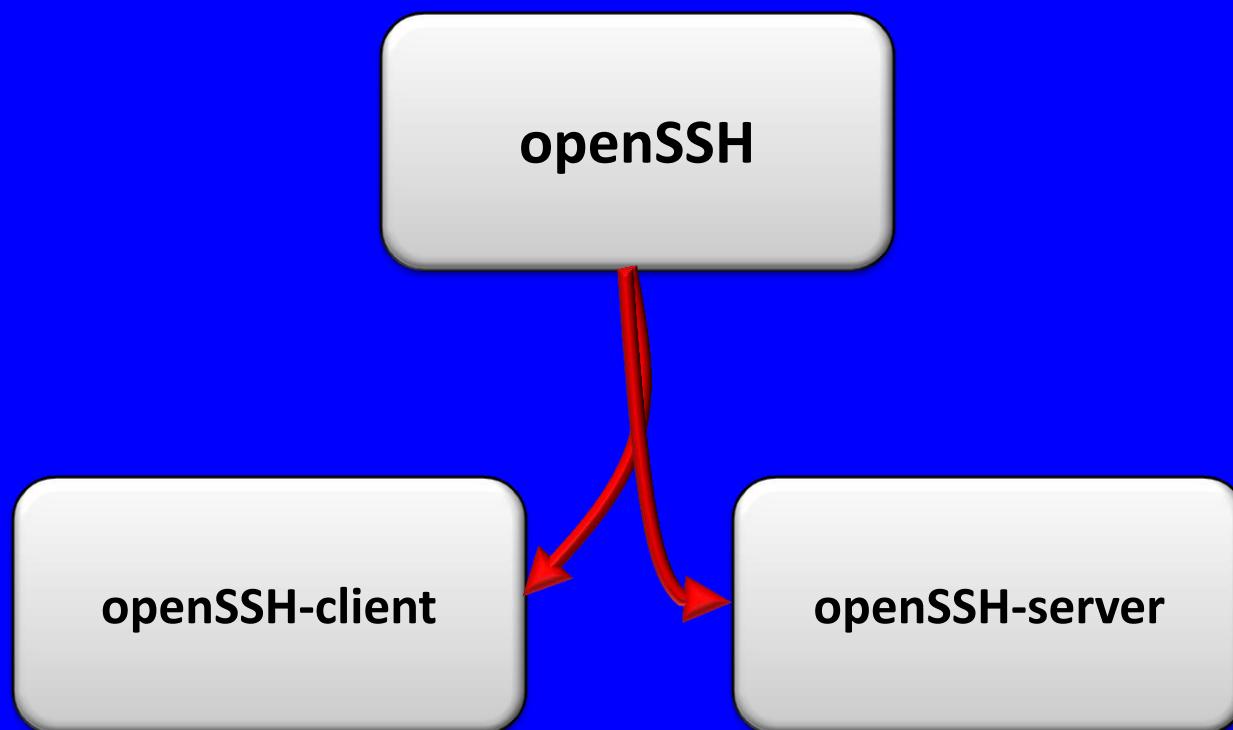
# Install SSH

# SSH

- Hadoop requires SSH access to manage its nodes:
  - Remote machines
  - Or local machine if used Hadoop on it.
- For single-node Hadoop setup: configure SSH access to localhost .

# OpenSSH

- The Ubuntu flavor of SSH is called OpenSSH.



# ssh-client already installed

```
x - root@chitrakant-System-Product-Name: ~
root@chitrakant-System-Product-Name:~# man ssh
root@chitrakant-System-Product-Name:~# ssh localhost
The authenticity of host 'localhost (127.0.0.1)' can't be established.
ECDSA key fingerprint is 6a:25:a0:8a:5d:7f:55:3c:97:b8:d5:7f:9e:a7:a3:80.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'localhost' (ECDSA) to the list of known hosts.
root@localhost's password:
Welcome to Ubuntu 11.10 (GNU/Linux 3.0.0-12-generic i686)

 * Documentation: https://help.ubuntu.com/

0 packages can be updated.
0 updates are security updates.

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

root@chitrakant-System-Product-Name:~# 1
```

## ❖ Install openssh online:

```
$ apt-get install openssh-client openssh-server
```

## ❖ Manually Install openSSH:

### ❖ Download openssh installer from their web site:

 **openssh-server\_5.8p1-7ubuntu1\_i386.deb**

### ❖ Create a directory to store openssh installer:

```
root@chitrakant-System-Product-Name: /  
root@chitrakant-System-Product-Name:~# cd / 2  
root@chitrakant-System-Product-Name:/# ls  
bin dev initrd.img lost+found opt run srv usr  
boot etc java media proc sbin sys var  
cdrom home lib mnt root selinux tmp vmlinuz  
root@chitrakant-System-Product-Name:/# mkdir openssh 3
```

## ❖ Copy openssh installer into openssh directory:

4

```
root@chitrakant-System-Product-Name:/# cd openssh/  
root@chitrakant-System-Product-Name:/openssh# ls  
root@chitrakant-System-Product-Name:/openssh# cp /root/Desktop/Required_Software  
s/other_tools/openssh/openssh-server_5.8p1-7ubuntu1_i386.deb .  
root@chitrakant-System-Product-Name:/openssh# ls  
openssh-server_5.8p1-7ubuntu1_i386.deb  
root@chitrakant-System-Product-Name:/openssh#
```

5

5.1

## ❖ Make openssh installer as executable:

```
root@chitrakant-System-Product-Name:/openssl  
root@chitrakant-System-Product-Name:/openssl# ls  
openssh-server_5.8p1-7ubuntu1_i386.deb 6  
root@chitrakant-System-Product-Name:/openssl# chmod 777 openssh-server_5.8p1-7ub  
untu1_i386.deb  
root@chitrakant-System-Product-Name:/openssl# ls  
openssh-server_5.8p1-7ubuntu1_i386.deb 7  
root@chitrakant-System-Product-Name:/openssl# 8
```

## ❖ Install openssh using the command: **dpkg -i**

```
root@chitrakant-System-Product-Name:/openssl  
root@chitrakant-System-Product-Name:/openssl# dpkg -i openssh-server_5.8p1-7ubun  
tu1_i386.deb 9
```

## ❖ Apply ssh command to connect with the localhost:

```
x - root@chitrakant-System-Product-Name: ~  
root@chitrakant-System-Product-Name:~# [REDACTED]  
root@chitrakant-System-Product-Name:~# ssh localhost [REDACTED] 10  
The authenticity of host 'localhost (127.0.0.1)' can't be established.  
ECDSA key fingerprint is 6a:25:a0:8a:5d:7f:55:3c:97:b8:d5:7f:9e:a7:a3:80.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added 'localhost' (ECDSA) to the list of known hosts.  
root@localhost's password: [REDACTED] 11 Provide password
```

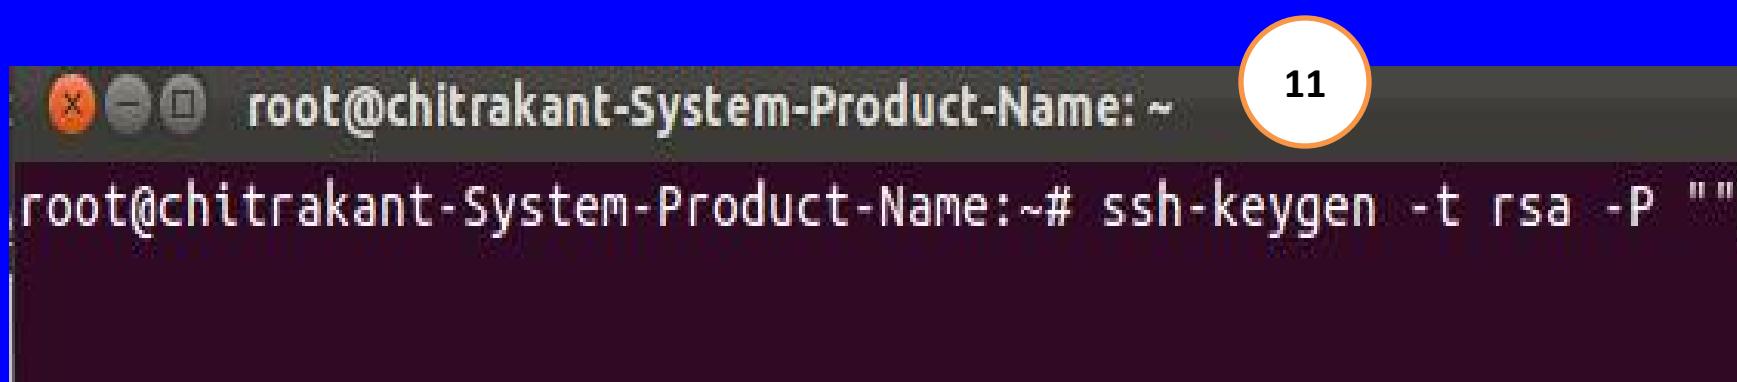
- ❖ Asks for the password and connected.
- ❖ But it should be able to connect without providing password.

# Generate an SSH key

- First, we have to generate an SSH key for the user account
- Here we use for root account.

## ❖ ssh-keygen program:

- The ssh-keygen program generates a public/private key pair.
- That generated key allows password-less access to any remote server which has our public key stored in its authorized\_keys file.



A screenshot of a terminal window titled "root@chitrakant-System-Product-Name: ~". The window shows the command "ssh-keygen -t rsa -P "" being typed at the prompt. A white circle with the number "11" is overlaid on the top right corner of the terminal window.

```
root@chitrakant-System-Product-Name: ~# ssh-keygen -t rsa -P ""
```

- It asks for the file name, but for simple case do not provide any file name.

```
root@chitrakant-System-Product-Name: ~
root@chitrakant-System-Product-Name:~# ssh-keygen -t rsa -P ""
Generating public/private rsa key pair.
Enter file in which to save the key (/root/.ssh/id_rsa):
```

12

Just press enter

```
x - root@chitrakant-System-Product-Name: ~
root@chitrakant-System-Product-Name:~# ssh-keygen -t rsa -P ""
Generating public/private rsa key pair.
Enter file in which to save the key (/root/.ssh/id_rsa):
Your identification has been saved in /root/.ssh/id_rsa.
Your public key has been saved in /root/.ssh/id_rsa.pub.
The key fingerprint is:
cb:27:f6:5b:5a:d7:11:ac:b7:48:02:3e:60:4a:40:71 root@chitrakant-System-Product-N
ame
The key's randomart image is:
+--[ RSA 2048]----+
| .+.E
|   o   .
|   . o .   o
|   . o o .   .
|   . S o . o o
|   . . . o o.o
|   = . o.....
|   . + + .
|   +
+-----+
root@chitrakant-System-Product-Name:~#
```

- ❖ Enable SSH access to our local machine with this newly created key:
- To access ssh command without requiring a password:
- Following steps are used for enabling public-key authentication, but instead of entering a passphrase when prompted, simply press Enter.

```
root@chitrakant-System-Product-Name:~# echo $HOME  
/root  
root@chitrakant-System-Product-Name:~# cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_keys
```

```
root@chitrakant-System-Product-Name:~# echo $HOME  
/root  
root@chitrakant-System-Product-Name:~# cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/a  
uthorized_keys  
root@chitrakant-System-Product-Name:~# ssh localhost  
Welcome to Ubuntu 11.10 (GNU/Linux 3.0.0-12-generic i686)  
  
 * Documentation: https://help.ubuntu.com/  
  
0 packages can be updated.  
0 updates are security updates.  
  
Last login: Tue Jan  1 07:14:18 2002 from localhost  
root@chitrakant-System-Product-Name:~# █
```

## **Step - 2**

## **Step – 3**

# **Disabling IPv6**

# Why?

- There's no practical point in enabling IPv6 on a box when we are not connected to any IPv6 network.
- Hence, simply disabled IPv6 on Ubuntu machine.

# How To Disable IPv6 Under Ubuntu 10.04/10.10/11.04

1. On Ubuntu (10.04/10.10/11.04), launch the terminal and issue this command to check whether IPv6 is enabled or not:

```
$ cat /proc/sys/net/ipv6/conf/all/disable_ipv6
```

0 --> Enabled

1 --> Disabled

```
root@chitrakant-System-Product-Name:~# cat /proc/sys/net/ipv6/conf/all/disable_ipv6
```

```
root@chitrakant-System-Product-Name:~# cat /proc/sys/net/ipv6/conf/all/disable_ipv6
```

```
root@chitrakant-System-Product-Name:~# nano /etc/sysctl.conf
root@chitrakant-System-Product-Name:~# nano /etc/sysctl.conf
root@chitrakant-System-Product-Name:~# sysctl -p
net.ipv6.conf.all.disable_ipv6 = 1
net.ipv6.conf.default.disable_ipv6 = 1
net.ipv6.conf.lo.disable_ipv6 = 1
root@chitrakant-System-Product-Name:~# cat /proc/sys/net/ipv6/conf/all/disable_i
pv6
1
root@chitrakant-System-Product-Name:~#
```

## 2. To disable IPv6, we need to edit the 'sysctl.conf' file.

- So, via the terminal, issue this command:

2

```
root@chitrakant-System-Product-Name:~# nano /etc/sysctl.conf
```



root@chitrakant-System-Product-Name: ~

GNU nano 2.2.6

File: /etc/sysctl.conf

```
# Do not send ICMP redirects (we are not a router)
#net.ipv4.conf.all.send_redirects = 0
#
# Do not accept IP source route packets (we are not a router)
#net.ipv4.conf.all.accept_source_route = 0
#net.ipv6.conf.all.accept_source_route = 0
#
# Log Martian Packets
#net.ipv4.conf.all.log_martians = 1
#
##### added to disable IPV6 #####
net.ipv6.conf.all.disable_ipv6 = 1
net.ipv6.conf.default.disable_ipv6 = 1
net.ipv6.conf.lo.disable_ipv6 = 1
#####
#
```

3

3. Issue now this command to reload your configuration:

```
root@chitrakant-System-Product-Name:~# sysctl -p
```

4

IPv6 is now disabled, you can check its status with the first command given above (1 --> disabled).

```
root@chitrakant-System-Product-Name:~# sysctl -p  
net.ipv6.conf.all.disable_ipv6 = 1  
net.ipv6.conf.default.disable_ipv6 = 1  
net.ipv6.conf.lo.disable_ipv6 = 1  
root@chitrakant-System-Product-Name:~#
```

```
root@chitrakant-System-Product-Name:~# cat /proc/sys/net/ipv6/conf/all/disable_ipv6
```

5

# **Step - 3**

## **Step – 4**

# **HADOOP installation and configuration**

# 1. Installation

- Download Hadoop from the Apache Download Mirrors.



❖ create directory to hold hadoop installer:

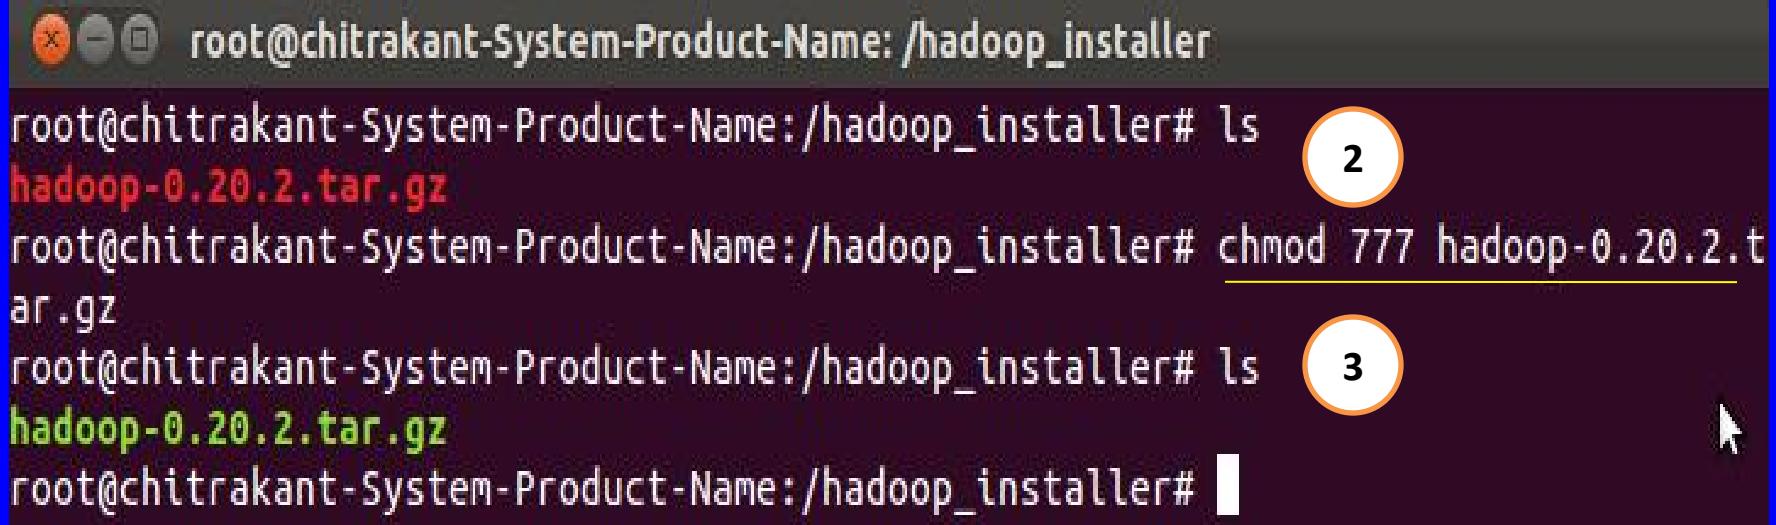
```
root@chitrakant-System-Product-Name: /  
root@chitrakant-System-Product-Name:# ls  
bin dev initrd.img lost+found openssh root selinux tmp vmlinuz  
boot etc java media opt run srv usr  
cdrom home lib mnt proc sbin sys var  
root@chitrakant-System-Product-Name:# mkdir hadoop_installer 1  
root@chitrakant-System-Product-Name:# ls  
bin etc java mnt root srv var  
boot hadoop_installer lib openssh run sys vmlinuz  
cdrom home lost+found opt sbin tmp  
dev initrd.img media proc selinux usr  
root@chitrakant-System-Product-Name:#
```

- ❖ Copy hadoop installer into the target folder and make it executable:

```
x - root@chitrakant-System-Product-Name: /hadoop_installer
root@chitrakant-System-Product-Name:/hadoop_installer# ls
hadoop-0.20.2.tar.gz
root@chitrakant-System-Product-Name:/hadoop_installer# chmod 777 hadoop-0.20.2.t
ar.gz
root@chitrakant-System-Product-Name:/hadoop_installer# ls
hadoop-0.20.2.tar.gz
root@chitrakant-System-Product-Name:/hadoop_installer#
```

2

3



❖ Uncompress it using tar command:

```
root@chitrakant-System-Product-Name:/hadoop_installer  
root@chitrakant-System-Product-Name:/hadoop_installer# ls  
hadoop-0.20.2.tar.gz 4  
root@chitrakant-System-Product-Name:/hadoop_installer# tar xzf hadoop-0.20.2.tar  
.gz
```

## ❖ Hadoop components in the destination folder:

```
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2
root@chitrakant-System-Product-Name:/hadoop_installer# ls
hadoop-0.20.2 hadoop-0.20.2.tar.gz
root@chitrakant-System-Product-Name:/hadoop_installer# cd hadoop-0.20.2/
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# ls
bin          docs           ivy          README.txt
build.xml    hadoop-0.20.2-ant.jar   ivy.xml     src
c++          hadoop-0.20.2-core.jar  lib          webapps
CHANGES.txt  hadoop-0.20.2-examples.jar librecordio
conf         hadoop-0.20.2-test.jar   LICENSE.txt
contrib      hadoop-0.20.2-tools.jar  NOTICE.txt
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2#
```

5

## 2. Update \$HOME/.bashrc

### ❖ Set Hadoop environment in .bashrc file:

```
root@chitrakant-System-Product-Name: ~
```

```
GNU nano 2.2.6
```

```
File: .bashrc
```

```
export JAVA_HOME="/java/jdk1.6.0_31"
export PATH=$PATH:$JAVA_HOME/bin

# or use symbolic link
# export JAVA_HOME="/java/jdk_6"

#####
##### Hadoop environment setting #####
#####

export HADOOP_HOME=/hadoop_installer/hadoop-0.20.2
# Add Hadoop bin/ directory to PATH
export PATH=$PATH:$HADOOP_HOME/bin

#####
# aliases for running Hadoop-related commands
unalias fs &> /dev/null
alias fs="hadoop fs"
unalias hls &> /dev/null
alias hls="fs -ls"
#####

#
```

6

### 3. Configuration

- **hadoop-env.sh**
- **conf/\*-site.xml**

```
x - root@chitrakant-System-Product-Name: /hadoop_installer/hadoop-0.20.2/conf 7
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2/conf# ls
capacity-scheduler.xml      hadoop-policy.xml    slaves
configuration.xsl           hdfs-site.xml       ssl-client.xml.example
core-site.xml                log4j.properties   ssl-server.xml.example
hadoop-env.sh                mapred-site.xml
hadoop-metrics.properties   masters
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2/conf#
```

# The Hadoop XML Configuration Files

- Hadoop's primary configuration files:
  1. Core-site.xml,
  2. hdfs-site.xml,
  3. mapred-site.xml.
- Hadoop uses a simple XML file format for configuration files

# Sample Hadoop XML configuration file

```
<configuration>
    <!-- Set 'some.property.name' to the value 'some-value'.
    <property>
        <name>some.property.name</name>
        <value>some-value</value>
    </property>
    <!--Set 'foo.bar.baz' to the value '42' and prevent it from being overridden by marking it final.-->
    <property>
        <name>foo.bar.baz</name>
        <value>42</value>
        <final>true</final>
    </property>
    <!-- Additional property elements... -->
</configuration>
```

❖ These files control the configuration of the common libraries used by:

- Hadoop
- HDFS
- MapReduce

## Example : Hadoop XML configuration file

```
<?xml version="1.0"?>
<!-- core-site.xml -->
<configuration>
    <property>
        <name>fs.default.name</name>
        <value>hdfs://namenode/</value>
        <final>true</final>
    </property>
</configuration>
```

# Significance of configuration files

- ❖ Overrides built-in default values:
- Properties defined in each of these three files override built in default values which are contained within the main Hadoop jar file.

# hadoop-env.sh

- Environment-specific settings store in this file.
- Configure JAVA\_HOME in this file.

```
$ nano /conf/hadoop-env.sh
```

- ❖ Set the JAVA\_HOME environment variable to the Sun JDK/JRE 6 directory

```
# The java implementation to use.  
export JAVA_HOME=/jdk1
```

root@chitrakant-System-Product-Name: /hadoop\_installer/hadoop-0.20.2/conf

GNU nano 2.2.6

File: hadoop-env.sh

```
# Set Hadoop-specific environment variables here.
```

```
# The only required environment variable is JAVA_HOME. All others are
# optional. When running a distributed configuration it is best to
# set JAVA_HOME in this file, so that it is correctly defined on
# remote nodes.
```

Uncomment and provide java  
installation path

```
# The java implementation to use. Required.
```

```
# export JAVA_HOME=/usr/lib/j2sdk1.5-sun
```

8

```
# Extra Java CLASSPATH elements. Optional.
```

```
# export HADOOP_CLASSPATH=
```

```
# The maximum amount of heap to use, in MB. Default is 1000.
```

```
# export HADOOP_HEAPSIZE=2000
```

# **conf/\*-site.xml**

- **Directory configuration:**
- Now configure the directory where Hadoop will store its data files.
- The network ports it listens to, etc.
- **HDFS:**
- This setup will use Hadoop's Distributed File System, HDFS, even it has single node.

# Create directory for Hadoop data

- Here, we use the directory **/app/hadoop/tmp** in this setup.

```
root@chitrakant-System-Product-Name: /app/hadoop/tmp
root@chitrakant-System-Product-Name:/# pwd
/
root@chitrakant-System-Product-Name:/# mkdir -p /app/hadoop/tmp 1
root@chitrakant-System-Product-Name:/# ls
app  cdrom  hadoop_installer  java      media    opt     run      srv  usr
bin  dev    home            lib       mnt     proc    sbin    sys  var
boot etc   initrd.img      lost+found  openssh  root    selinux tmp  vmlinuz
root@chitrakant-System-Product-Name:/# cd app
root@chitrakant-System-Product-Name:/app# ls 2
hadoop
root@chitrakant-System-Product-Name:/app# cd hadoop/
root@chitrakant-System-Product-Name:/app/hadoop# ls 3
tmp
root@chitrakant-System-Product-Name:/app/hadoop# cd tmp/
root@chitrakant-System-Product-Name:/app/hadoop/tmp# ls
root@chitrakant-System-Product-Name:/app/hadoop/tmp# █
```

# **core-site.xml**

- Contains system-level Hadoop configuration items:
  - The HDFS URL,
  - The Hadoop temporary directory.

- ❖ Edit the file **\$HADOOP\_HOME/conf/core-site.xml** and make sure it looks like the following:

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<configuration>
    <property>
        <name>hadoop.tmp.dir</name>
        <value>/usr/local/hadoop/tmp</value>
    </property>
    <property>
        <name>fs.default.name</name>
        <value>hdfs://localhost:8020</value>
    </property>
</configuration>
```

Hadoop created tmp directory we  
create our own **/app/hadoop/tmp**

- ❖ The property **fs.default.name** is a HDFS filesystem URI

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<configuration>
    <property>
        <name>hadoop.tmp.dir</name>
        <value>/usr/local/hadoop/tmp</value>
    </property>
    <property>
        <name>fs.default.name</name>
        <value>hdfs://localhost:8020</value>
    </property>
</configuration>
```

- Host is the namenode's hostname or IP address

- Port is the port that the namenode will listen on for RPCs.

- ❖ If no port is specified, the default of 8020 is used

root@chitrakant-System-Product-Name: /hadoop\_installer/hadoop-0.20.2/conf

GNU nano 2.2.6

File: core-site.xml

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<!-- Put site-specific property overrides in this file. --&gt;

&lt;configuration&gt;
    &lt;property&gt;
        &lt;name&gt;hadoop.tmp.dir&lt;/name&gt;
        &lt;value&gt;/app/hadoop/tmp&lt;/value&gt;
        &lt;description&gt;A base for other temporary directories.&lt;/description&gt;
    &lt;/property&gt;
    &lt;property&gt;
        &lt;name&gt;fs.default.name&lt;/name&gt;
        &lt;value&gt;hdfs://localhost:54310&lt;/value&gt;
        &lt;description&gt;
            The name of the default file system.
            A URI whose scheme and authority determine
            the FileSystem implementation. The uri's
            scheme determines the config property
            (fs.SCHEME.impl) naming the FileSystem
            implementation class. The uri's authority
            is used to determine the host, port, etc.
            for a filesystem.
        &lt;/description&gt;
    &lt;/property&gt;
&lt;/configuration&gt;</pre>
```

# conf/mapred-site.xml

- ❖ In **mapred-site.xml** we specify the hostname and port of the JobTracker.

hostname (or IP address)

Port number on which the jobtracker  
listens for RPC communication

```
<!-- In: conf/mapred-site.xml -->
<property>
  <name>mapred.job.tracker</name>
  <value>localhost:54311</value>
  <description>
    The host and port that the MapReduce job tracker runs at.
    If "local", then jobs are run in-process as a single map and
    reduce task.
  </description>
</property>
```

## **mapred.job.tracker**

- **mapred.job.tracker provides the locations of:**
  - namenode,
  - datanodes
  - clients
- **for tasktrackers and MapReduce clients.**

GNU nano 2.2.6

File: mapred-site.xml

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<!-- Put site-specific property overrides in this file. --&gt;

&lt;configuration&gt;
    &lt;property&gt;
        &lt;name&gt;mapred.job.tracker&lt;/name&gt;
        &lt;value&gt;localhost:54311&lt;/value&gt;
        &lt;description&gt;
            The host and port that the MapReduce job tracker runs
            at. If "local", then jobs are run in-process as a single map
            and reduce task.
        &lt;/description&gt;
    &lt;/property&gt;
&lt;/configuration&gt;</pre>
```

# hdfs-site.xml

- Contains HDFS settings :
  - The default file replication count,
  - The block size,
  - Whether permissions are enforced.
- ❖ In **hdfs-site.xml** we specify the default replication factor for HDFS, which should only be **one** because we're running on only one node.

- ❖ Edit the file **\$HADOOP\_HOME/conf/hdfs-site.xml** and make sure it looks like the following:

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <!-- specify this so that running 'hadoop namenode -format' formats the right dir -->
    <name>dfs.name.dir</name>
    <value>/usr/local/hadoop/cache/hadoop/dfs/name</value>
  </property>
</configuration>
```

GNU nano 2.2.6

File: hdfs-site.xml

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<!-- Put site-specific property overrides in this file. --&gt;

&lt;configuration&gt;
    &lt;property&gt;
        &lt;name&gt;dfs.replication&lt;/name&gt;
        &lt;value&gt;1&lt;/value&gt;
        &lt;description&gt;
            Default block replication.
            The actual number of replications can be specified when the file is created.
            The default is used if replication is not specified in create time.
        &lt;/description&gt;
    &lt;/property&gt;
&lt;/configuration&gt;</pre>
```

# **Formatting the HDFS file-system via the NameNode**

- **Hadoop File System:**
- Implemented on top of the local file-system of our “cluster”.
- ( Single node cluster: on top of one local machine’s file system )
- ( Multi node cluster: on top of all machine’s file system )

- **Format the Hadoop File System:**
- Format only at first time when set up a Hadoop cluster.
- Format when system is running : data will be lost, not recommended.
- Format the file system: simply initializes the directory specified by `dfs.name.dir` variable.

```
$ /usr/local/hadoop/bin/hadoop namenode -format
```

Command for format

# doop\_installer/hadoop-0.20.2



```
/hadoop_installer/hadoop-0.20.2# bin/hadoop namenode -format
```

```
02/01/01 08:42:28 INFO namenode.NameNode: STARTUP_MSG:
```

```
*****
```

```
STARTUP_MSG: Starting NameNode
```

```
STARTUP_MSG: host = chitrakant-System-Product-Name/127.0.1.1
```

```
STARTUP_MSG: args = [-format]
```

```
STARTUP_MSG: version = 0.20.2
```

```
STARTUP_MSG: build = https://svn.apache.org/repos/asf/hadoop/common/branches/b  
'chrисdo' on Fri Feb 19 08:07:34 UTC 2010
```

```
*****
```

```
02/01/01 08:42:28 INFO namenode.FSNamesystem: fsOwner=root,root
```

```
02/01/01 08:42:28 INFO namenode.FSNamesystem: supergroup=supergroup
```

```
02/01/01 08:42:28 INFO namenode.FSNamesystem: isPermissionEnabled=true
```

```
02/01/01 08:42:28 INFO common.Storage: Image file of size 94 saved in 0 seconds.
```

```
02/01/01 08:42:29 INFO common Storage directory /app/hadoop/tmp/dfs/name has been successfully  
.
```

```
02/01/01 08:42:29 INFO namenode.NameNode: SHUTDOWN_MSG:
```

```
*****
```

```
SHUTDOWN_MSG: Shutting down NameNode at chitrakant-System-Product-Name/127.0.1.1
```

```
*****
```

```
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# 
```

```
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# 
```

# Starting single-node cluster

- Run the command:

```
$ /usr/local/hadoop/bin/start-all.sh
```

- This will startup:
  1. Namenode,
  2. Datanode,
  3. Jobtracker
  4. Tasktracker
  5. SecondaryNamenode

## ❖ Use jps to list out running hadoop processes

```
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# jps
4965 DataNode
5328 TaskTracker
5183 JobTracker
5119 SecondaryNameNode
5493 Jps
4787 NameNode
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# █
```

- Check with netstat if Hadoop is listening on the configured ports.

```
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2          8:45 AM Guest
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# netstat -plten | grep java
tcp      0      0 127.0.0.1:54311          0.0.0.0:*          LISTEN     0          348590  5183/java
tcp      0      0 0.0.0.0:50090           0.0.0.0:*          LISTEN     0          349131  5119/java
tcp      0      0 0.0.0.0:50060           0.0.0.0:*          LISTEN     0          349126  5328/java
tcp      0      0 0.0.0.0:44173           0.0.0.0:*          LISTEN     0          348460  5119/java
tcp      0      0 0.0.0.0:50030           0.0.0.0:*          LISTEN     0          348977  5183/java
tcp      0      0 0.0.0.0:40082           0.0.0.0:*          LISTEN     0          349262  4965/java
tcp      0      0 0.0.0.0:53173           0.0.0.0:*          LISTEN     0          347024  4787/java
tcp      0      0 0.0.0.0:50070           0.0.0.0:*          LISTEN     0          348724  4787/java
tcp      0      0 0.0.0.0:50010           0.0.0.0:*          LISTEN     0          348837  4965/java
tcp      0      0 0.0.0.0:50075           0.0.0.0:*          LISTEN     0          350066  4965/java
tcp      0      0 127.0.0.1:52926          0.0.0.0:*          LISTEN     0          350318  5328/java
tcp      0      0 0.0.0.0:48035           0.0.0.0:*          LISTEN     0          348587  5183/java
tcp      0      0 0.0.0.0:50020           0.0.0.0:*          LISTEN     0          350149  4965/java
tcp      0      0 127.0.0.1:54310          0.0.0.0:*          LISTEN     0          347025  4787/java

root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2#
```

# Stopping single-node cluster

```
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# bin/stop-all.sh
stopping jobtracker
localhost: stopping tasktracker
stopping namenode
localhost: stopping datanode
localhost: stopping secondarynamenode
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# █
```

- ❖ To stop all the daemons running on the machine.

# **Step - 4**

# Step – 5

## Installing eclipse

## ❖ Download eclipse installer



```
root@chitrakant-System-Product-Name: /eclipse
root@chitrakant-System-Product-Name:/# cd eclipse/
root@chitrakant-System-Product-Name:/eclipse# cp eclipse-jee-helios-SR2-linux-gtk.tar.gz .
```

2

## ❖ Uncompress the eclipse installer file

```
root@chitrakant-System-Product-Name:/eclipse
```

3

```
root@chitrakant-System-Product-Name:/eclipse# tar xzf eclipse-jee-helios-SR2-linux-gtk.tar.gz
```

## ❖ create directory eclipse in /opt directory

```
root@chitrakant-System-Product-Name: /eclipse_installer  
root@chitrakant-System-Product-Name: /opt# ls 4  
eclipse
```

## ❖ move uncompressed eclipse into eclipse directory of /opt directory

```
root@chitrakant-System-Product-Name: /eclipse_installer  
root@chitrakant-System-Product-Name: /opt# cd eclipse/ 5  
root@chitrakant-System-Product-Name: /opt/eclipse# ls  
about_files  configuration  eclipse.ini  icon.xpm      p2  
about.html    dropins       epl-v10.html  libcairo-swt.so  plugins  
artifacts.xml  eclipse       features     notice.html  readme
```

## ❖ Provide execution permission to the eclipse executable

```
root@chitrakant-System-Product-Name: /opt/eclipse
root@chitrakant-System-Product-Name:/opt/eclipse# ls
about_files      configuration   eclipse.ini    icon.xpm      p2
about.html        dropins        epl-v10.html  libcairo-swt.so  plugins
artifacts.xml    eclipse       features       notice.html  readme
root@chitrakant-System-Product-Name:/opt/eclipse# chmod 777 eclipse
```

6

## ❖ Create an executable in home path:

 root@chitrakant-System-Product-Name: /opt/eclipse

7

```
root@chitrakant-System-Product-Name: /opt/eclipse# touch /usr/bin/eclipse  
root@chitrakant-System-Product-Name: /opt/eclipse# chmod 777 /usr/bin/eclipse  
root@chitrakant-System-Product-Name: /opt/eclipse# nano /usr/bin/eclipse
```

8

- Write following steps in the eclipse file and exit from the editor

```
#!/bin/sh
```

```
ECLIPSE_HOME="/opt/eclipse" $ECLIPSE_HOME/eclipse $*
```

root@chitrakant-System-Product-Name: /opt/eclipse

GNU nano 2.2.6

File: /usr/bin/eclipse

```
#!/bin/sh
```

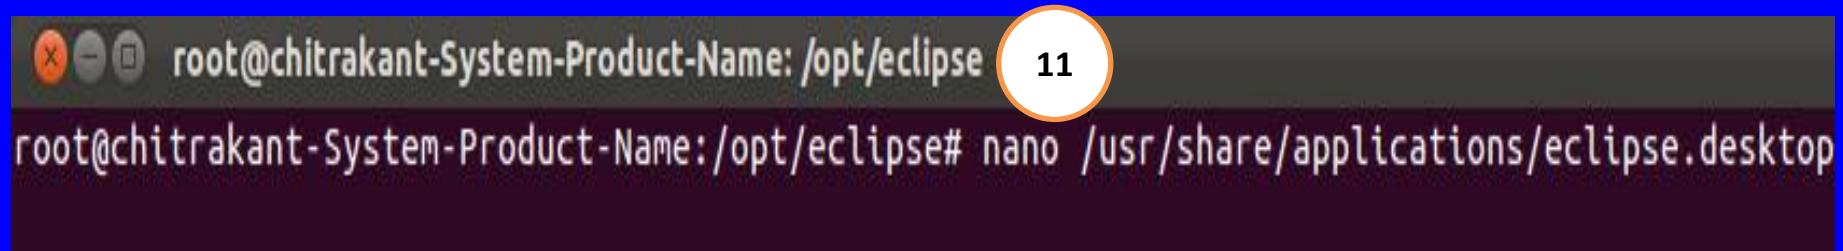
```
export ECLIPSE_HOME="/opt/eclipse"
$ECLIPSE_HOME/eclipse $*
```

9

## ❖ Make eclipse executable everywhere by creating a symlink

```
root@chitrakant-System-Product-Name: /opt/eclipse 10
root@chitrakant-System-Product-Name:/opt/eclipse# ln -s /usr/bin/eclipse /bin/eclipse
```

## ❖ create menu icon on the launcher



```
root@chitrakant-System-Product-Name: /opt/eclipse 11
root@chitrakant-System-Product-Name:/opt/eclipse# nano /usr/share/applications/eclipse.desktop
```

- Write following contents into eclipse.Desktop file and save

```
[Desktop Entry]
Encoding=UTF-8
Name=Eclipse
Comment=Eclipse IDE
Exec=eclipse
Icon=/opt/eclipse/icon.xpm
Terminal=false
Type=Application
Categories=GNOME;Application;Development;
StartupNotify=true
```

root@chitrakant-System-Product-Name: /opt/eclipse

GNU nano 2.2.6

File: /usr/share/applications/eclipse.desktop

```
[Desktop Entry]
Encoding=UTF-8
Name=Eclipse
Comment=Eclipse IDE
Exec=eclipse
Icon=/opt/eclipse/icon.xpm
Terminal=false
Type=Application
Categories=GNOME;Application;Development;
StartupNotify=true
```

12

## ❖ Specifying the JVM

```
root@chitrakant-System-Product-Name: /opt/eclipse
```

```
root@chitrakant-System-Product-Name: /opt/eclipse# nano eclipse.ini
```

```
root@chitrakant-System-Product-Name: /opt/eclipse# nano eclipse.ini
```

```
root@chitrakant-System-Product-Name: /opt/eclipse
GNU nano 2.2.6          File: eclipse.ini

-startup
plugins/org.eclipse.equinox.launcher_1.1.1.R36x_v20101122_1400.jar
--launcher.library
plugins/org.eclipse.equinox.launcher.gtk.linux.x86_1.1.2.R36x_v20101019_1345
-product
org.eclipse.epp.package.jee.product
--launcher.defaultAction
openFile
-showsplash
org.eclipse.platform
--launcher.XXMaxPermSize
256m
--launcher.defaultAction
openFile
-vmargs
-Dosgi.requiredJavaVersion=1.5
-XX:MaxPermSize=256m
-Xms40m
-Xmx512m
```

root@chitrakant-System-Product-Name: /opt/eclipse

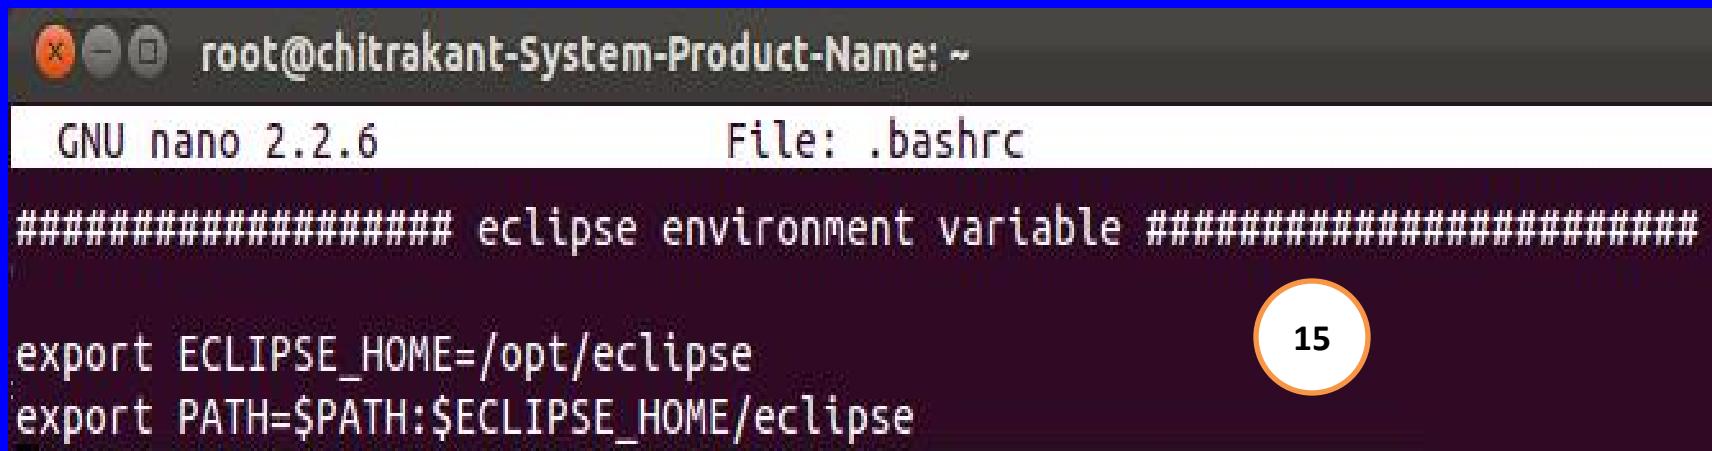
GNU nano 2.2.6

File: eclipse.ini

Modified

```
-startup
plugins/org.eclipse.equinox.launcher_1.1.1.R36x_v20101122_1400.jar
--launcher.library
plugins/org.eclipse.equinox.launcher.gtk.linux.x86_1.1.2.R36x_v20101019_1345
-product
org.eclipse.epp.package.jee.product
--launcher.defaultAction
openFile
-showsplash
org.eclipse.platform
--launcher.XXMaxPermSize
256m
--launcher.defaultAction
openFile
-vmargs
/java/jdk1.6.0_31/bin/java
-Dosgi.requiredJavaVersion=1.5
-XX:MaxPermSize=256m
-Xms40m
```

## ❖ Set eclipse environment variables in the .bashrc file



root@chitrakant-System-Product-Name: ~

GNU nano 2.2.6 File: .bashrc

```
##### eclipse environment variable #####
export ECLIPSE_HOME=/opt/eclipse
export PATH=$PATH:$ECLIPSE_HOME/eclipse
```

15

## ❖ Set eclipse launcher on the panel

Step 1 : make directory **jre** inside **/opt/eclipse** directory

Step 2 : make directory **bin** inside **/opt/eclipse/jre** directory

Step 3 : create symbolic link for the **/jdk/java/bin/java** inside **/opt/eclipse/jre/bin**

x - root@chitrakant-System-Product-Name: /opt/eclipse 16

root@chitrakant-System-Product-Name: /opt/eclipse# mkdir jre  
root@chitrakant-System-Product-Name: /opt/eclipse/jre# mkdir bin 17

x - root@chitrakant-System-Product-Name: /opt/eclipse/jre/bin 18

root@chitrakant-System-Product-Name: /opt/eclipse/jre# cd bin  
root@chitrakant-System-Product-Name: /opt/eclipse/jre/bin# ln -s /java/jdk1.6.0\_31/bin/java java

19

x - root@chitrakant-System-Product-Name: /opt/eclipse/jre/bin 20

root@chitrakant-System-Product-Name: /opt/eclipse/jre# cd bin  
root@chitrakant-System-Product-Name: /opt/eclipse/jre/bin# ls  
java 21



Required\_Softwares



Tutorials



# **Step - 5**

```
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# bin/hadoop dfs -put conf input  
root@chitrakant-System-Product-Name:/hadoop_installer/hadoop-0.20.2# bin/hadoop jar hadoop-*-examples.jar grep input output 'dfs[a-z.]+'  
02/01/01 00:26:18 INFO mapred.FileInputFormat: Total input paths to process : 13  
02/01/01 00:26:18 INFO mapred.JobClient: Running job: job_200201010003_0001  
02/01/01 00:26:19 INFO mapred.JobClient: map 0% reduce 0%  
02/01/01 00:26:30 INFO mapred.JobClient: map 15% reduce 0%  
02/01/01 00:26:36 INFO mapred.JobClient: map 30% reduce 0%  
02/01/01 00:26:39 INFO mapred.JobClient: map 46% reduce 10%  
02/01/01 00:26:42 INFO mapred.JobClient: map 61% reduce 10%  
02/01/01 00:26:45 INFO mapred.JobClient: map 76% reduce 10%  
02/01/01 00:26:48 INFO mapred.JobClient: map 92% reduce 15%  
02/01/01 00:26:51 INFO mapred.JobClient: map 100% reduce 15%  
02/01/01 00:26:57 INFO mapred.JobClient: map 100% reduce 100%  
02/01/01 00:26:59 INFO mapred.JobClient: Job complete: job_200201010003_0001  
02/01/01 00:26:59 INFO mapred.JobClient: Counters: 18  
02/01/01 00:26:59 INFO mapred.JobClient: Job Counters  
02/01/01 00:26:59 INFO mapred.JobClient: Launched reduce tasks=1  
02/01/01 00:26:59 INFO mapred.JobClient: Launched map tasks=13  
02/01/01 00:26:59 INFO mapred.JobClient: Data-local map tasks=13  
02/01/01 00:26:59 INFO mapred.JobClient: FileSystemCounters  
02/01/01 00:26:59 INFO mapred.JobClient: FILE_BYTES_READ=158  
02/01/01 00:26:59 INFO mapred.JobClient: HDFS_BYTES_READ=19146  
02/01/01 00:26:59 INFO mapred.JobClient: FILE_BYTES_WRITTEN=804  
02/01/01 00:26:59 INFO mapred.JobClient: HDFS_BYTES_WRITTEN=280  
02/01/01 00:26:59 INFO mapred.JobClient: Map-Reduce Framework  
02/01/01 00:26:59 INFO mapred.JobClient: Reduce input groups=7  
02/01/01 00:26:59 INFO mapred.JobClient: Combine output records=7  
02/01/01 00:26:59 INFO mapred.JobClient: Map input records=580  
02/01/01 00:26:59 INFO mapred.JobClient: Reduce shuffle bytes=230  
02/01/01 00:26:59 INFO mapred.JobClient: Reduce output records=7  
02/01/01 00:26:59 INFO mapred.JobClient: Spilled Records=14  
02/01/01 00:26:59 INFO mapred.JobClient: Map output bytes=193  
02/01/01 00:26:59 INFO mapred.JobClient: Map input bytes=19146  
02/01/01 00:26:59 INFO mapred.JobClient: Combine input records=10
```

# **Step - III**

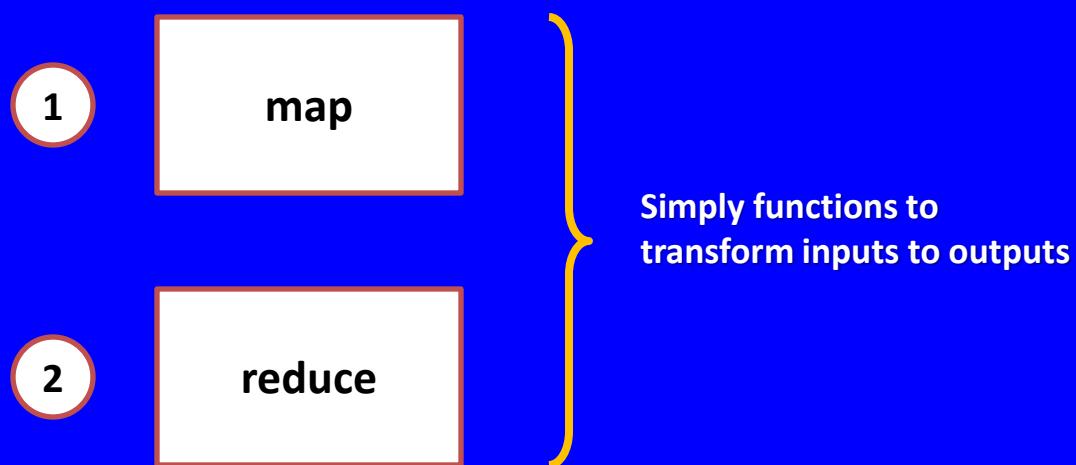
# Hadoop : Fourth step

## MapReduce Programming

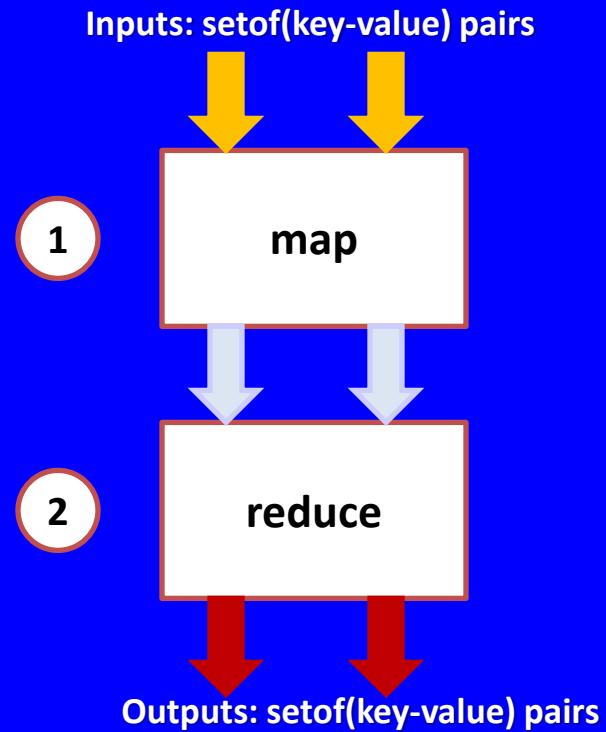
# Understanding MapReduce Program

# MapReduce

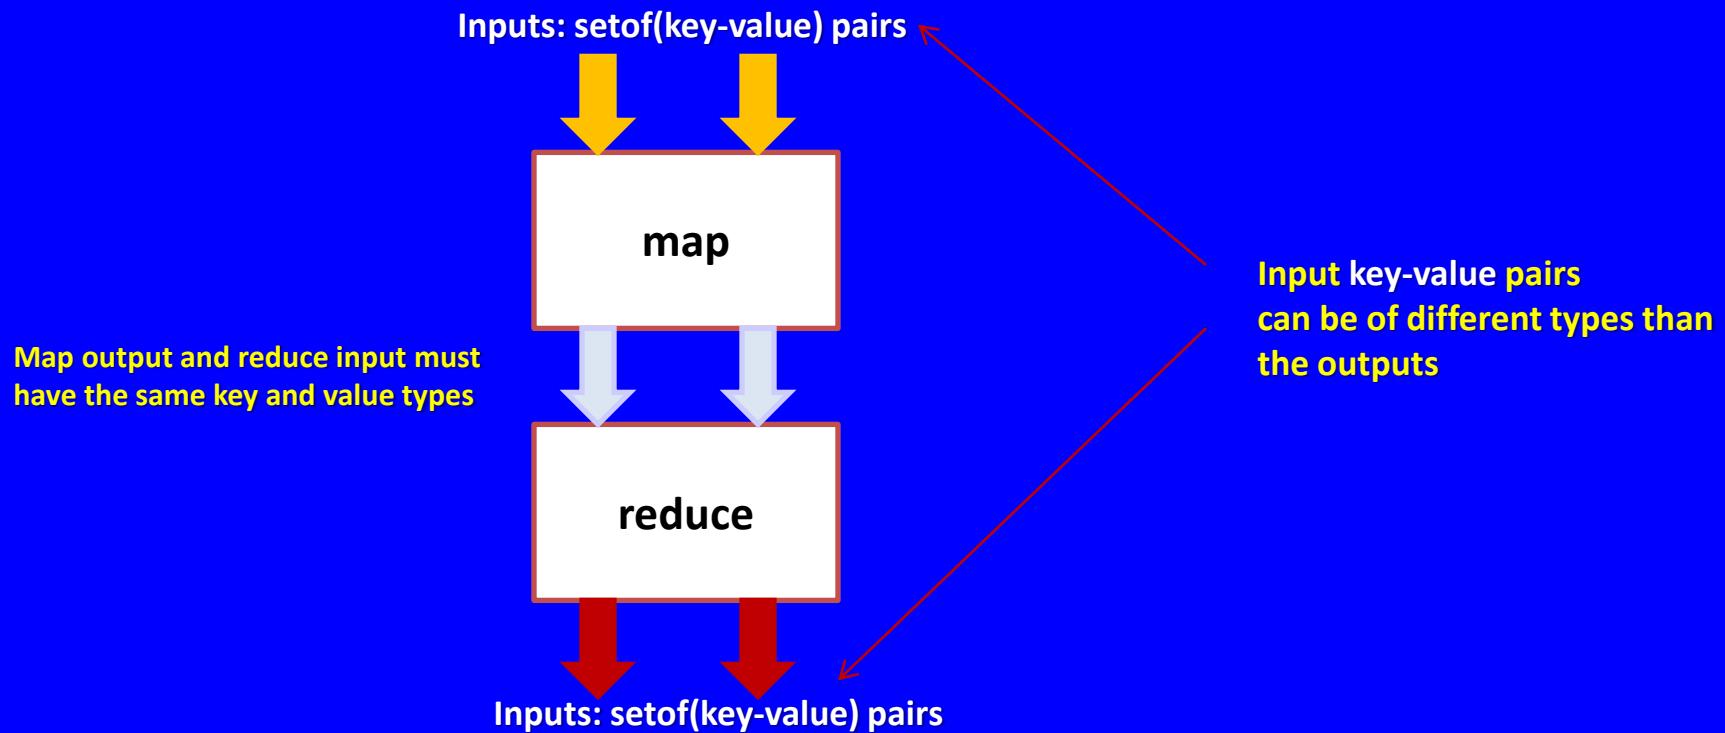
- A MapReduce job proceeds in two phases



- Inputs and outputs are key – value pairs



- MapReduce is typed



# Sample Input File

(0,	" book	2007	20	18	15")
(20	" book	2008	22	20	12")
(40,	"doctor	2007	545525	366136	57313")
(72,	"doctor	2008	668666	446034	72694")

- Input comprises lines of text
- Represented as key-value pairs in Hadoop

Key

Value

- |  |                                 |
|--|---------------------------------|
| ✓ The offset of the line within the file | ❖ The value is the line of text |
|--|---------------------------------|

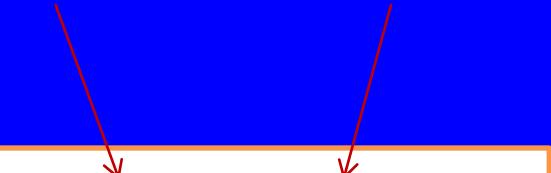
- Ignore other fields and the key from input file.
- Only word and the number of occurrences required.

(0,	" book	2007	20	18	15")
(20	" book	2008	22	20	12")
(40,	"doctor	2007	545525	366136	57313")
(72,	"doctor	2008	668666	446034	72694")

- The map emits →

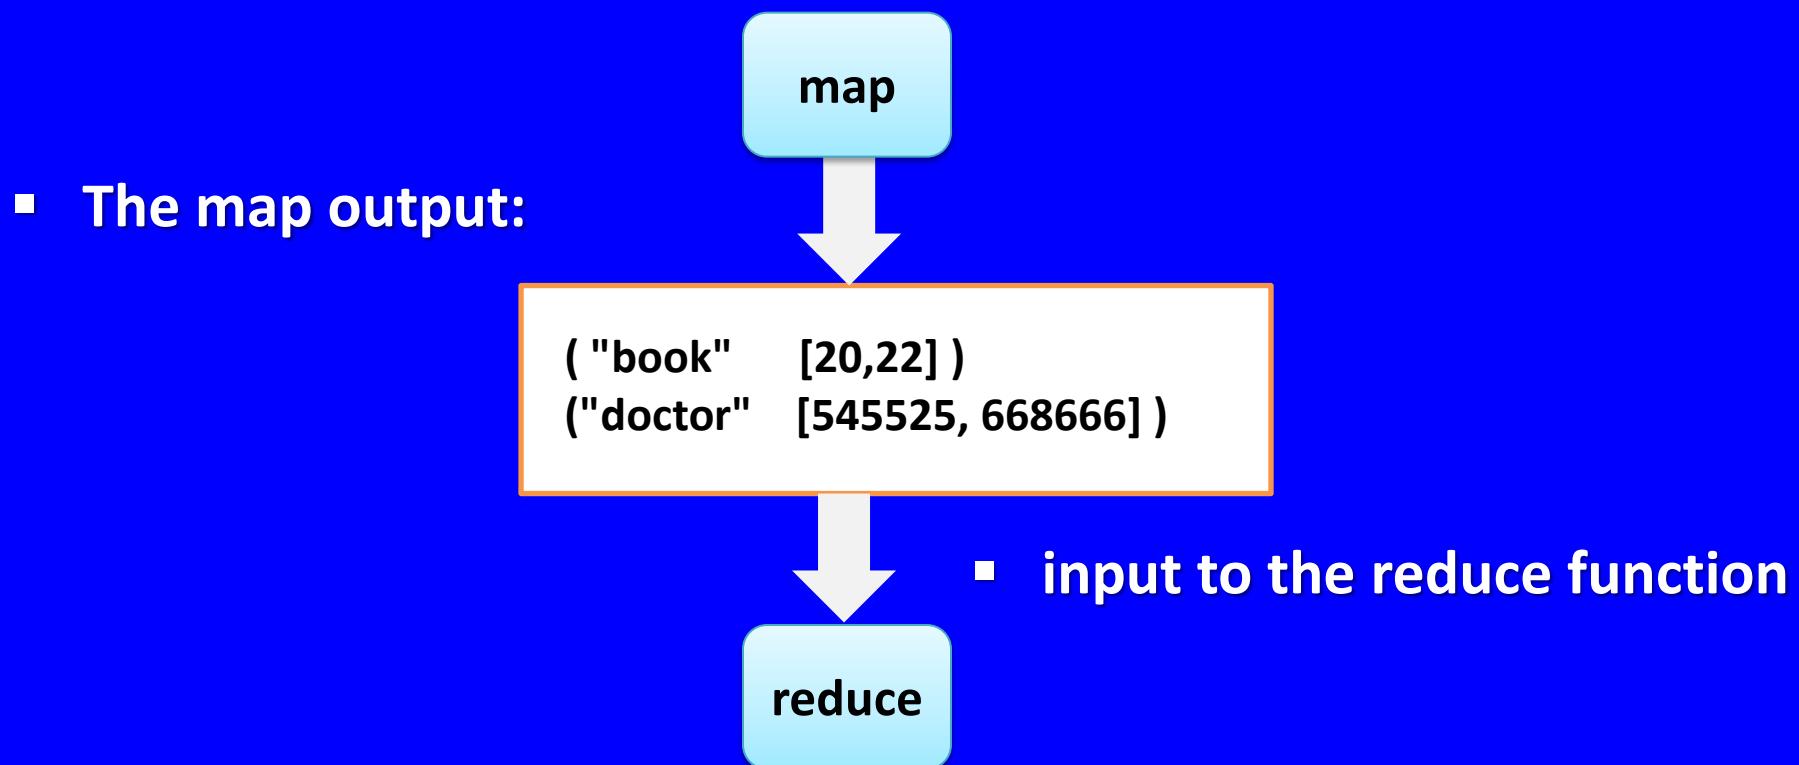
Key: word

Value : No. of occurrences



( "book" 20 )  
 (" book" 22)  
 ("doctor" 545525)  
 ("doctor" 668666)

- **Map does** → The values for a given key are brought together so that the reduce function processes them as a group



**Reduce does →** The reduce function iterates over the values for each key and does whatever processing it likes.

- This example → An aggregate count (simply sum the values)

```
( "book" [42] )  
( "doctor" [1214191] )
```

- Final result → The number of times each word occurred in the whole dataset

# Java MapReduce

- Write the MapReduce job to find the total count for each word.

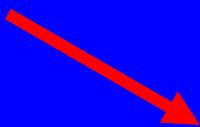
# map function

- Represented in Java by an instance of
  - ✓ **org.apache.hadoop.mapreduce.Mapper**
- **Decision about input key-value pairs and output key-value pairs types**

```
public class Mapper< KEYIN, VALUEIN, KEYOUT, VALUEOUT >
```

Declaration of the Mapper class

- **TextInputFormat** → Because we are processing text



Determines the input types as **LongWritable** and **Text**

(in package: [org.apache.hadoop.io](#) )

- Wrappers around standard Java types (`long` and `String`)

- TextInputFormat presents the input to our mapper as (LongWritable, Text) pairs

Key → The offset within the file

(0,	" book	2007	20	18	15")
(20	" book	2008	22	20	12")
(40,	"doctor	2007	545525	366136	57313")
(72,	"doctor	2008	668666	446034	72694")

value → The content of the line

# Mapper

- Extract the word and the number of occurrences, and ignore everything else
- Mapper output is (word, count) pairs, of type (Text, LongWritable)

```
public class CountMapper extends Mapper<LongWritable, Text, Text, LongWritable>
```

- Our task → Write the implementation of the map() method

```
package myPackage;
```

```
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;  
import java.io.IOException;
```

Import Library Packages

Every Map class must extend Mapper

```
public class CountMapper extends  
    Mapper<LongWritable, Text, Text, LongWritable> {
```

```
}
```

```
package myPackage;
```

```
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;  
import java.io.IOException;
```

Import Library Packages

Every Map class must extend Mapper

```
public class CountMapper extends  
    Mapper<LongWritable, Text, Text, LongWritable> {
```

```
protected void map(LongWritable key, Text value, Context context)  
    throws IOException, InterruptedException {
```

```
}
```

Every Mapper class must  
provide definition of  
map function

```
package myPackage;

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import java.io.IOException;

public class CountMapper extends
    Mapper<LongWritable, Text, Text, LongWritable> {

    private Text word = new Text();
    private LongWritable count = new LongWritable();

    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
    }

}
```

**Instance variables**  
To store map output  
key and value

```

package myPackage;

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import java.io.IOException;

public class CountMapper extends
    Mapper<LongWritable, Text, Text, LongWritable> {

    private Text word = new Text();
    private LongWritable count = new LongWritable();

    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        /* Body of map() → splits the tab – separated input line into fields
         * and uses the first field as the word and the third as the count
        */
    }
}

```

(0,	"book	2007	20	18	15")
(20,	"book	2008	22	20	12")
(40,	"doctor	2007	545525	366136	57313")
(72,	"doctor	2008	668666	446034	72694")
[0]	[1]	[2]	[3]	[4]	[5]

```
package myPackage;

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import java.io.IOException;

public class CountMapper extends
    Mapper<LongWritable, Text, Text, LongWritable> {

    private Text word = new Text();
    private LongWritable count = new LongWritable();

    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        context.write(word, count);
        /*The map output is written using the write method in Context*/
    }
}
```

```
public class CountMapper extends
    Mapper<LongWritable, Text, Text, LongWritable> {

    private Text word = new Text();
    private LongWritable count = new LongWritable();

    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        /* value is tab separated values: word, year, occurrences, #books, #pages
           we project out (word, occurrences) so we can sum over all years */

        String[] split = value.toString().split("\t+");
        word.set(split[0]);
        if (split.length > 2) {
            try {
                count.set(Long.parseLong(split[2]));
                context.write(word, count);
            } catch (NumberFormatException e) {
                // cannot parse - ignore
            }
        } //if
    }
}
```

- Running through our dataset, the map output looks like this:

```
( "book"    20 )  
(" book"   22)  
("doctor"  545525)  
("doctor"  668666)
```

# Reduce

```
( "book"  [20,22] )  
("doctor"  [545525, 668666])
```

```
package myPackage;
```

```
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Reducer;  
import java.io.IOException;
```

Import Library Packages

Every Reducer class must extend Reducer

```
public class CountReducer extends  
    Reducer<Text, LongWritable, Text, LongWritable> {
```

```
protected void reduce(KEY key, Iterable<LongWritable>values,  
    Context context) throws IOException, InterruptedException {
```

```
}
```

Every Reducer class must  
provide definition of  
reduce function

- Notice that the `reduce()` method signature is different from `map()`
- It has an iterator over the values, rather than a single value.
- This reflects the grouping that the framework performs on the values for a key.

```
public class CountReducer extends  
    Reducer<Text, LongWritable, Text, LongWritable> {  
  
    protected void reduce(KEY key, Iterable<LongWritable>values,  
        Context context) throws IOException, InterruptedException {  
  
    }  
}
```

```
public class CountReducer<KEY> extends Reducer<KEY, LongWritable,  
    KEY, LongWritable> {
```

```
    private LongWritable result = new LongWritable();
```

```
    public void reduce(KEY key, Iterable<LongWritable> values,  
        Context context) throws IOException, InterruptedException {
```

```
        long sum = 0;  
  
        for (LongWritable val : values) {  
            sum += val.get();  
        }  
  
        result.set(sum);  
  
        context.write(key, result);
```



It sums the values, then writes the total out using the same key as the input

```
}
```

- The output of the reducer will be:

```
( "book"    [42] )  
("doctor"   [1214191])
```

## Code the MapReduce Application/Driver class (CountJob.java)

- Define a driver class.
- Create a new client job, configuration object and advertise Mapper and Reducer classes.
- The driver class is responsible for setting our MapReduce job to run in Hadoop.
- We specify job name, data type of input/output and names of mapper and reducer classes.

```
package myPackage;

import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.reduce.LongSumReducer;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
```

- CountJob extends Hadoop's Configured class and implements the Tool interface's run() method
- This allows us to pass in configuration options at runtime



```
public class CountJob extends Configured implements Tool {
```

```
    public int run(String[] args) throws Exception {
```

```
    }
```

```
}
```

- The code to launch a job is in the run() method.
- It is managed by the Job class from the org.apache.hadoop.mapreduce package.

```
public class CountJob extends Configured implements Tool {
```

1    public static void main(String[] args) throws Exception {

}

2    public int run(String[] args) throws Exception {

}

}

```
public class CountJob extends Configured implements Tool {  
  
    public static void main(String[] args) throws Exception {  
        int rc = ToolRunner . run( new CountJob(), args );  
        System.exit(rc);  
    }  
}
```

```
public class CountJob extends Configured implements Tool {
```

```
    public int run(String[] args) throws Exception {
```

1

```
        /* The code to launch a job is in the run() method */  
        /* it is managed by the Job class from the org.apache.hadoop.mapreduce package. */
```

2

```
        /* The Job instance specifies various things about the job:  
           A name for display purposes,  
           The mapper and reducer classes  
           The job output types, (Have to match the mapper and reducer output types)  
        */
```

3

```
        setMapOutputKeyClass()  
        setMapOutputValueClass()
```

Should call when map output types are  
different from reduce output types

```
}
```

```
}
```

```
public class CountJob extends Configured implements Tool {
```

```
    public int run(String[] args) throws Exception {
```

2

```
        setJarByClass()
```

tell Hadoop to look for the local JAR that contains the **CountJob** class

4

```
        FileInputFormat . addInputPath( job, new Path( args[0] ));  
        FileOutputFormat . setOutputPath( job, new Path( args[1] ));
```



- Tell the job what input data to process and where to place the output.
- Do this via the static APIs on **FileInputFormat** and **FileOutputFormat**.
- Use command-line arguments to specify the file paths.

```
}
```

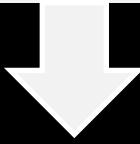
```
}
```

```
public class CountJob extends Configured implements Tool {
```

```
    public int run(String[] args) throws Exception {
```

5

- Launches the job and waits for it to complete.
- While it is running, it prints the progress on the console.



```
        return job.waitForCompletion( true ) ? 0 : 1;  
    }
```

```
}
```

```
public class CountJob extends Configured implements Tool {
```

```
    public int run(String[] args) throws Exception {
```

1

```
        Job job = new Job( getConf() );  
        job . setJarByClass( getClass() );
```

4

```
        FileInputFormat . addInputPath( job, new Path( args[0] ) );  
        FileOutputFormat . setOutputPath( job, new Path( args[1] ) );
```

2

```
        job . setJobName( getClass() .getSimpleName( ) );  
        job . setMapperClass( CountMapper.class );  
        job . setCombinerClass( CountReducer.class );  
        job . setReducerClass( CountReducer.class );
```

3

```
        job . setOutputKeyClass( Text . class );  
        job . setOutputValueClass( LongWritable . class );
```

5

```
    return job . waitForCompletion ( true ) ? 0 : 1 ;
```

```
}
```

```
}
```

```
public class CountJob extends Configured implements Tool {
```

```
    public static void main(String[] args) throws Exception {
        int rc = ToolRunner . run( new CountJob(), args );
        System.exit(rc);
    }
    public int run(String[] args) throws Exception {
        Job job = new Job( getConf() );
        job . setJarByClass( getClass() );
        FileInputFormat . addInputPath( job, new Path( args[0] ) );
        FileOutputFormat . setOutputPath( job, new Path( args[1] ) );
        job . setJobName( getClass() .getSimpleName( ) );
        job . setMapperClass( CountMapper.class );
        job . setCombinerClass( CountReducer.class );
        job . setReducerClass( CountReducer.class );
        job . setOutputKeyClass( Text . class );
        job . setOutputValueClass( LongWritable . class );
    } return job . waitForCompletion ( true ) ? 0 : 1 ;
}
```

# Compiling the CountJob program

## CountJob

- Folder contains CountJob program

CountMapper.java

- Contains the map function implementation

CountReducer.java

- Contains the reduce function implementation

CountJob.java

- Contains the code coordinating the execution of the map and reduce functions

- Execute the following commands in the **CountJob** folder:

compiles the program using the classes developed by Hadoop (**hadoop-core-1.0.4.jar**)

- CountJob\$ **javac -cp hadoop-core-1.0.4.jar \*.java**
- CountJob\$ **jar cvf CountJob.jar \*.class**

Creates a jar file called **CountJob.jar** to be used for running the Count program in Hadoop

# Running the CountJob program in Hadoop

- Perform following task in the folder containing Hadoop installation

```
hadoop $ bin/start-all.sh
```

Starts the Hadoop services

```
hadoop $ ssh localhost
```

Establishes a secure  
connection with the machine

```
hadoop $ mkdir input
```

Creates the directory to put input.txt file

- Copy the CountJob.jar and the input.txt file into the Hadoop installation directory

```
hadoop $ ls
```

CountJob.jar

Input.txt

.....

```
hadoop-examples-1.0.4.jar  
hadoop-minicluster-1.0.4.jar  
hadoop-test-1.0.4.jar  
hadoop-tools-1.0.4.jar
```

- Now prepare the input for the CountJob program:

Creates a directory called input in the HDFS

```
hadoop$ bin/hadoop dfs -mkdir input
```

```
hadoop$ bin/hadoop dfs -put input.txt input
```

Copy input.txt into the input folder in HDFS

- Finally execute the following commands:

Run the CountJob program in Hadoop

```
hadoop$ bin/hadoop jar CountJob.jar CountJob input output
```

HDFS folder where the input files reside

HDFS folder that will contain the output files

Copies the **output** folder from  
HDFS to local machine

```
hadoop$ bin/hadoop dfs -get output output
```

- ❖ The result of the CountJob program in a file (probably) called part-00000

## **Step - IV**

# Assignment-1:

- Write Character count mapper, reducer and driver classes using Hadoop APIs for the sample input file as shown below:

**SampleInputFile.txt**

```
l m n m n l l m m l l n n m l l l m n l m l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
l m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
l m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n  
m n m n l l m m l l n n m l l l m n l m l m n l m m m n l m n n l m n l m n l n n m l n m m m l l m n m n m l n m n m l m n l m n l l n m n
```

**Sample Output format**

I	236
m	111
n	134

## Assignment-2:

- Write a MapReduce program using Hadoop APIs for following task:  
Consider an input file containing integer numbers.  
Provide this file as an input to the program and the program finds maximum number of two consecutive columns and produces output.

SampleInputFile.txt

```
8,4,5,6,10,9,2,3,5,6  
1,2,4,3,7,6,8,9,4,5  
3,5,2,3,7,5,8,9,9,0  
. . . . .  
5,7,7,4,8,6,5,7,9,0
```

Sample Output format

```
1st&2ndColumn 8  
2nd&3rdColumn 6  
. . . .  
9th&10thColumn 6
```

Thank You!