

Linux Kernel Module Programming

Kernel Modules

- Kernel modules are piece of code, that can be loaded and unloaded from kernel on demand.
- Kernel modules offers an easy way to extend the functionality of the base kernel without having to rebuild or recompile the kernel again. Most of the drivers are implemented as a Linux kernel modules. When those drivers are not needed, we can unload only that specific driver, which will reduce the kernel image size.
- The kernel modules will have a **.ko** extension. On a normal linux system, the kernel modules will reside inside
`/lib/modules/<kernel_version>/kernel/` directory.

Examples:

- Typical modules:
 - Device drivers
 - File system drivers
 - System calls

Advantages

- There is no necessity to rebuild the kernel, when a new kernel option is added.
- Modules help find system problems (if system problem caused a module just don't load it).
- Modules are much faster to maintain and debug.
- Modules once loaded are in as much fast as kernel.

Utilities to manipulate kernel modules

- **lsmod**
- **insmod**
- **modinfo**
- **rmmmod**
- **modprobe**

lsmod

- Lists modules that are loaded already.

```
tushar@tushar-laptop ~ $ lsmod
Module                Size  Used by
nls_utf8              12493   1
udf                   83847   1
nls_iso8859_1         12617   1
crc_itu_t             12627   1 udf
rfcomm                53664   8
bnep                  18895   2
binfmt_misc           13140   1
hid_generic           12492   0
dm_multipath          22402   0
scsi_dh               14458   1 dm_multipath
usbhid                47070   0
hid                   87604   2 hid_generic,usbhid
uvcvideo              71309   0
videobuf2_vmalloc     13048   1 uvcvideo
videobuf2_memops      13170   1 videobuf2_vmalloc
videobuf2_core        39258   1 uvcvideo
```

modinfo

- Display module information.

```
tushar@tushar-laptop ~ $ modinfo usb_storage
filename:      /lib/modules/3.13.0-37-generic/kernel/drivers/usb/storage/usb-storage.ko
license:      GPL
description:   USB Mass Storage driver for Linux
author:       Matthew Dharm <mdharm-usb@one-eyed-alien.net>
srcversion:   13955DAA5B7302244B5FD1E
alias:        usb:v*p*d*dc*dsc*dp*ic08isc06ip50in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc05ip50in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc04ip50in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc03ip50in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc02ip50in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc01ip50in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc06ip00in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc05ip00in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc04ip00in*
alias:        usb:v*p*d*dc*dsc*dp*ic08isc03ip00in*
```

insmod

- Insert module into kernel.
- Syntax:
-insmod <module_name>.ko

rmmod

- Removes module from kernel. You cannot remove a module which is already used by any program.
- Syntax:
-rmmod <module_name>.ko

Kernel Module Implementation

- The kernel considers only modules that have been loaded into RAM by the *insmod* program and for each of them allocates memory area containing:
 - A module object.
 - Null terminated string that represents module's name.
 - The code that implements the functions of the module.

Linux Kernel Module Programming

- Write a **hello_proc.c program**
- Create a Makefile
- The program and Makefile should be kept in a single folder.
- Change directory to this folder and execute following:
 - **make**
 - **insmod hello_proc.ko**
 - **dmesg (see the kernel buffer contents, reads the kernel log file /var/log/syslog)**
 - **lsmod**
 - **cat /proc/hello_proc**

Files created after building the module

- **hello.o**
 - Module object file before linking.
- **hello.mod.c**
 - Contains module's information.
- **hello.mod.o**
 - After compilation and linking of hello.mod.c.
- **modules.order**
 - The order in which two or three modules get linked.
- **Modules.symvers**
 - Symbol versions if any.
- **hello.ko**
 - A module kernel object file after linking hello.o and hello.mod.o

Example: hello.c

```
#include <linux/module.h> // included for all kernel modules
#include <linux/kernel.h> // included for KERN_INFO
#include <linux/init.h> // included for __init and __exit macros
MODULE_LICENSE("GPL");
MODULE_AUTHOR("Tushar B Kute");
MODULE_DESCRIPTION("A Simple Hello World module");
static int __init hello_init(void)
{
    printk(KERN_INFO "Hello world!\n");
    return 0; // Nonzero
    return means that the module couldn't be
    loaded.
}
static void __exit hello_cleanup(void)
{
    printk(KERN_INFO "Good Bye.\n");
}
module_init(hello_init);
module_exit(hello_cleanup);
```

Makefile

- **Obj-m += hello.o**

all:

**make C /lib/modules/\$(shell
uname r)/build M=\$(PWD)modules**

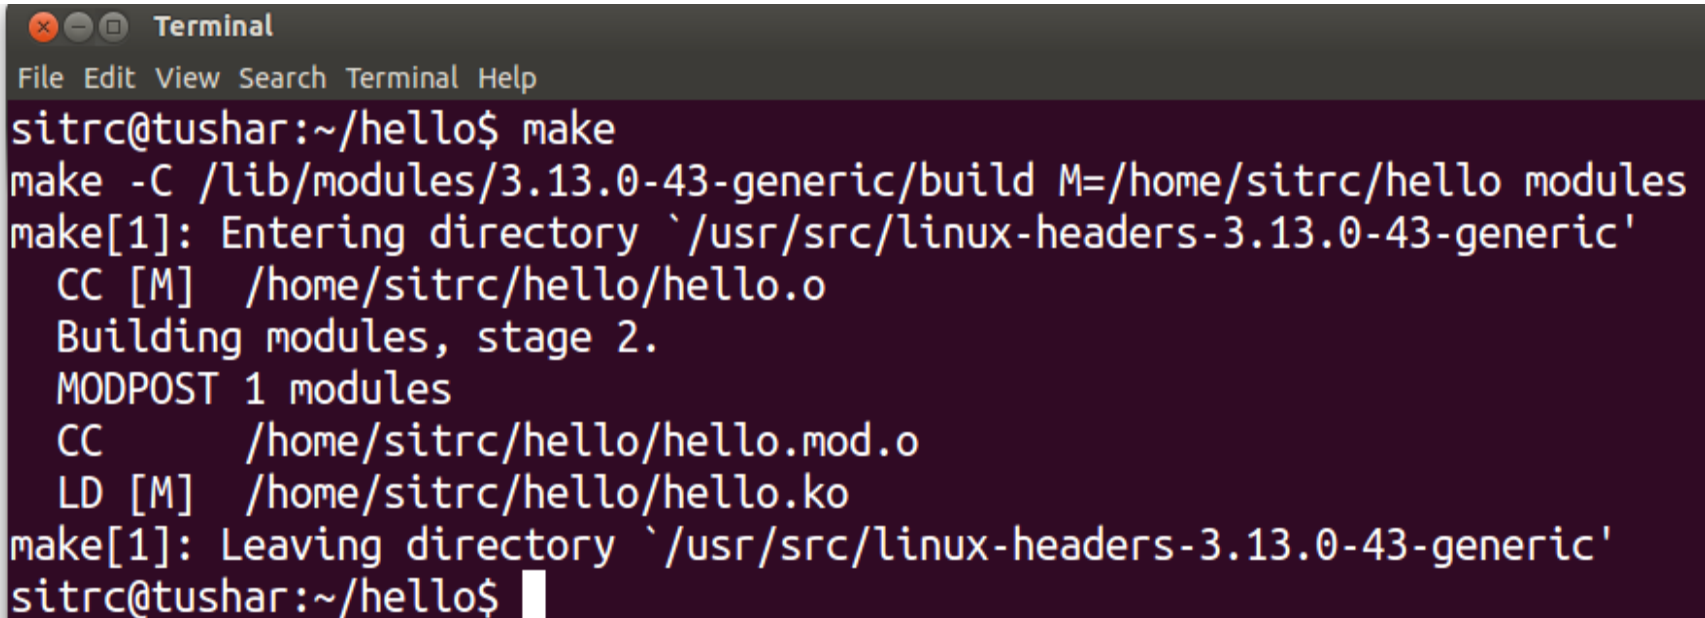
clean:

**make C/lib/modules/\$(shell uname
r)/build M=\$(PWD)**

clean

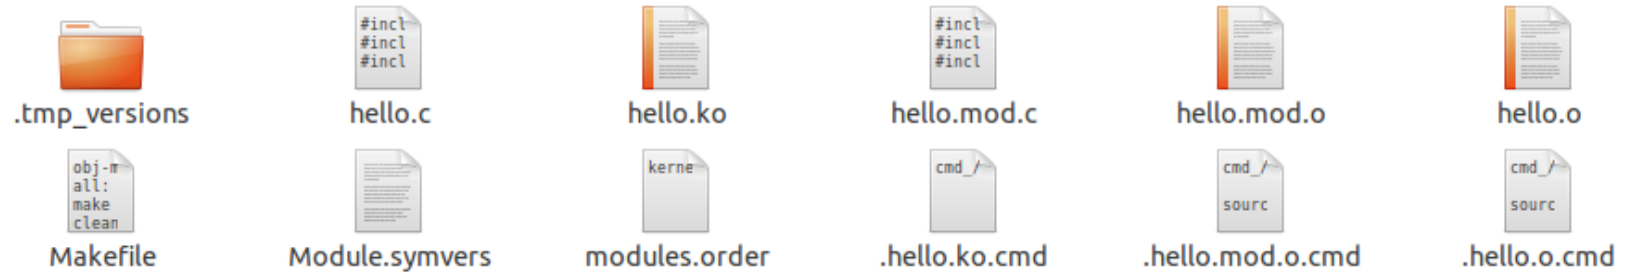
Compile the program

- **make**



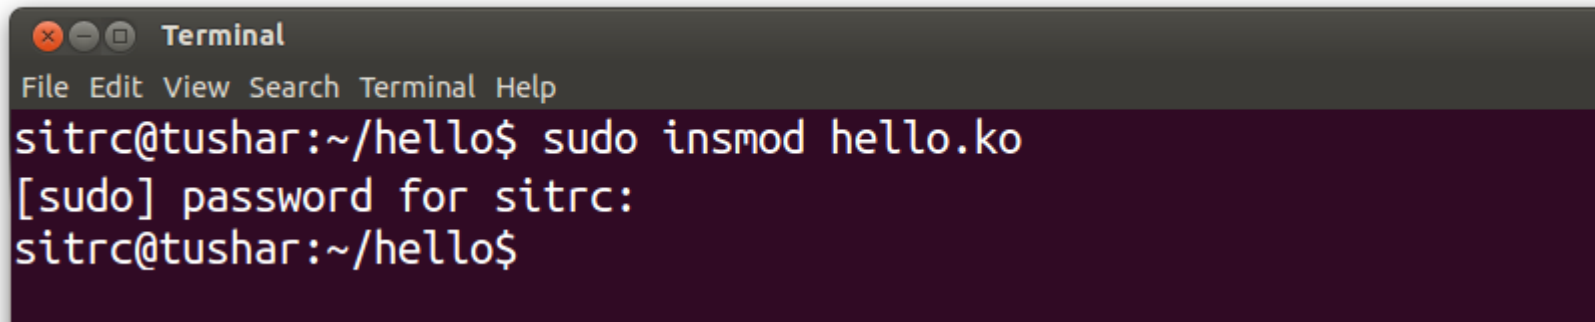
```
Terminal
File Edit View Search Terminal Help
sitrc@tushar:~/hello$ make
make -C /lib/modules/3.13.0-43-generic/build M=/home/sitrc/hello modules
make[1]: Entering directory `/usr/src/linux-headers-3.13.0-43-generic'
  CC [M]  /home/sitrc/hello/hello.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC      /home/sitrc/hello/hello.mod.o
  LD [M]  /home/sitrc/hello/hello.ko
make[1]: Leaving directory `/usr/src/linux-headers-3.13.0-43-generic'
sitrc@tushar:~/hello$
```

File generated



Insert the module

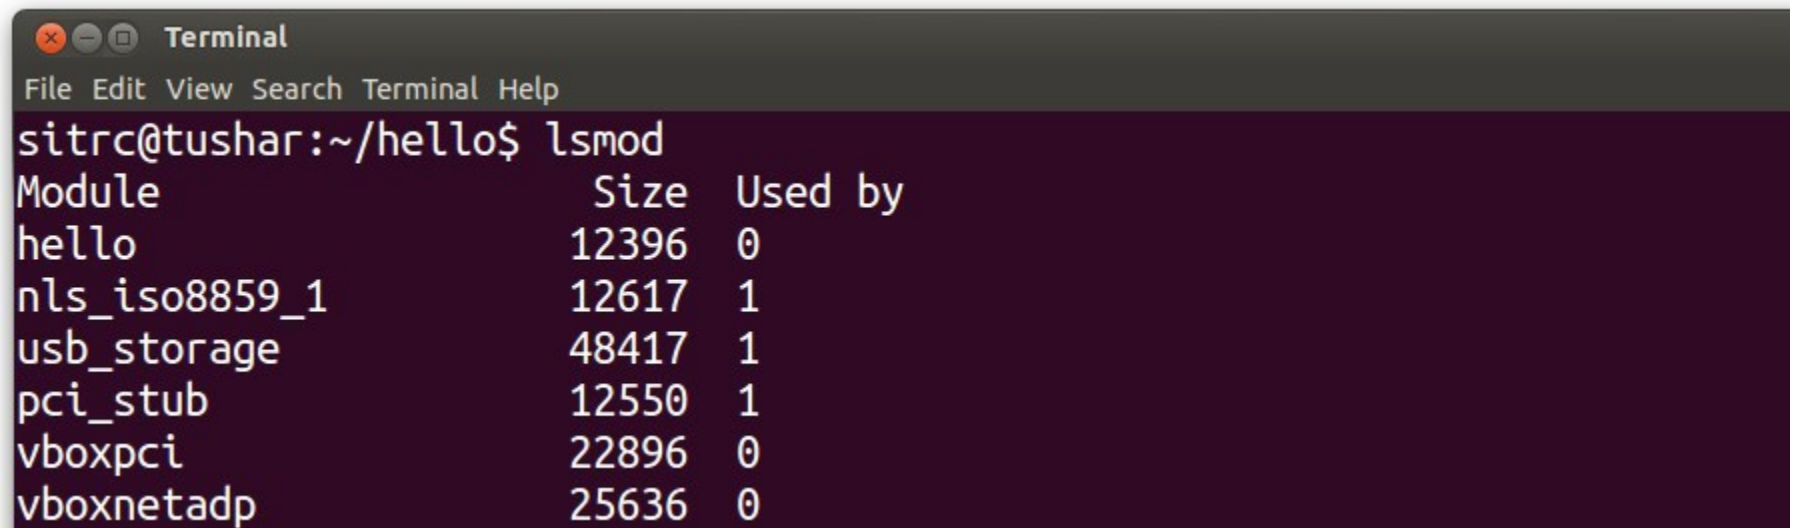
- **insmod hello.ko**

A screenshot of a macOS Terminal window. The title bar at the top says "Terminal" with standard window control buttons (close, zoom, full screen). Below the title bar is a menu bar with "File", "Edit", "View", "Search", "Terminal", and "Help". The main area of the terminal has a dark purple background and white text. It shows a user named "sitrc" at a machine named "tushar" in the directory "~/hello". The user has entered the command "sudo insmod hello.ko". The terminal shows the prompt "[sudo] password for sitrc:" followed by the command being executed. The prompt then returns to "sitrc@tushar:~/hello\$".

```
Terminal
File Edit View Search Terminal Help
sitrc@tushar:~/hello$ sudo insmod hello.ko
[sudo] password for sitrc:
sitrc@tushar:~/hello$
```

Check the module in the list

- **lsmod**



A terminal window titled "Terminal" with a menu bar containing "File", "Edit", "View", "Search", "Terminal", and "Help". The prompt is "sitrc@tushar:~/hello\$". The command "lsmod" has been executed, resulting in a table of loaded kernel modules.

Module	Size	Used by
hello	12396	0
nls_iso8859_1	12617	1
usb_storage	48417	1
pci_stub	12550	1
vboxpci	22896	0
vboxnetadp	25636	0

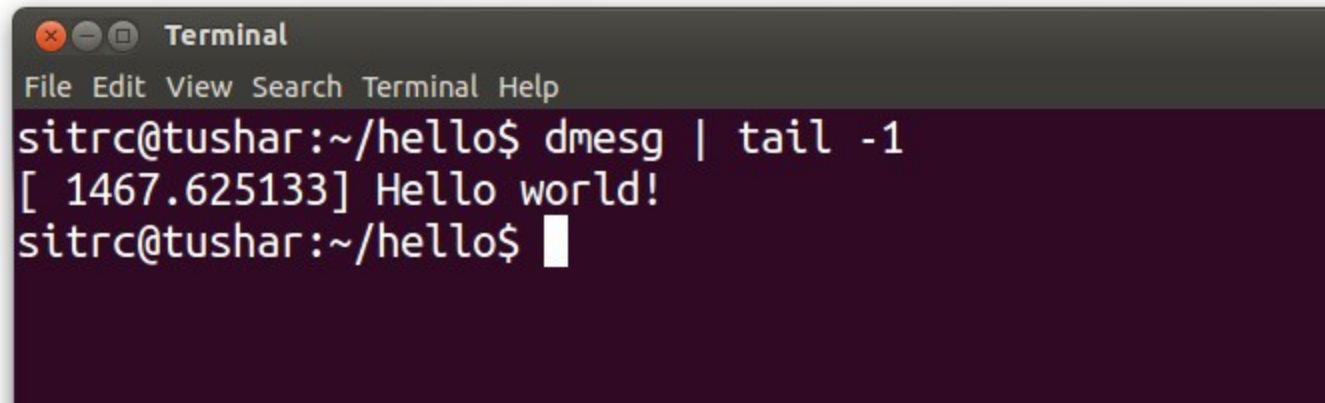
Check kernel ring buffer

- **dmesg**

```
[ 340.234241] sd 4:0:0:0: [sdb] Write cache: disabled, read cache: enabled, do
[ 340.239509]   sdb: sdb1
[ 340.243618] sd 4:0:0:0: [sdb] Attached SCSI removable disk
[ 341.180282] FAT-fs (sdb1): Volume was not properly unmounted. Some data may
.
[ 344.999200] systemd-hostnamed[2932]: Warning: nss-myhostname is not installe
ame might make it unresolveable. Please install nss-myhostname!
[ 672.696406] systemd-hostnamed[3458]: Warning: nss-myhostname is not installe
ame might make it unresolveable. Please install nss-myhostname!
[ 1017.496450] perf samples too long (2523 > 2500), lowering kernel.perf_event_
[ 1467.625133] Hello world!
sitrc@tushar:~/hello$
```

Check kernel ring buffer

- **dmesg | tail 1**



```
Terminal
File Edit View Search Terminal Help
sitrc@tushar:~/hello$ dmesg | tail -1
[ 1467.625133] Hello world!
sitrc@tushar:~/hello$
```

Removing the module

- **rmmod hello.ko**
- **dmesg or**
- **dmesg | tail 1**

What is dmesg?

- The dmesg command is used to write the kernel messages in Linux and other Unix-like operating systems to standard output (which by default is the display screen).
- dmesg obtains its data by reading the kernel ring buffer. A buffer is a portion of a computer's memory that is set aside as a temporary holding place for data that is being sent to or received from an external device, such as a hard disk drive (HDD), printer or keyboard.
- A ring buffer is a buffer of fixed size for which any new data added to it overwrites the oldest data in it.

What is printk?

- The kernel print function, `printk()`, behaves almost identically to the C library `printf()` function.
- `printk()` is simply the name of the kernel's formatted print function. It is callable from just about anywhere in the kernel at any time.
- The major difference between `printk()` and `printf()` is the capability of the former to specify a loglevel.
- The kernel uses the loglevel to decide whether to print the message to the console. The kernel displays all messages with a loglevel below a specified value on the console.

printk : example

- `printk(KERN_WARNING "This is a warning!\n");`
- `printk(KERN_DEBUG "This is a debug notice!\n");`
- `printk("I did not specify a loglevel!\n");`